Create Filegroup

```
CREATE DATABASE [Hekaton]
ON PRIMARY
(NAME = N'Hekaton_data', FILENAME = N'C:\Data\Hekaton_data.mdf'),
FILEGROUP [Hekaton_InMemory] CONTAINS MEMORY_OPTIMIZED_DATA
(NAME = N'Hekaton_mem', FILENAME = N'C:\Data\Mem\Hekaton_Lun1')
LOG ON
(NAME = N'Hekaton_log', FILENAME = N'C:\Data\Log\Hekaton_log.ldf')
```

FileGroup Container

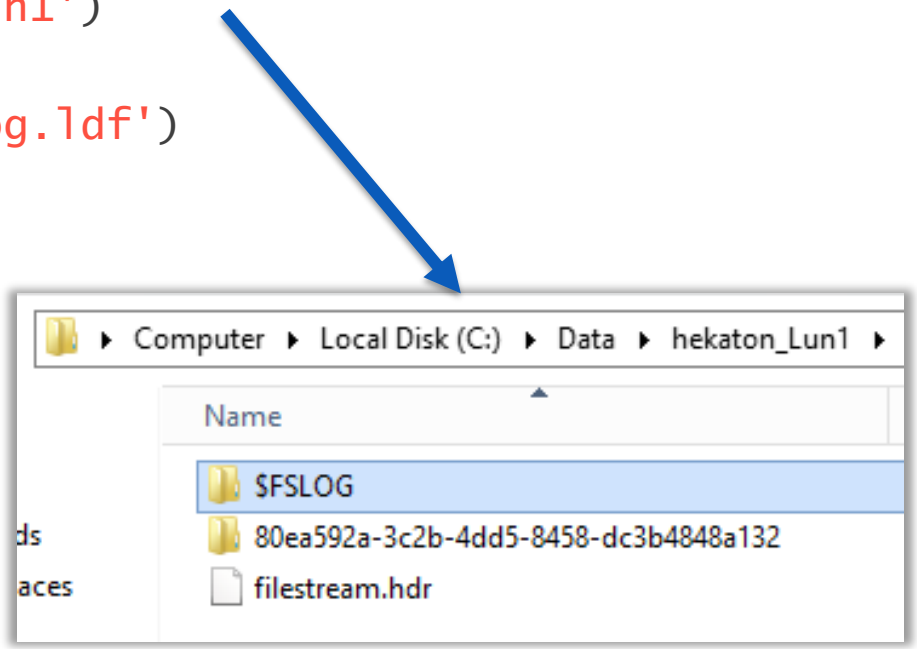


```
ALTER DATABASE [Hekaton]
ADD FILE (NAME = N'Hekaton_mem',
FILENAME = N'C:\Data\Mem\Hekaton_Lun2')
TO FILEGROUP [Hekaton_InMemory]
```

Create Memory Optimized Table

```
CREATE TABLE [Customer] (
 [CustomerID] INT NOT NULL
 PRIMARY KEY NONCLUSTERED HASH WITH (BUCKET_COUNT = 1000000),
 [AddressID] INT NOT NULL INDEX [IxName] HASH WITH (BUCKET_COUNT =
1000000),
 [LName] NVARCHAR(250) COLLATE Latin1_General_100_BIN2 NOT NULL
 INDEX [IXLName] NONCLUSTERED (LName)
)
WITH (MEMORY_OPTIMIZED = ON, DURABILITY = SCHEMA_AND_DATA);
```

Hash Index

Collation BIN2

Range Index

This table is memory optimized

This table is durable

Memory Optimized Table Limitations

Optimized for high-throughput OLTP

- No DML triggers
- No XML and no CLR data types

Optimized for in-memory

- Rows are at most 8060 bytes – no off row data
- No Large Object (LOB) types like varchar(max)

Scoping limitations

- No FOREIGN KEY and no CHECK constraints
- No schema changes (ALTER TABLE) – need to drop/recreate table
- No add/remove index – need to drop/recreate table
- No Computed Columns
- No Cross-Database Queries

Data Access to Memory Optimized Tables

Interpreted T-SQL Access

- Access both memory and disk based tables
- Less performant
- Virtually full T-SQL surface
- When to use
  - Ad-hoc queries
  - Reporting-style queries
  - Speeding up app migration

Natively Compiled Procs

- Access only memory optimized tables
- Maximum performance
- Limited T-SQL surface area
- When to use
  - OLTP-style operations
  - Optimize performance critical business logic
  - More the logic embedded, better the performance improvement

Natively Compiled Procedures

```
CREATE PROCEDURE [dbo].[InsertOrder](@id INT, @date DATETIME)
WITH
NATIVE_COMPILATION,
SCHEMABINDING,
EXECUTE AS OWNER
AS
BEGIN ATOMIC
WITH
(TRANSACTION ISOLATION LEVEL = SNAPSHOT,
LANGUAGE = N'us_english')
        -- Insert T-SQL here…
END
```

This proc is natively compiled

Native procs must be schema-bound

Execution context is required

Atomic blocks

- Create a transaction if there is none
- Otherwise, create a savepoint

Session settings are fixed at create time

Natively Compiled Procedure Performance

To avoid the server having to map parameter names and convert types:

- Use ordinal (nameless) parameters, do not use named parameters
- Match parameter types passed.
- Use Extended Event ‘natively_compiled_proc_slow_parameter_passing’ to identify.

name	timestamp	parameter_name	reason	sql_text
natively_compiled_proc_slow_parameter_passing	2013-12-18 14:30:18.8776070	@order_id	named_parameters	DECLARE @order_id BIGINT = 84...
natively_compiled_proc_slow_parameter_passing	2013-12-18 14:31:08.6299601		parameter_conversion	DECLARE @order_id BIGINT = '84...
natively_compiled_proc_slow_parameter_passing	2013-12-18 14:31:15.5519728		parameter_conversion	DECLARE @order_id BIGINT = '84...

Memory Optimized TVP and Table Variables

Usage Scenarios

- Storing Intermediate Result Sets
- Replacement for table variables
- Passing TVP to Natively compiled procedures

Advantages

- Always in memory, never spill to disk
- No Tempdb utilization
- Data access is efficient

Memory Optimized TVP

```
CREATE TYPE InMemTVP AS TABLE
(ProductID INT NOT NULL PRIMARY KEY NONCLUSTERED HASH(ProductID)
WITH (BUCKET_COUNT = 100))
WITH (MEMORY_OPTIMIZED = ON)
GO

CREATE PROCEDURE [dbo].[GetProductDetails_Native](@tvp dbo.InMemTVP READONLY)
WITH
NATIVE_COMPILATION,
SCHEMABINDING,
EXECUTE AS SELF
AS
BEGIN ATOMIC WITH (TRANSACTION ISOLATION LEVEL = SNAPSHOT, LANGUAGE = N'English')
SELECT a.ProductID, Name, ListPrice FROM Production.Product_inmem a
INNER JOIN @tvp b ON a.ProductID = b.ProductID
END
GO
```

Memory Optimized Table Creation

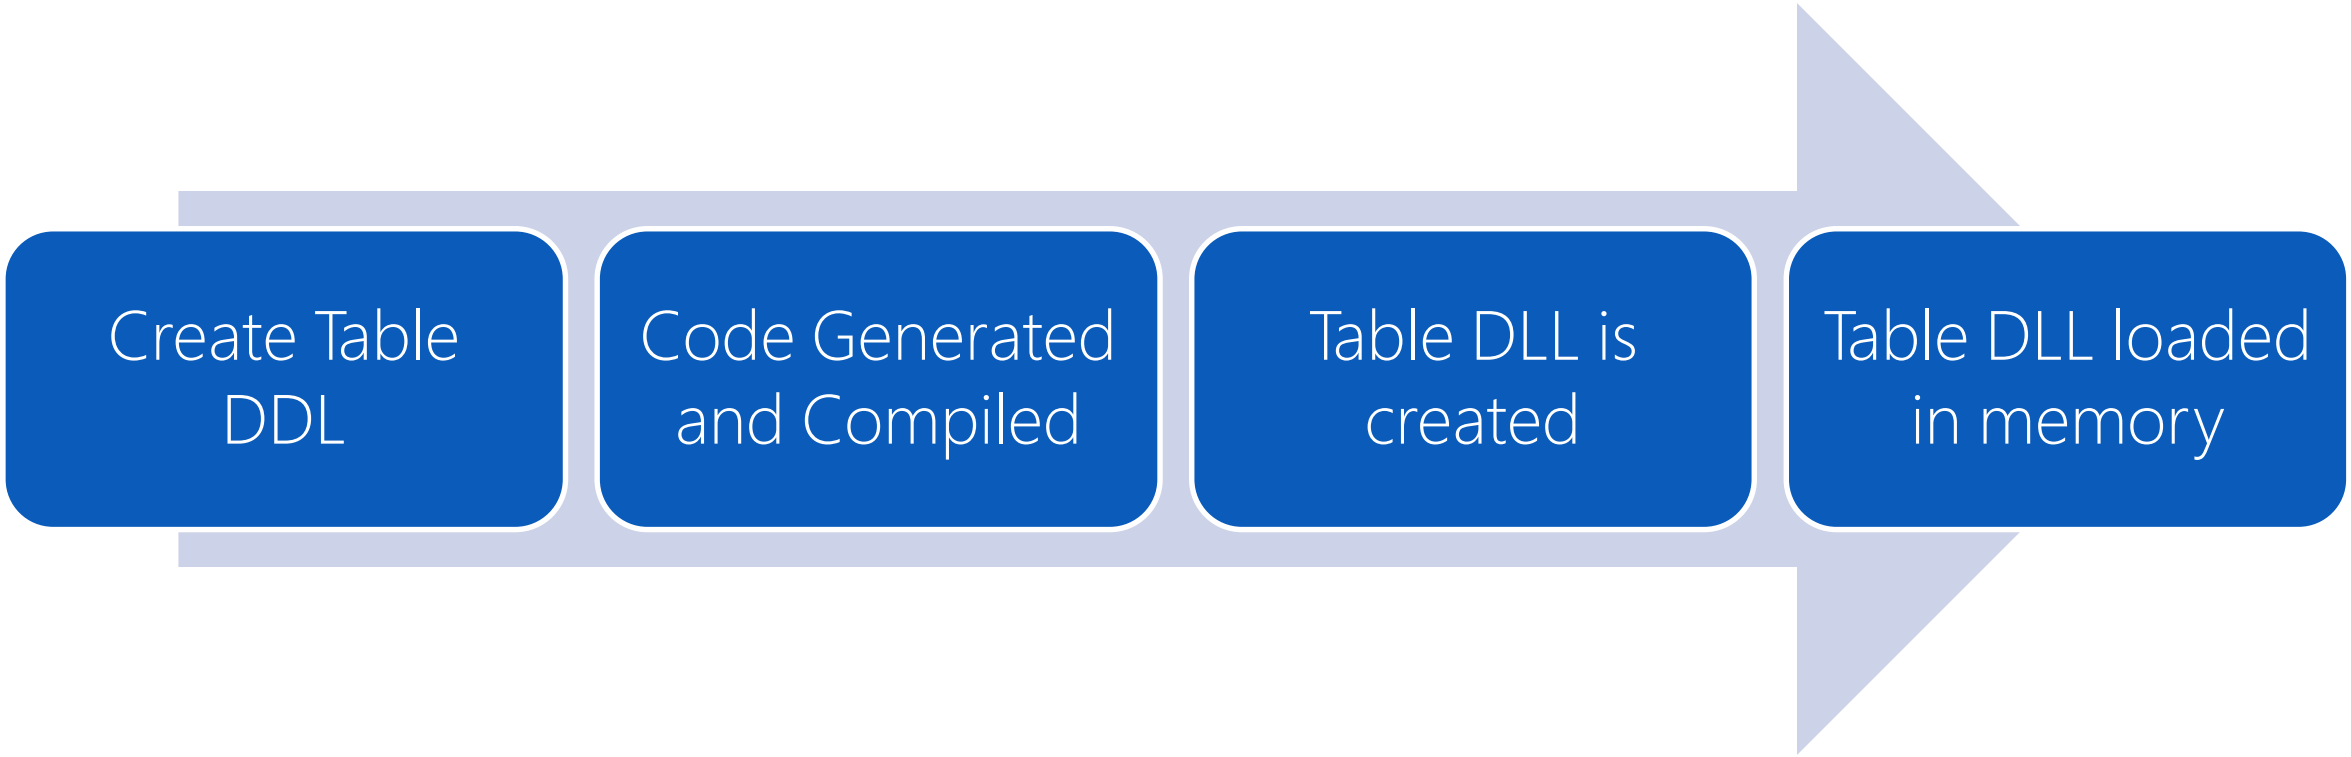

Native Compilation Files

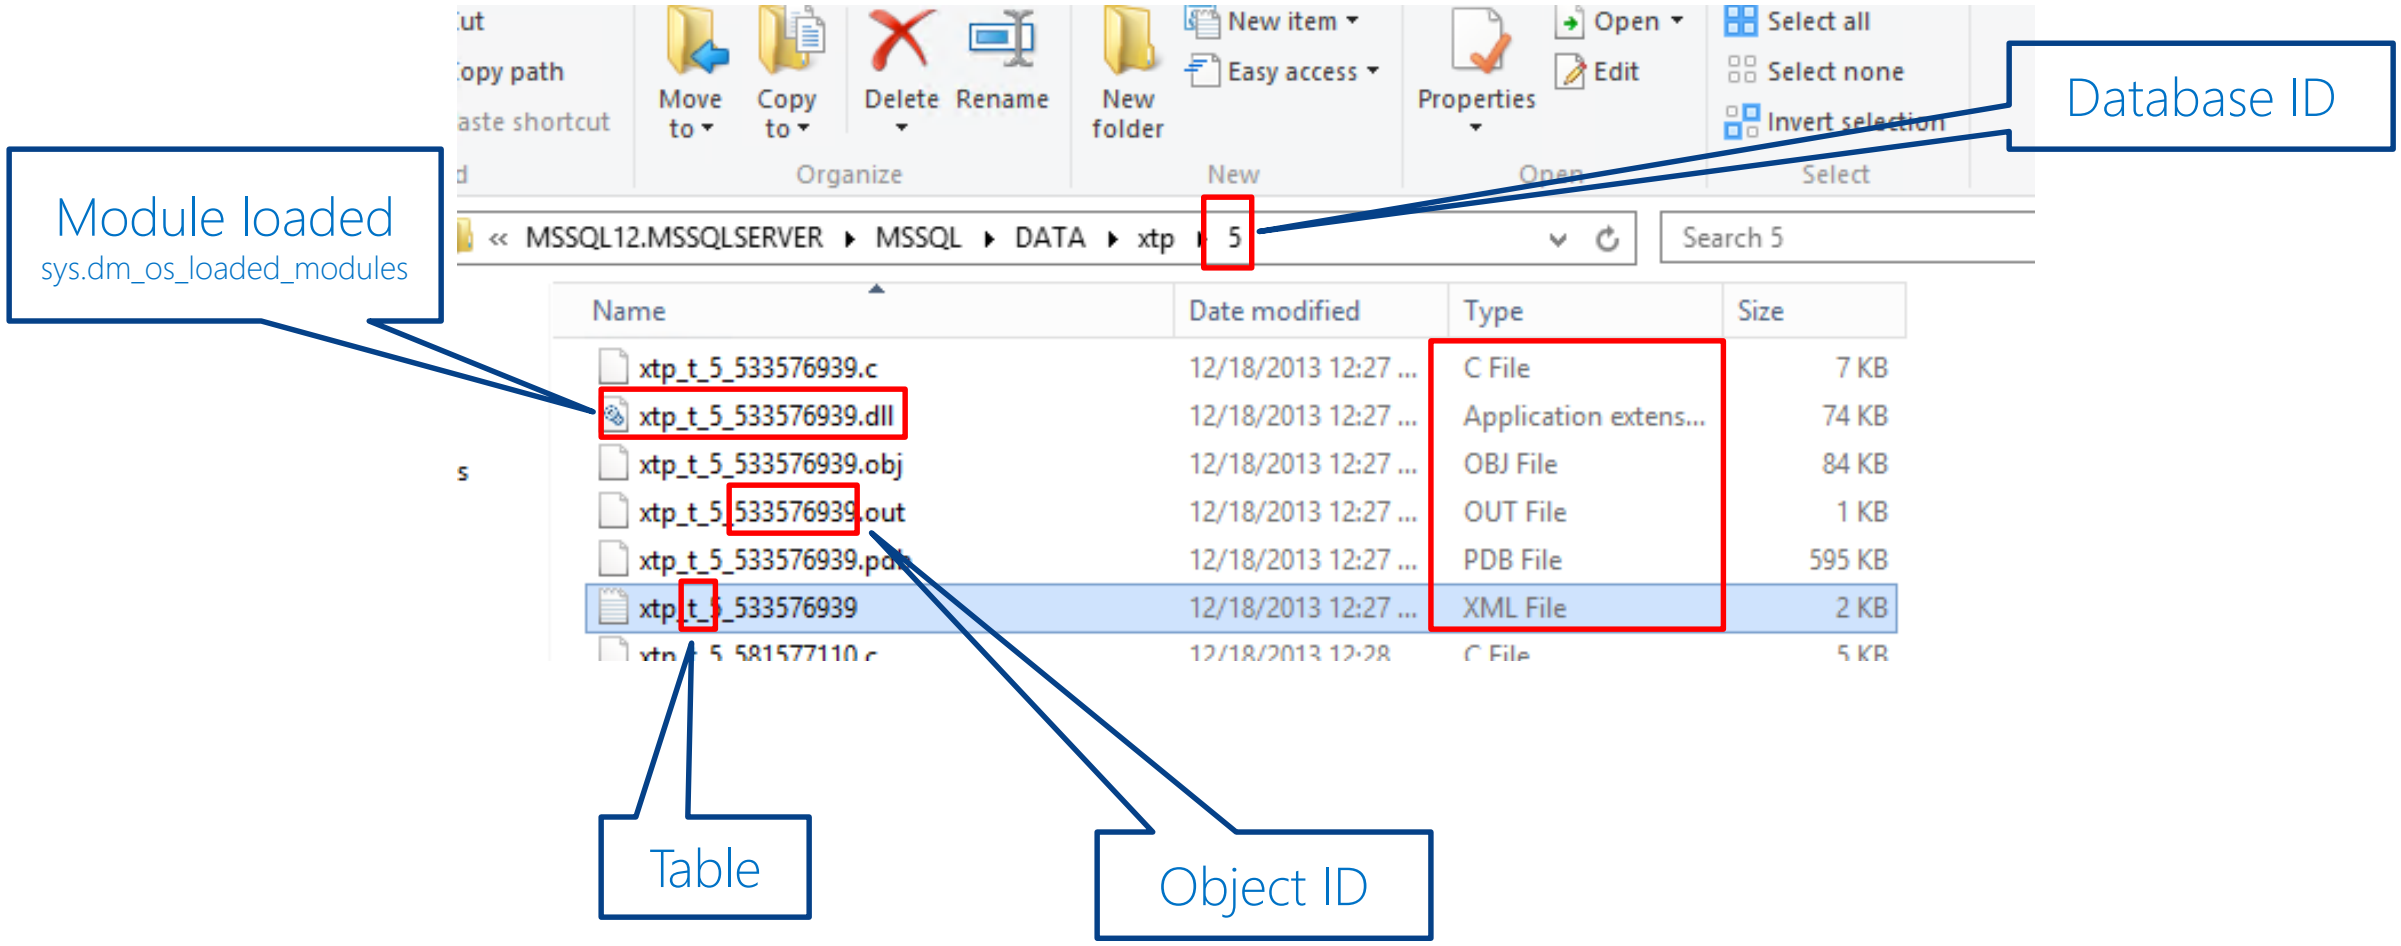

Natively Compiled Procedure

Queries are optimized at procedure create time

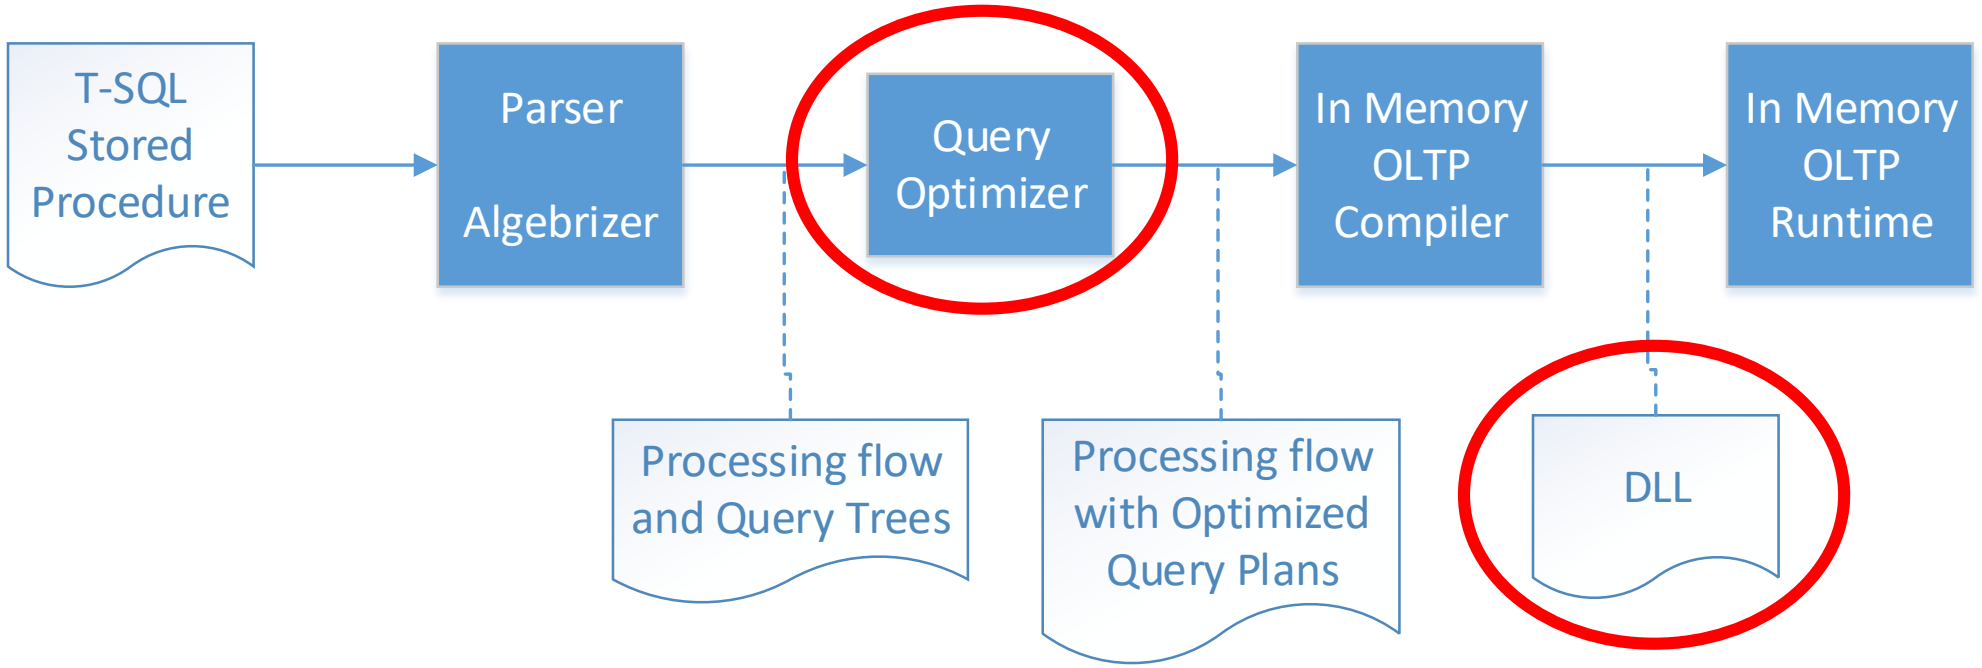

DLLs are **not** part of the plan cache

Compilation Process

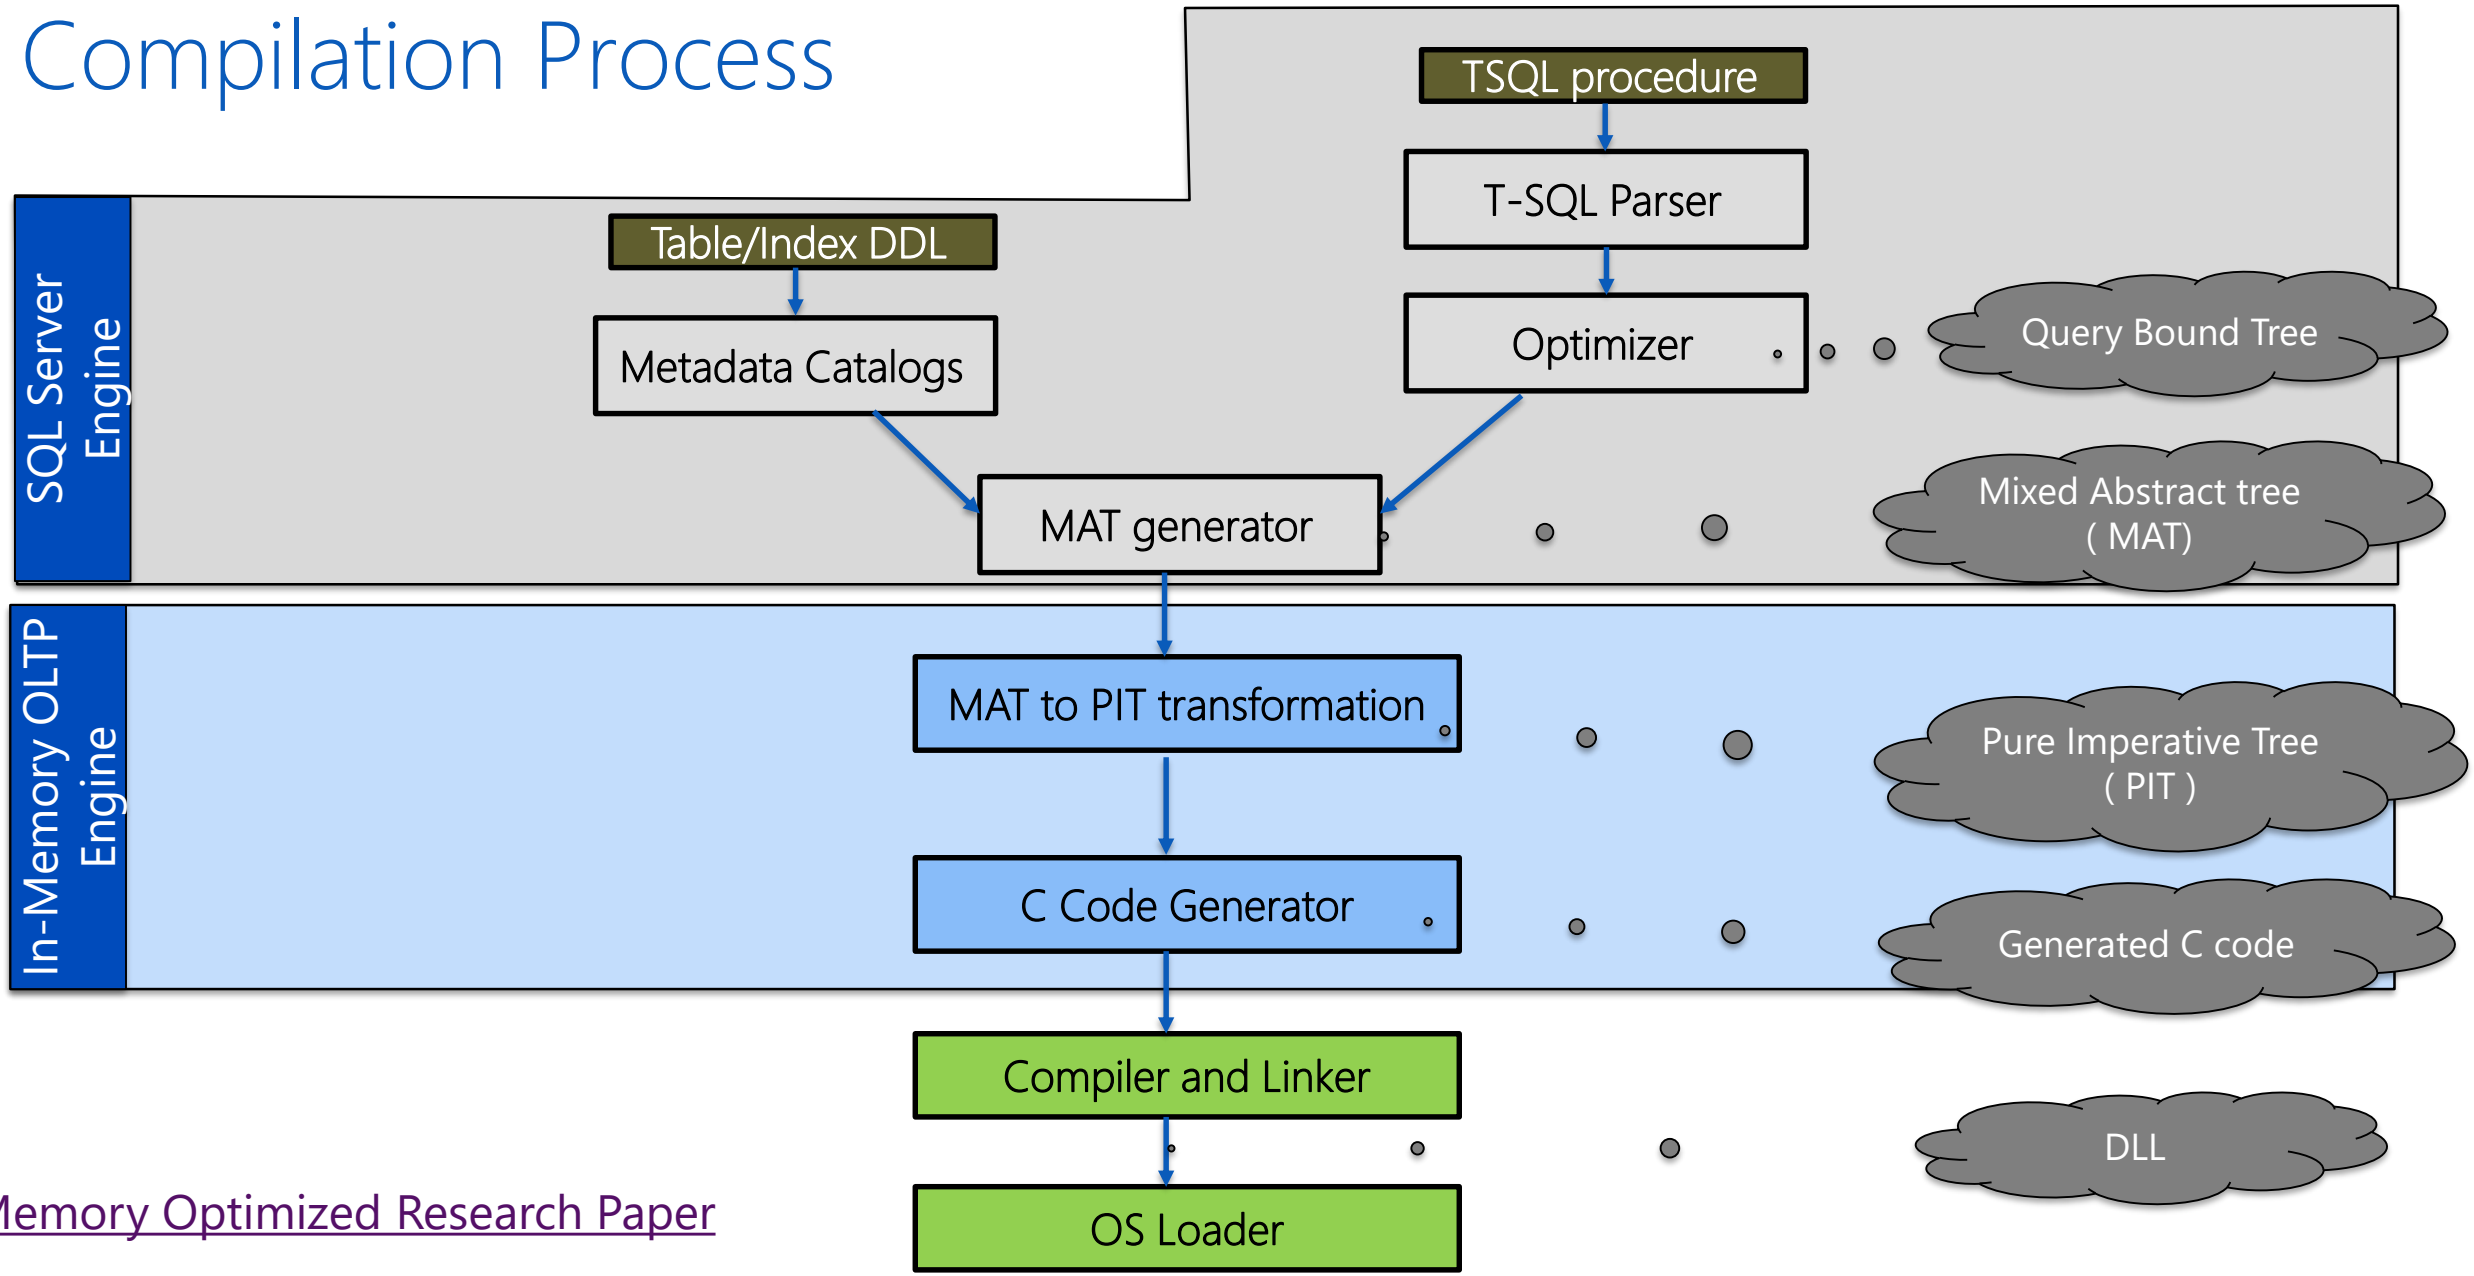

[Memory Optimized Research Paper]

Natively Compiled Procedure Execution

Query operators and functions are baked into DLL

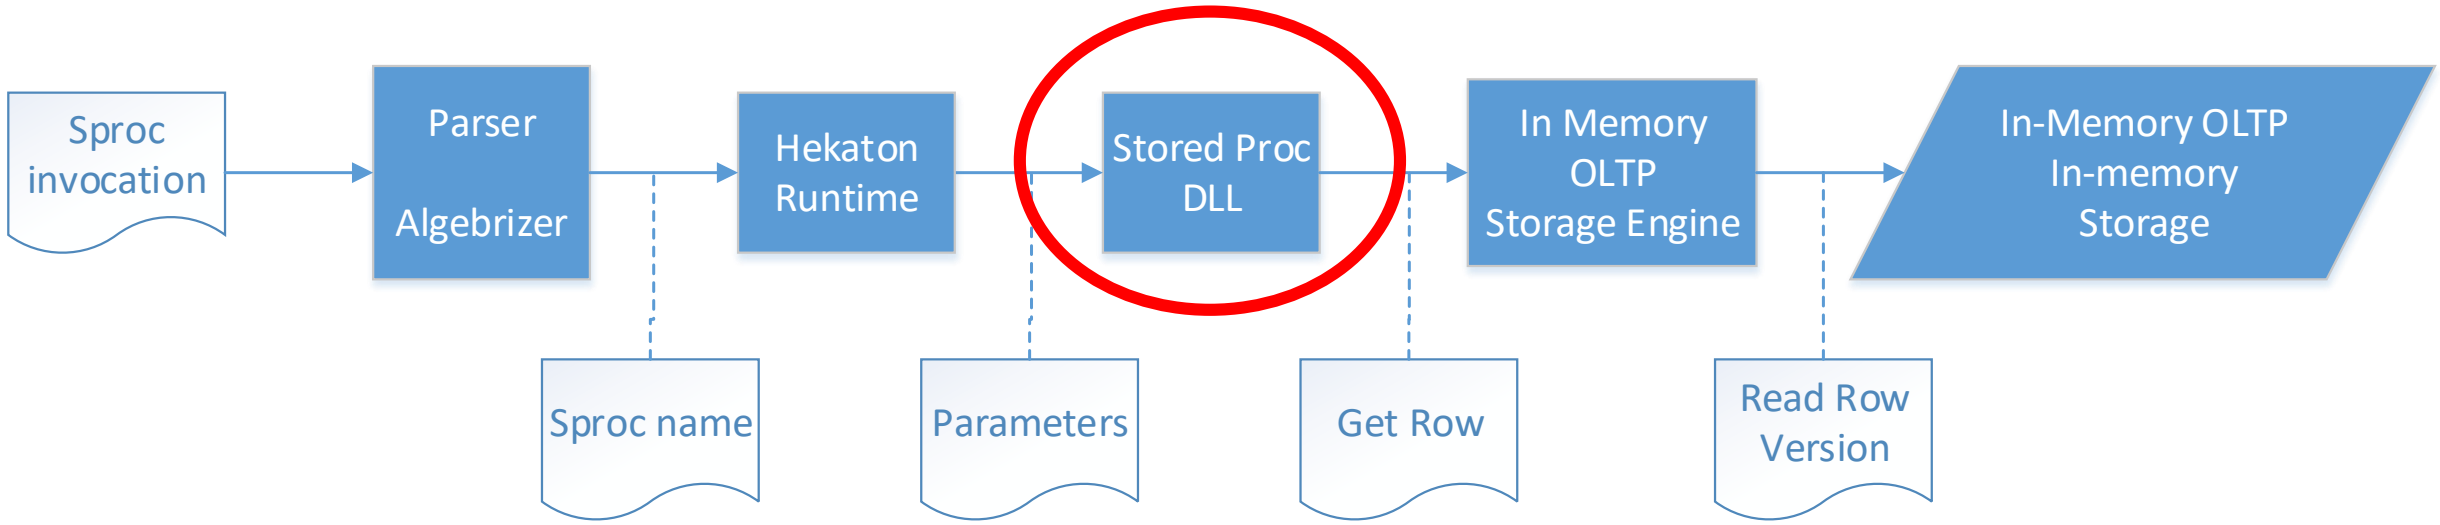

Other Considerations

Parameter sniffing

- For Natively compiled procedures, UNKNOWN parameter value assumed
- Still at play for Interpreted T-SQL

Recompilation for Natively compiled procedures

- No automatic recompilation, requires a drop and recreate
- Recompiled on first execution after restart of the Server or failover

In-Memory OLTP Structures

Rows

- Row structure is optimized for memory access
- There are no Pages
- Rows are versioned and there are no in-place updates

Indexes

- There is no clustered index, only non-clustered indexes
- Indexes point to rows, access to rows is via an index
- Indexes do not exist on disk, only in memory, recreated during recovery
- Hash indexes for point lookups
- Range indexes for ordered scans and range scans

In-Memory Row Format

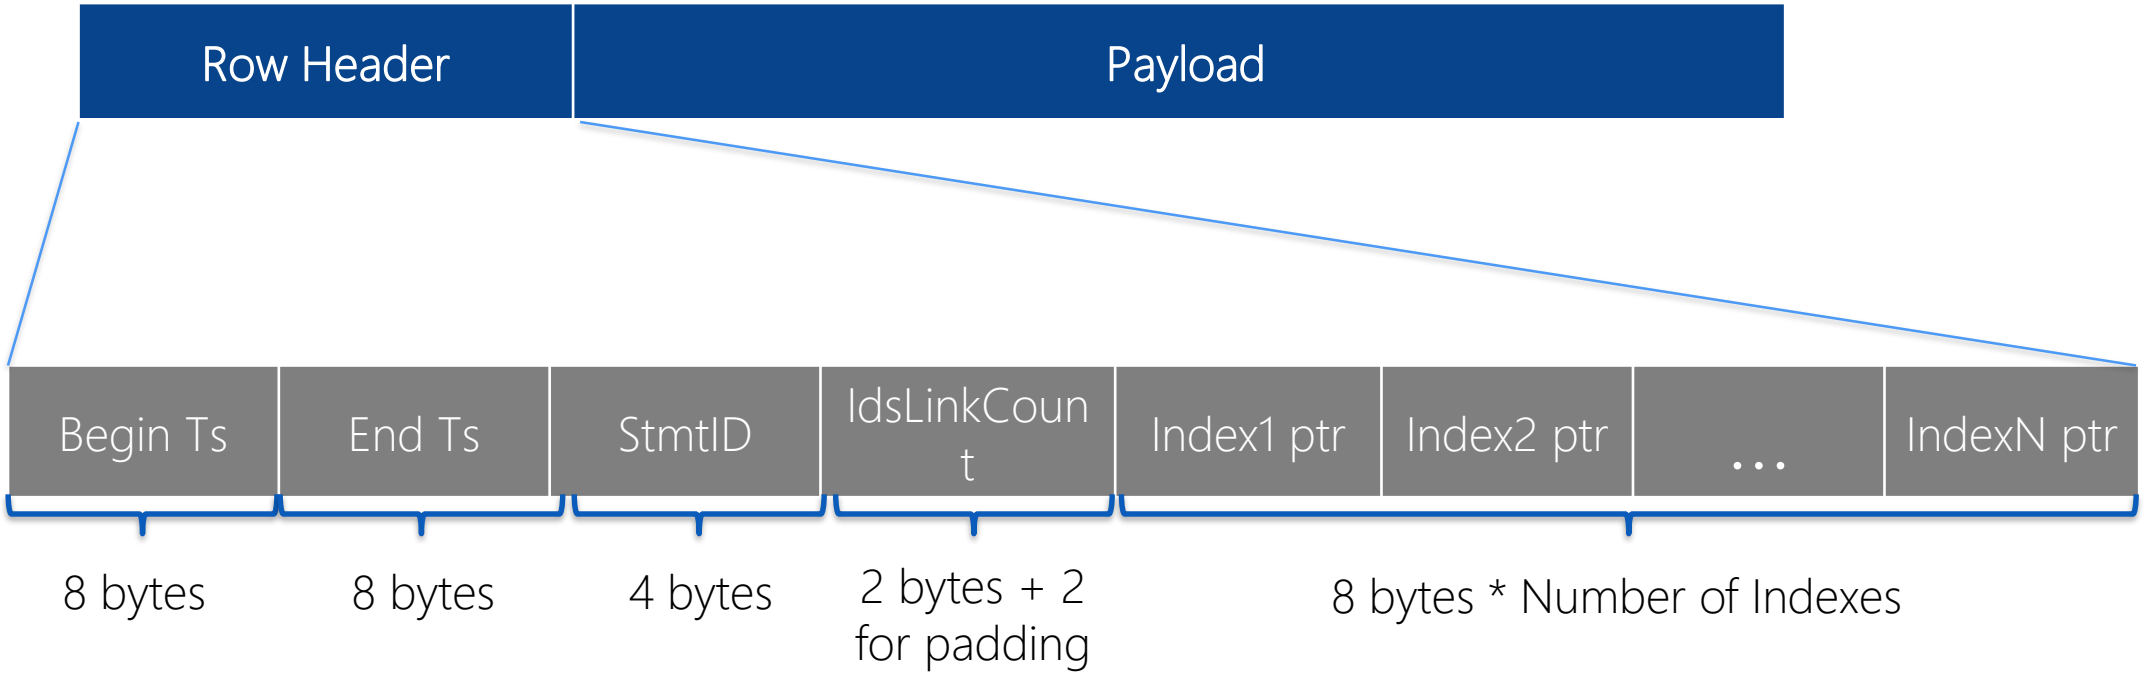

- Begin/End timestamp determines row's version validity and visibility
- No concept of data pages, only rows exist
- Row size limited to 8060 bytes (@table create time) to allow data to be moved to disk-based table
- Not every SQL table schema is supported (Ex: LOB and SqlVariant)

Hash Indexes

Hash index with (bucket_count=8):

Hash function *f*:

- Maps values to buckets
- Built into the system

Hash mapping:

Array of 8-byte Memory pointers

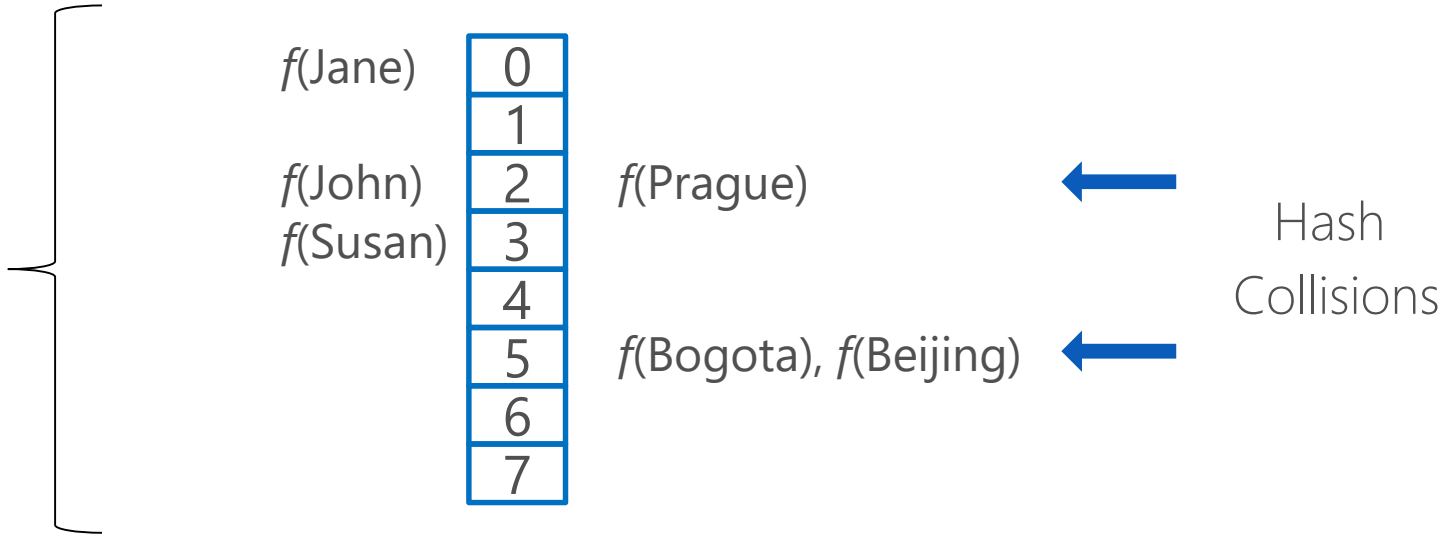

Key Lookup B-Tree vs. Hash Index

Non clustered index

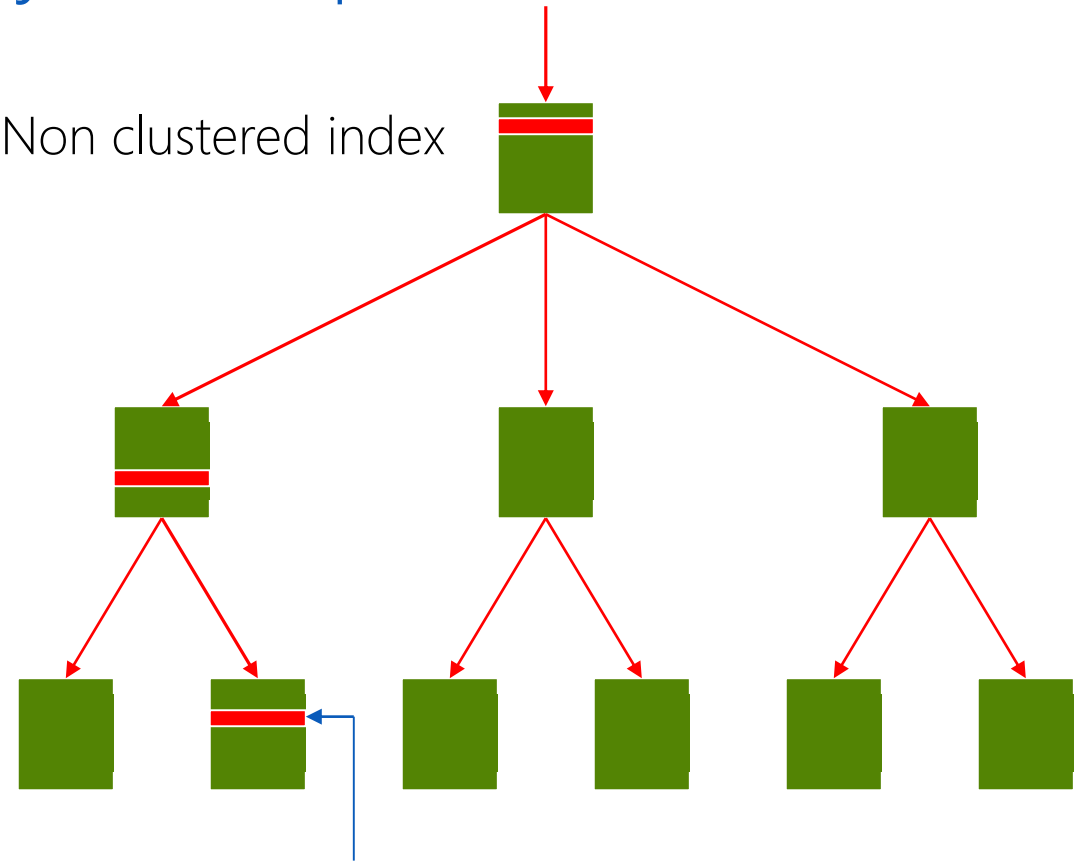

Matching index record

Hash index on Name

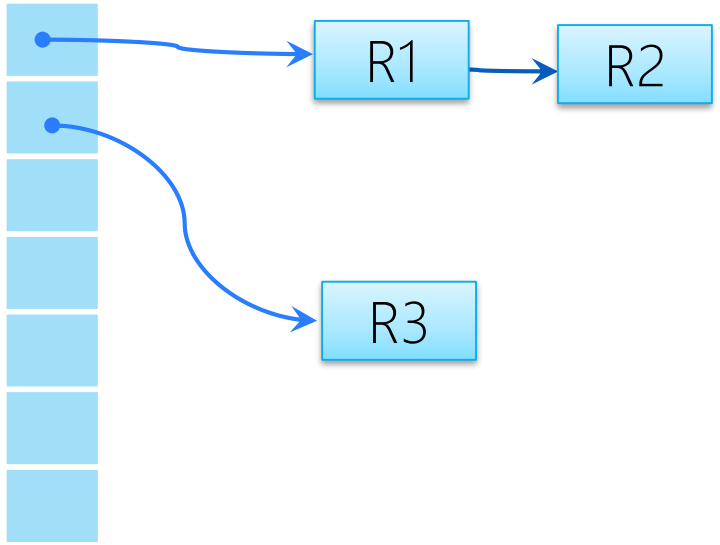

Hash Index Traversal

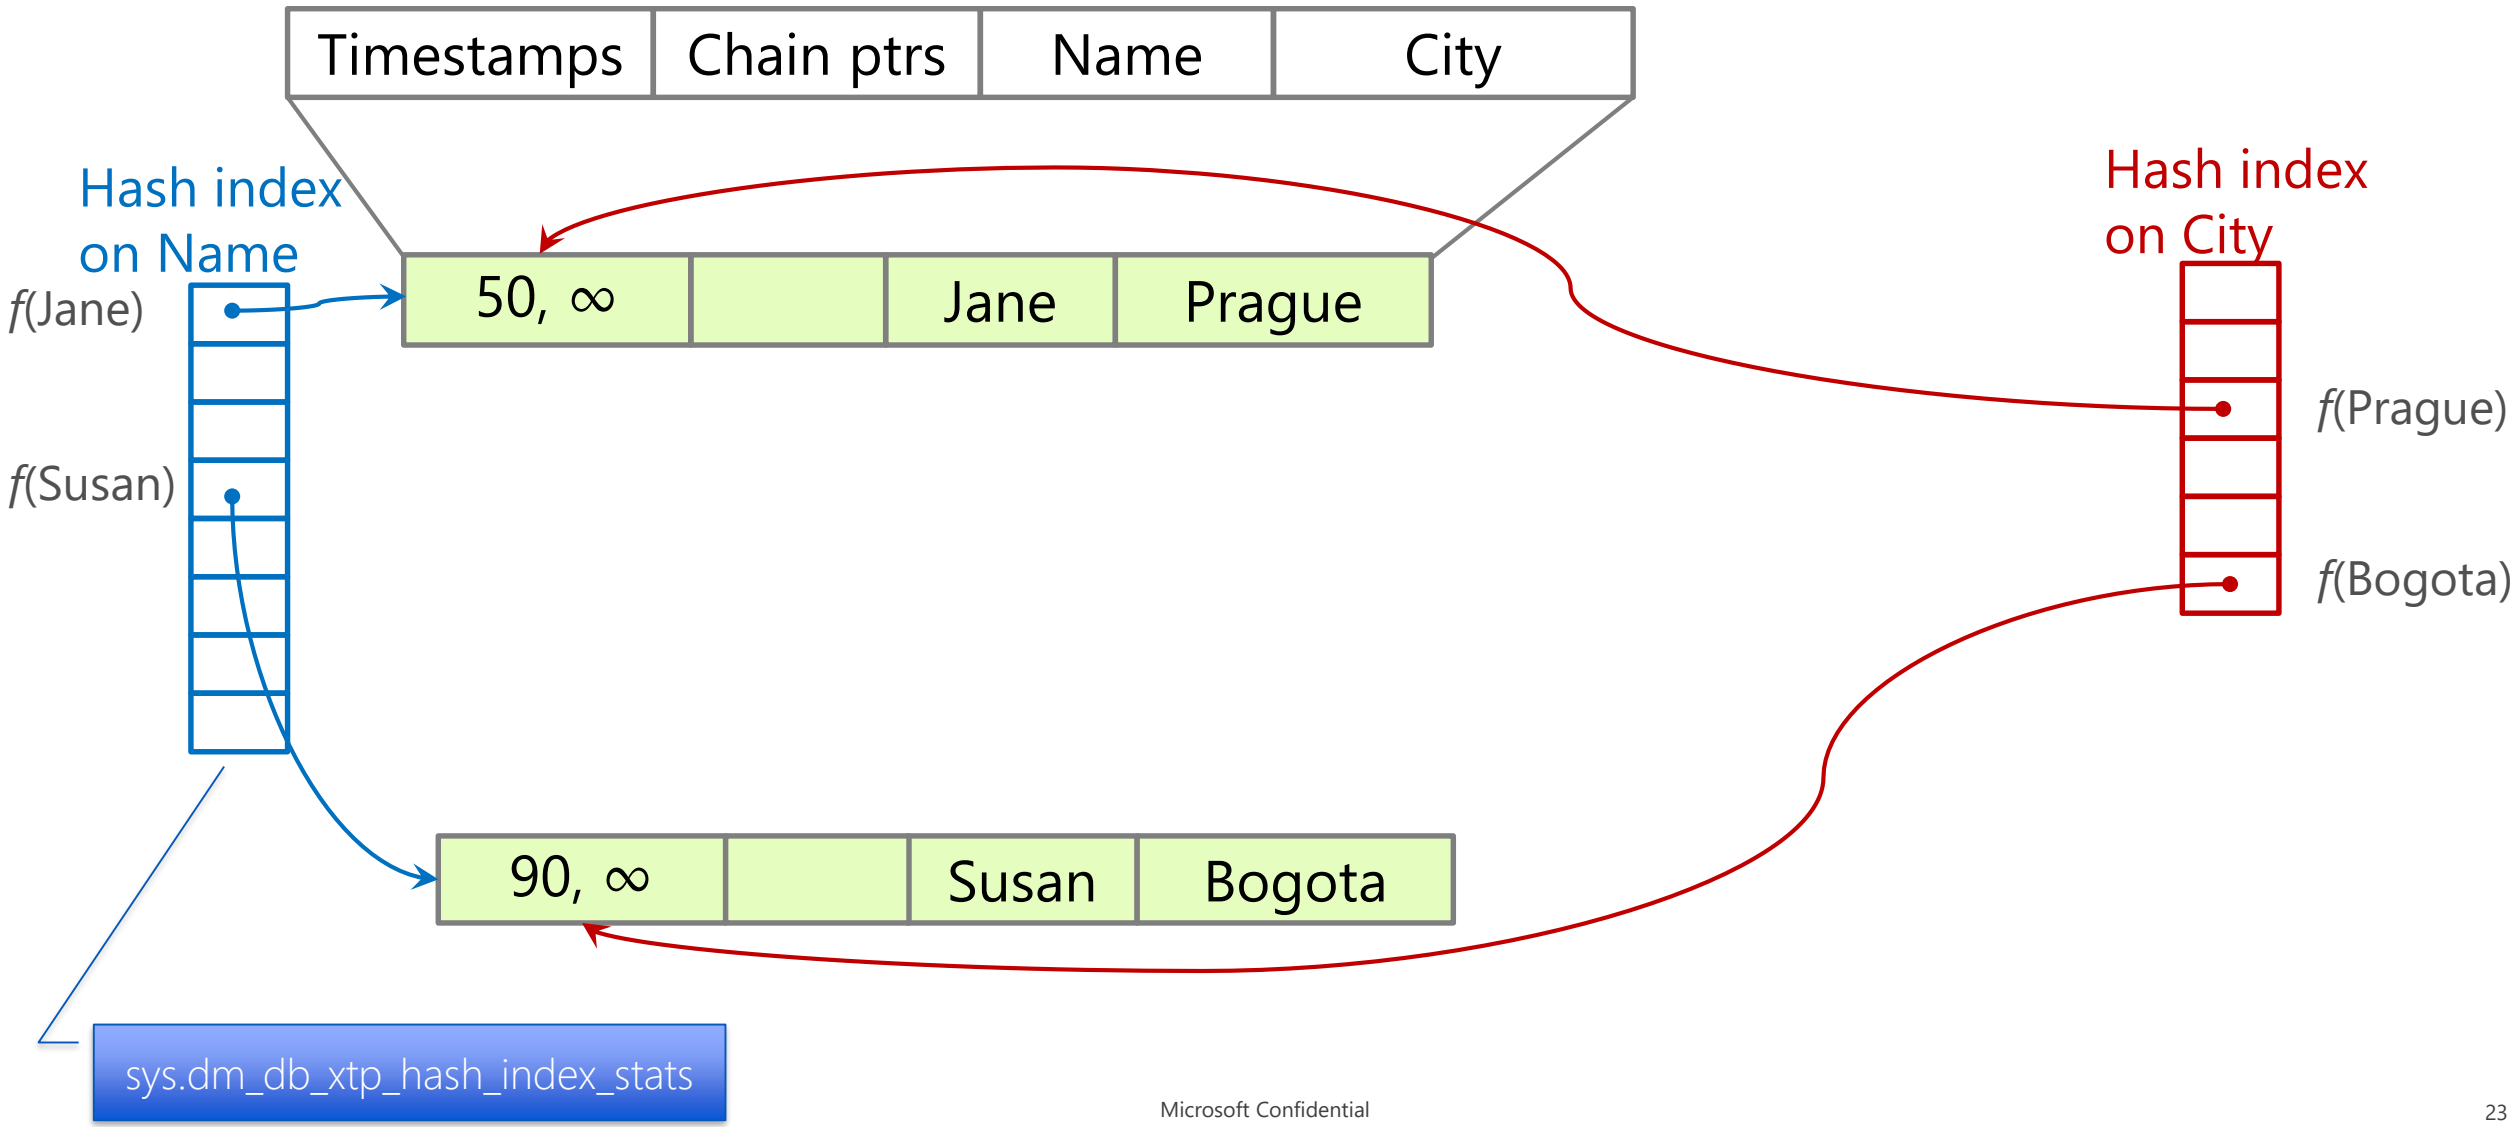

Memory Optimized Table Insert

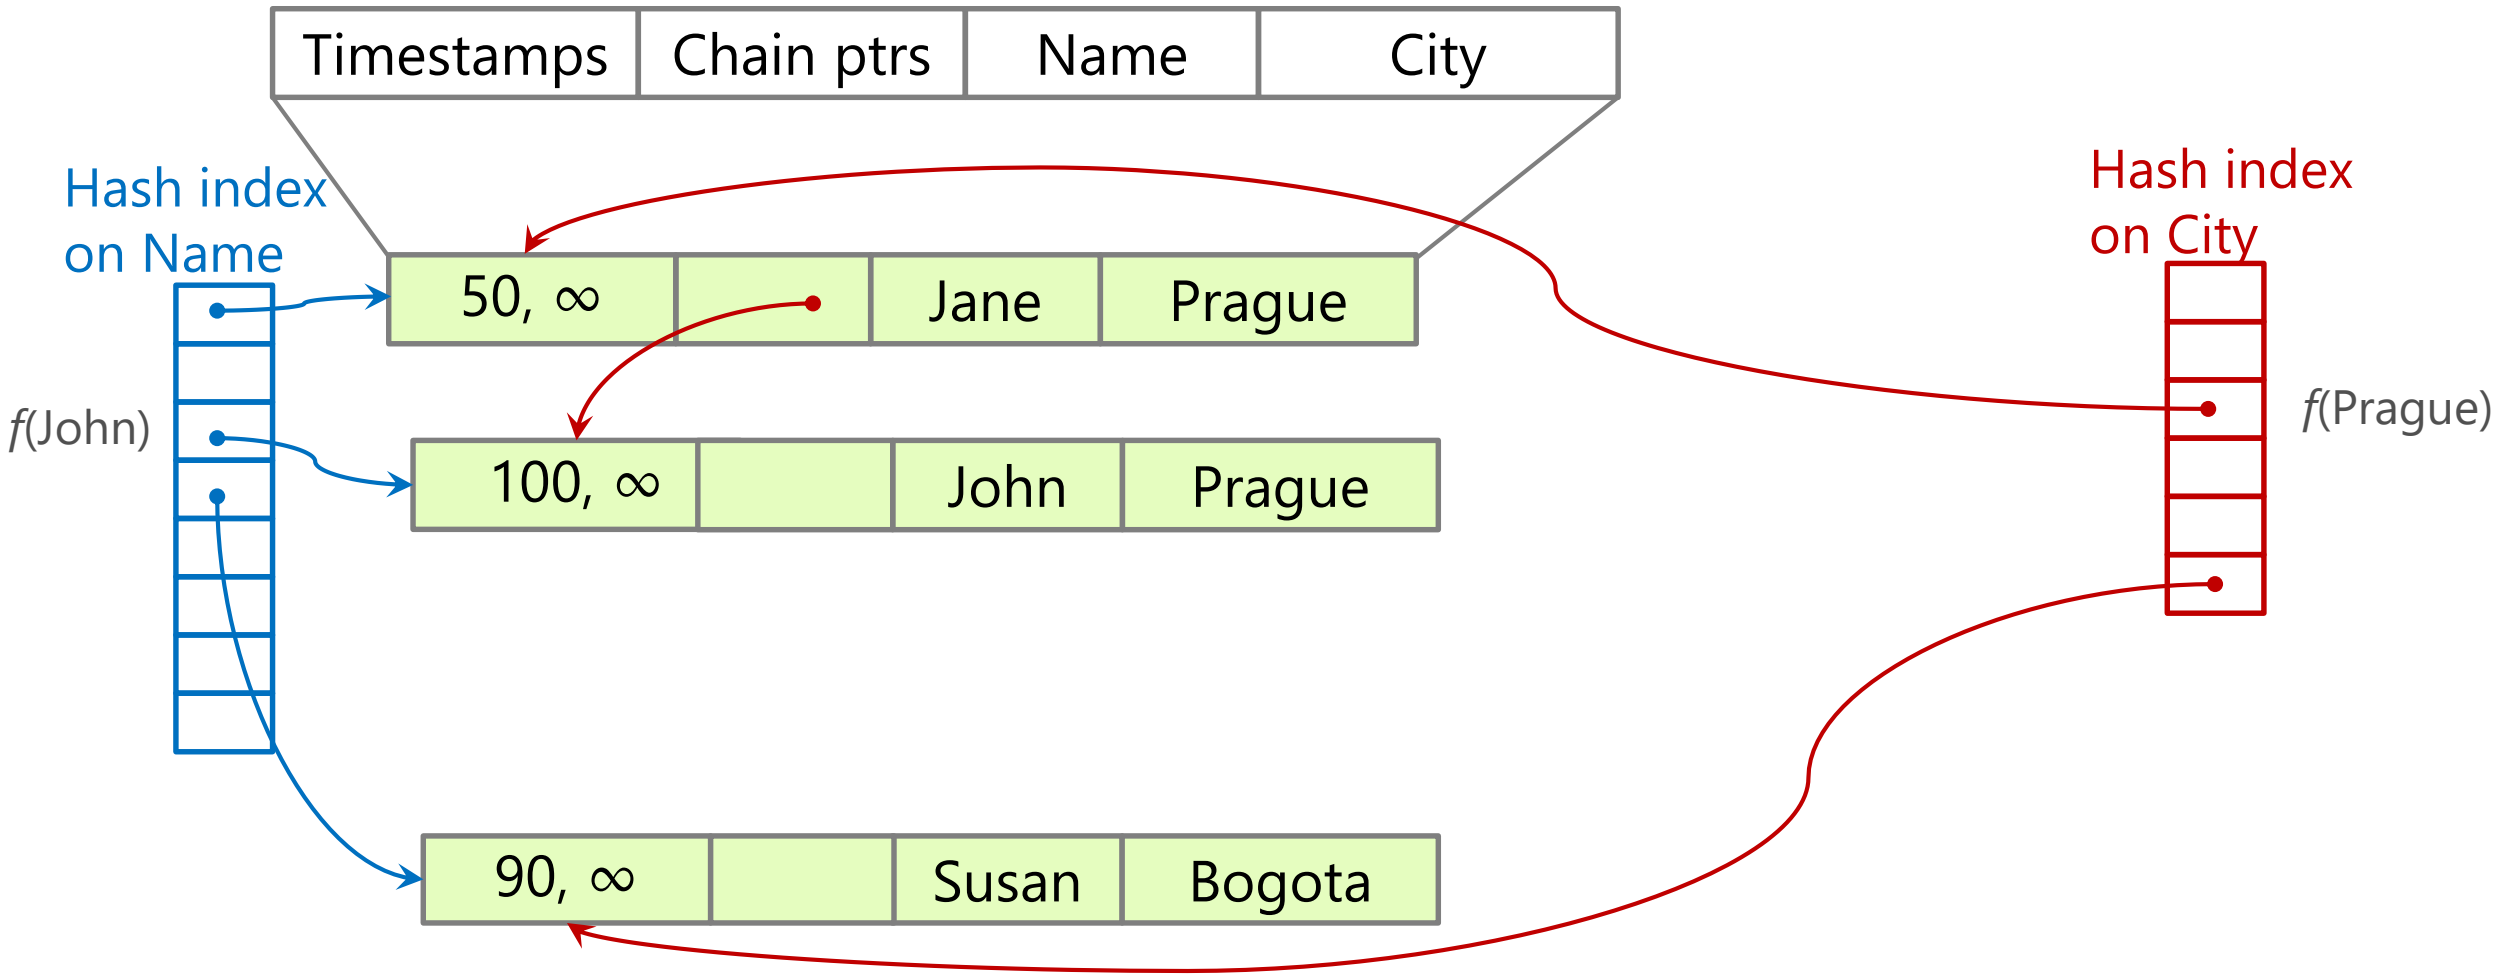

T100: INSERT (John, Prague)

Memory Optimized Table Delete

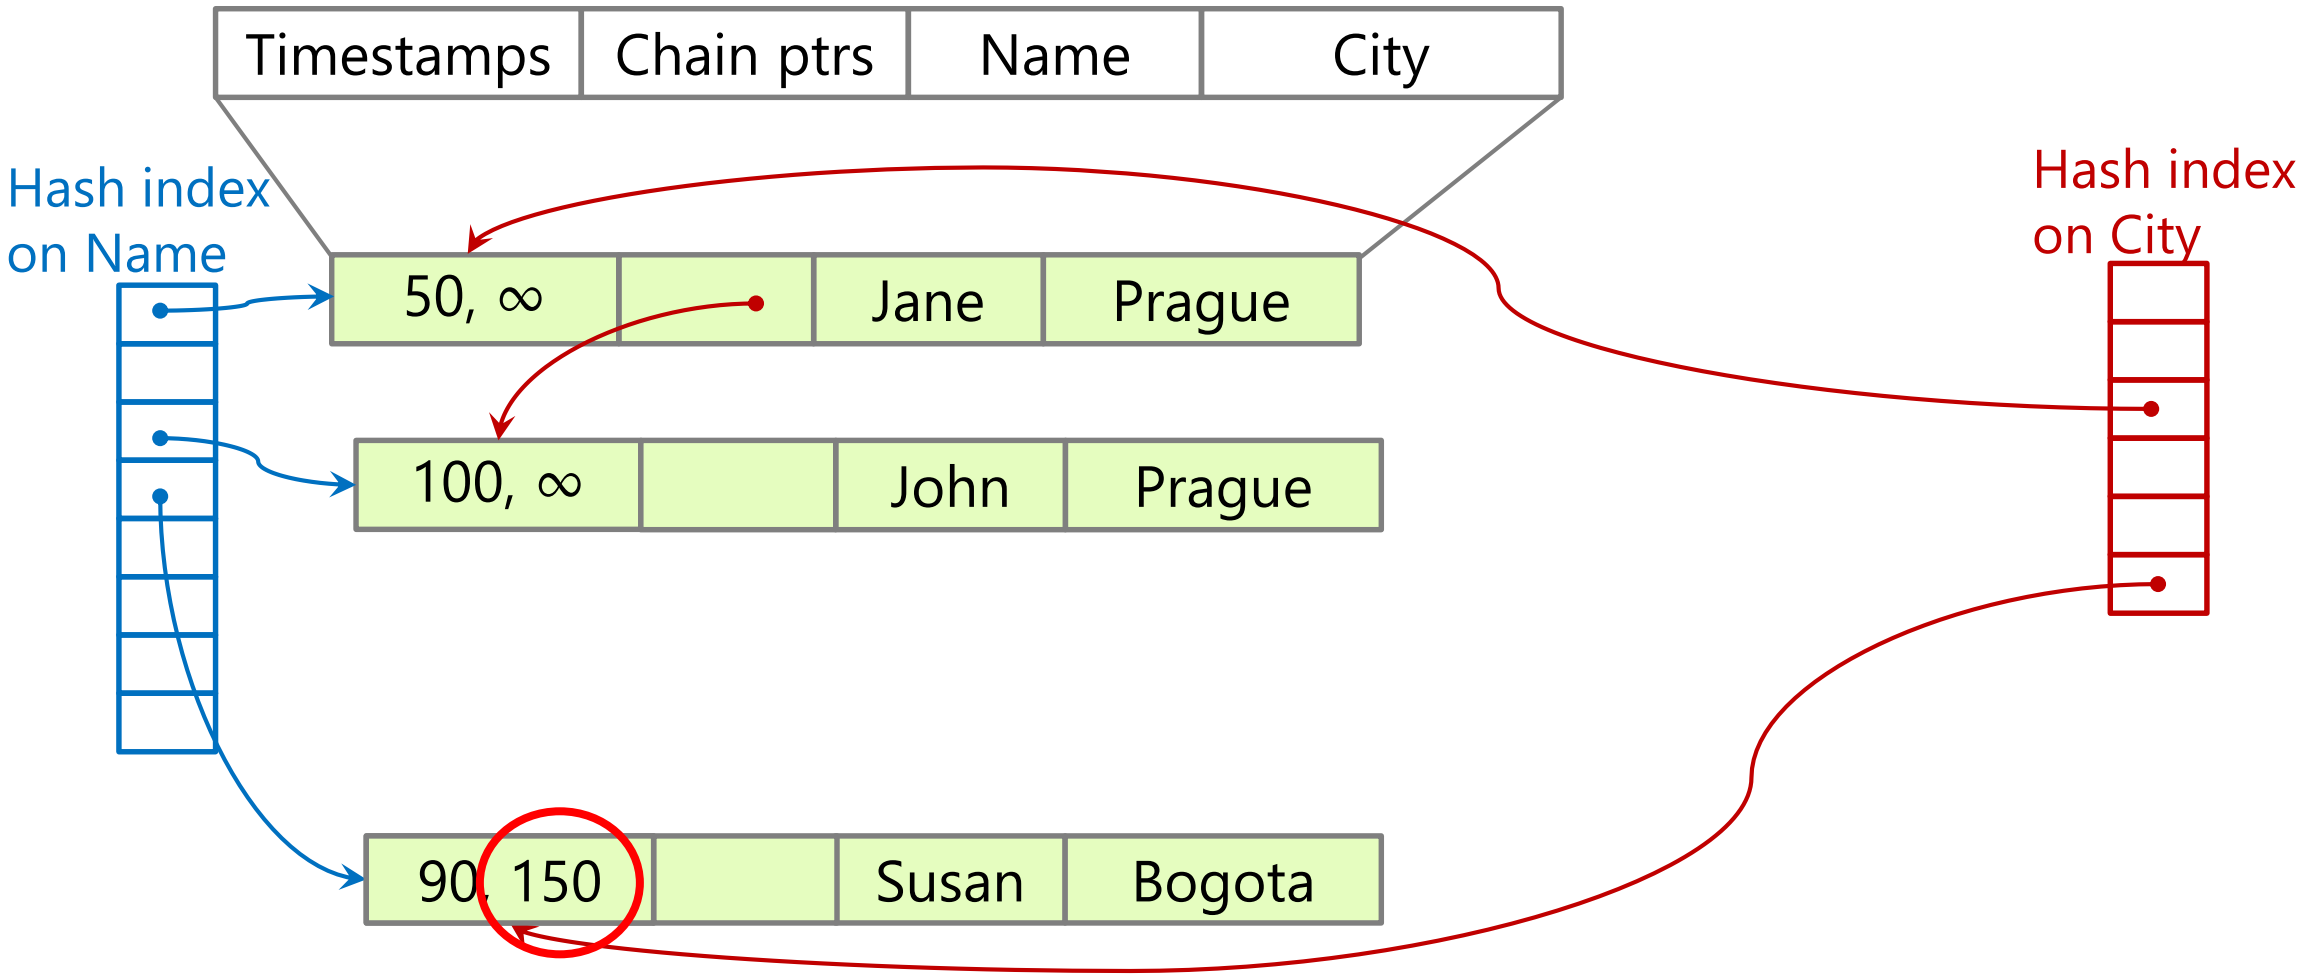

T150: DELETE (Susan, Bogota)

Memory Optimized Table Update

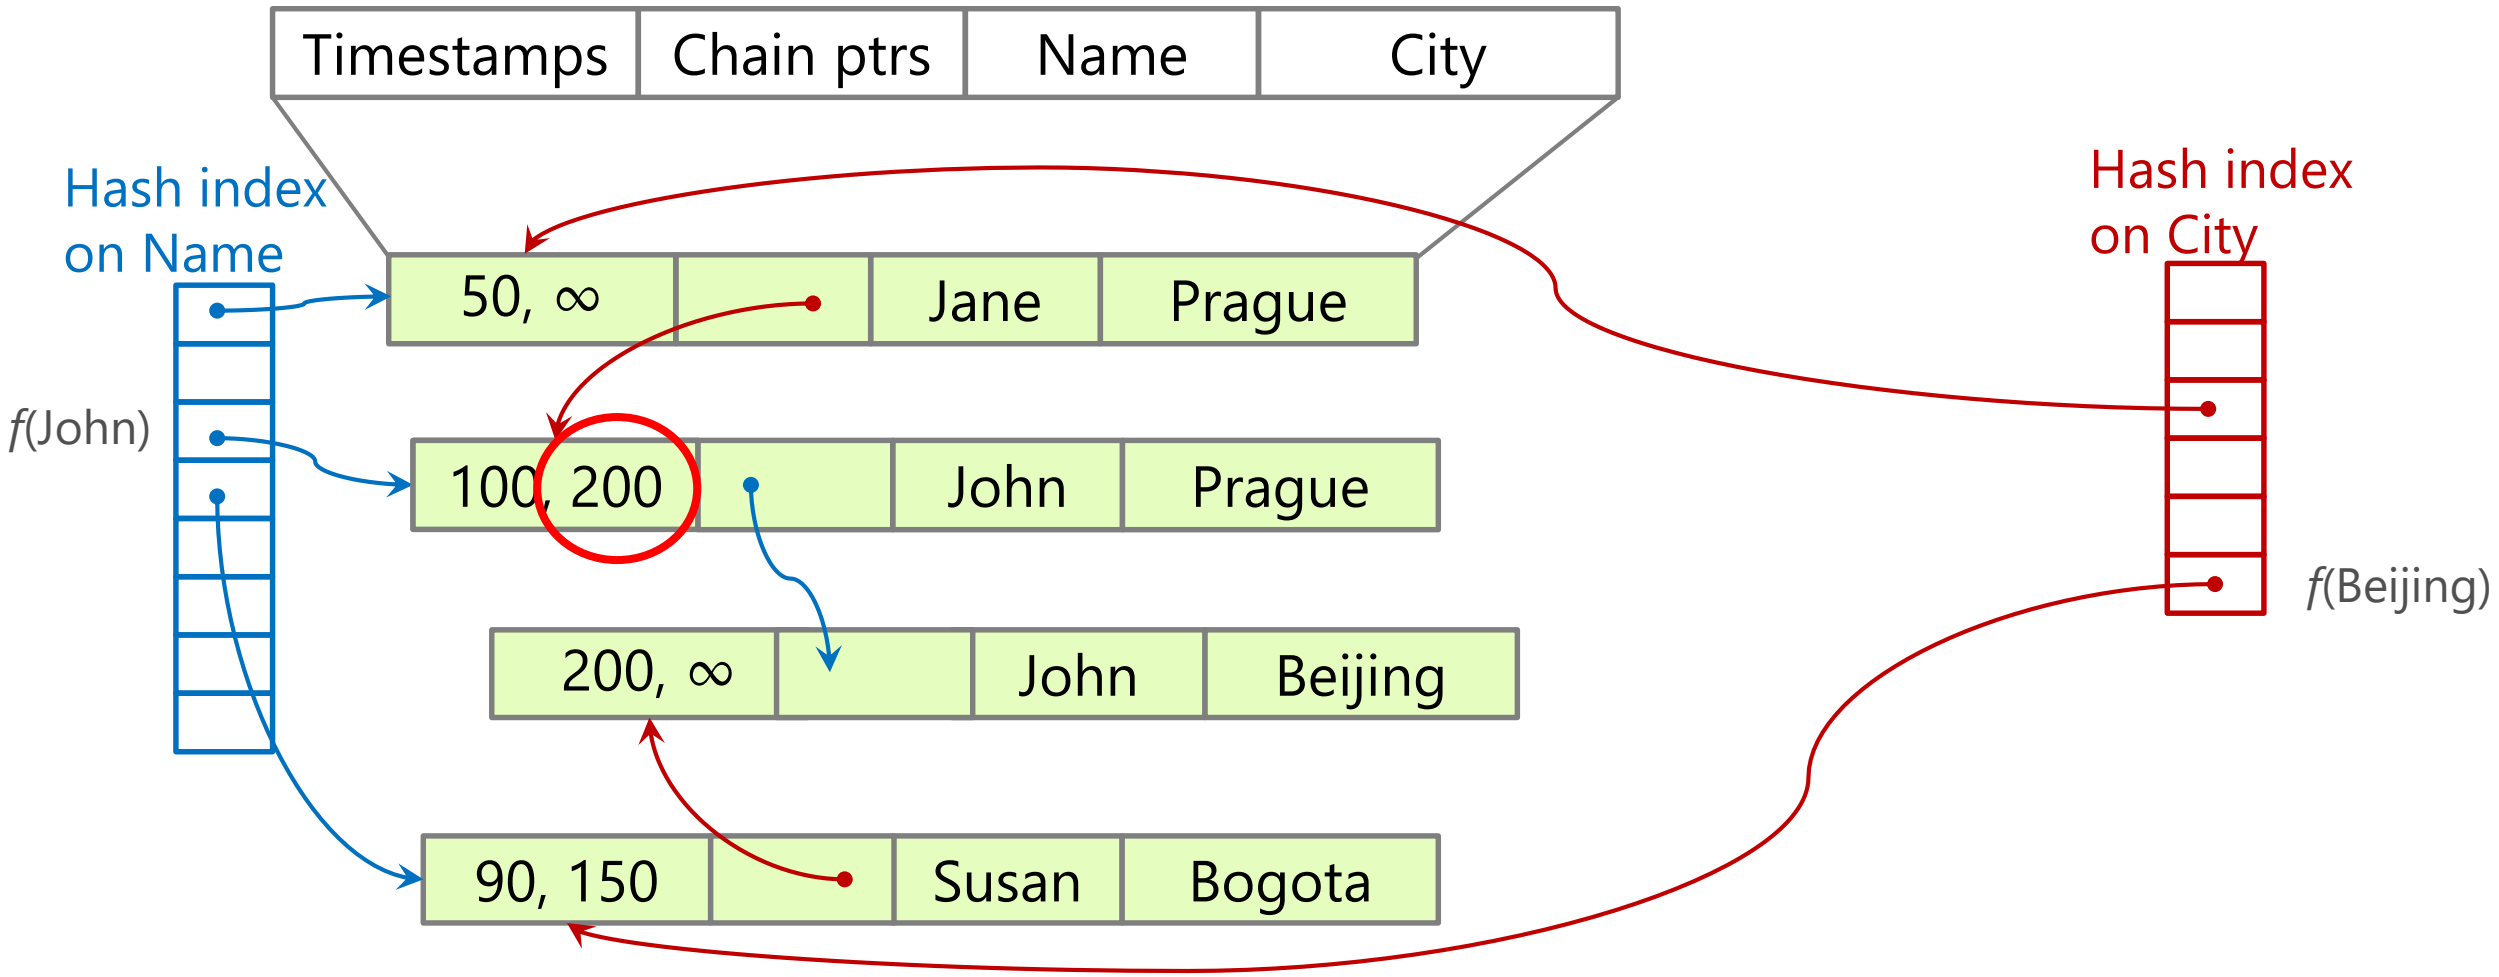

T200: UPDATE (John, Prague) to (John, Beijing)

Garbage Collection

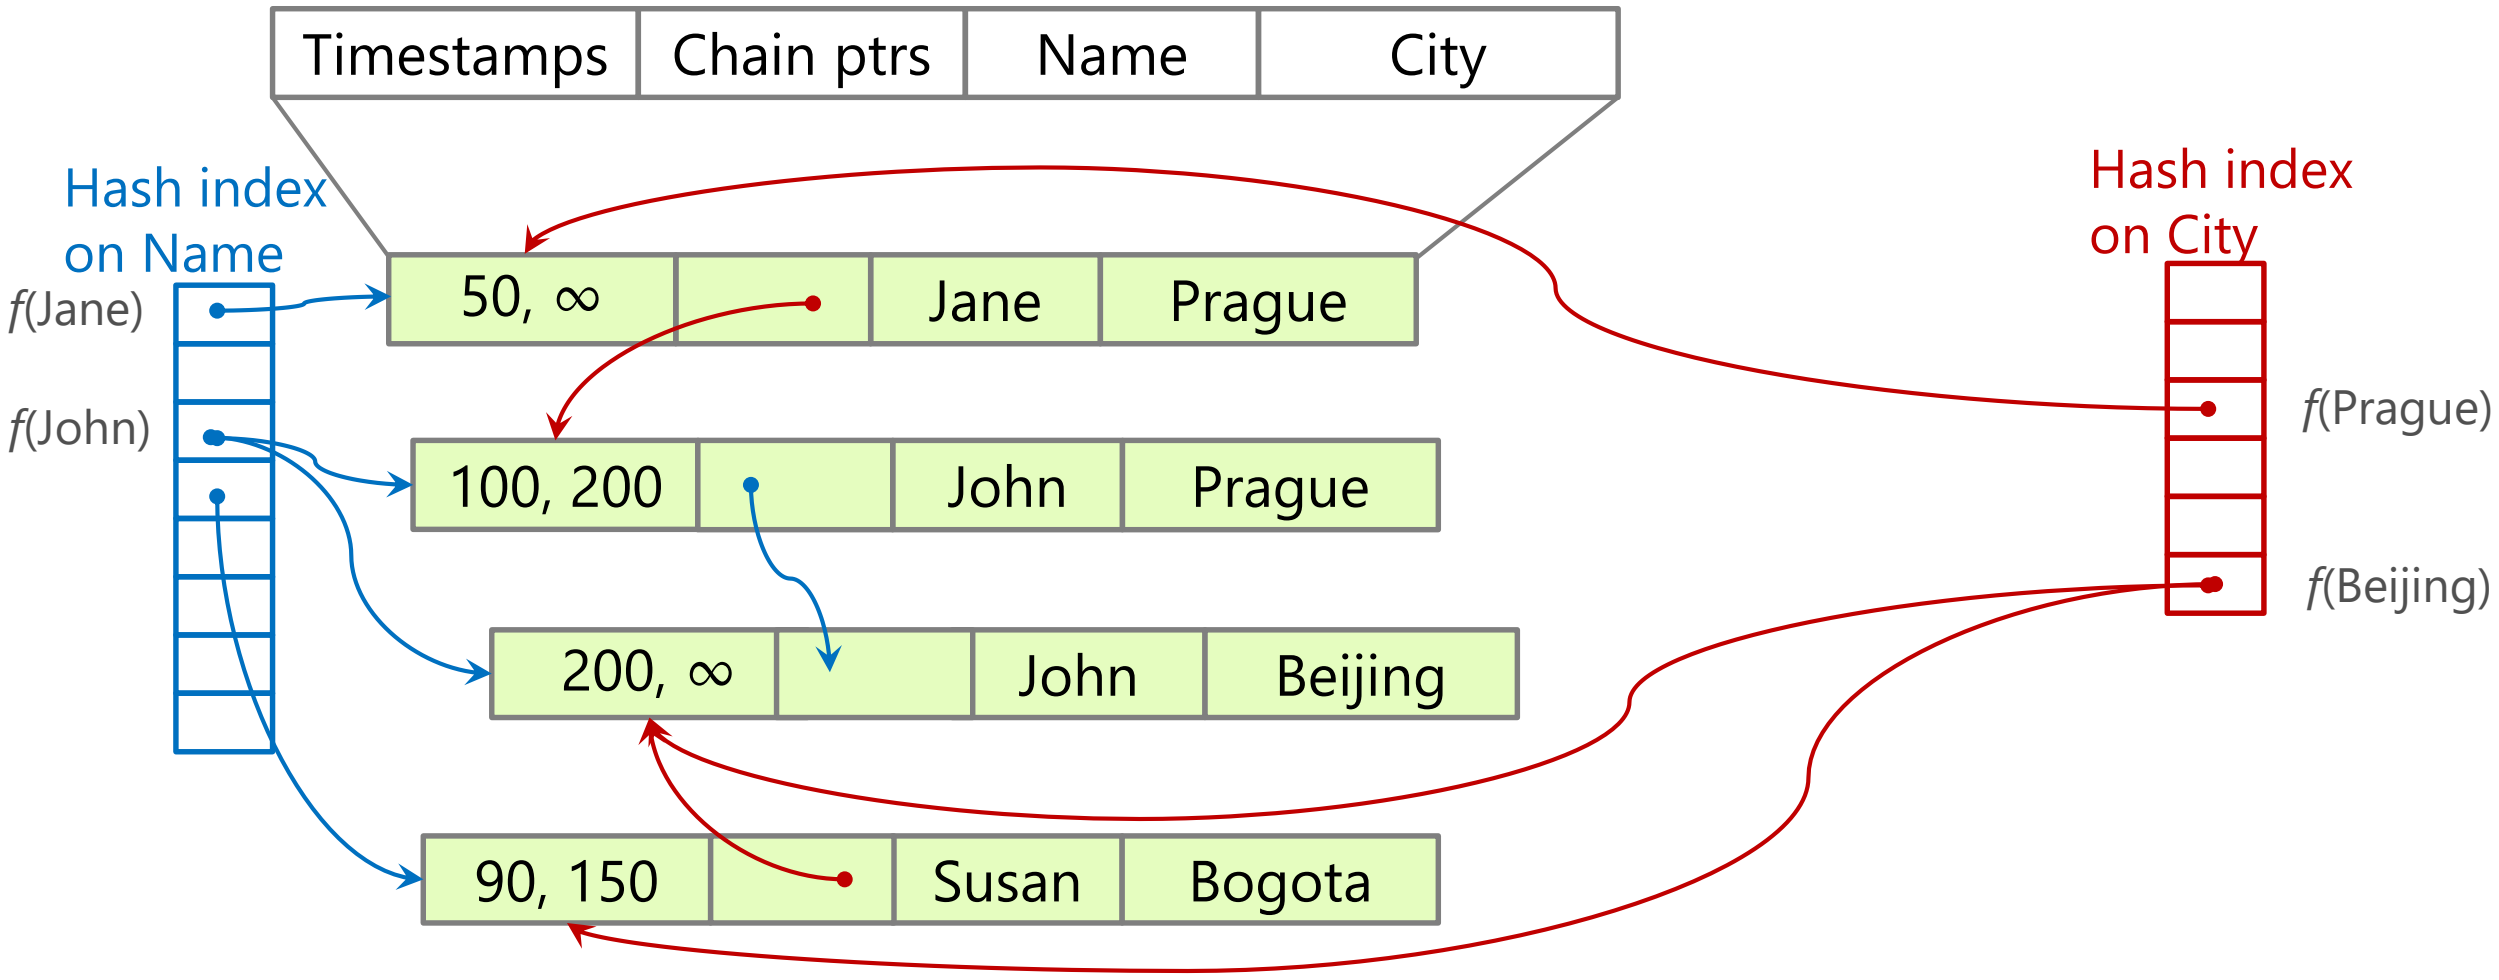

T250: ***Garbage collection***

Non-Clustered (Range) Index

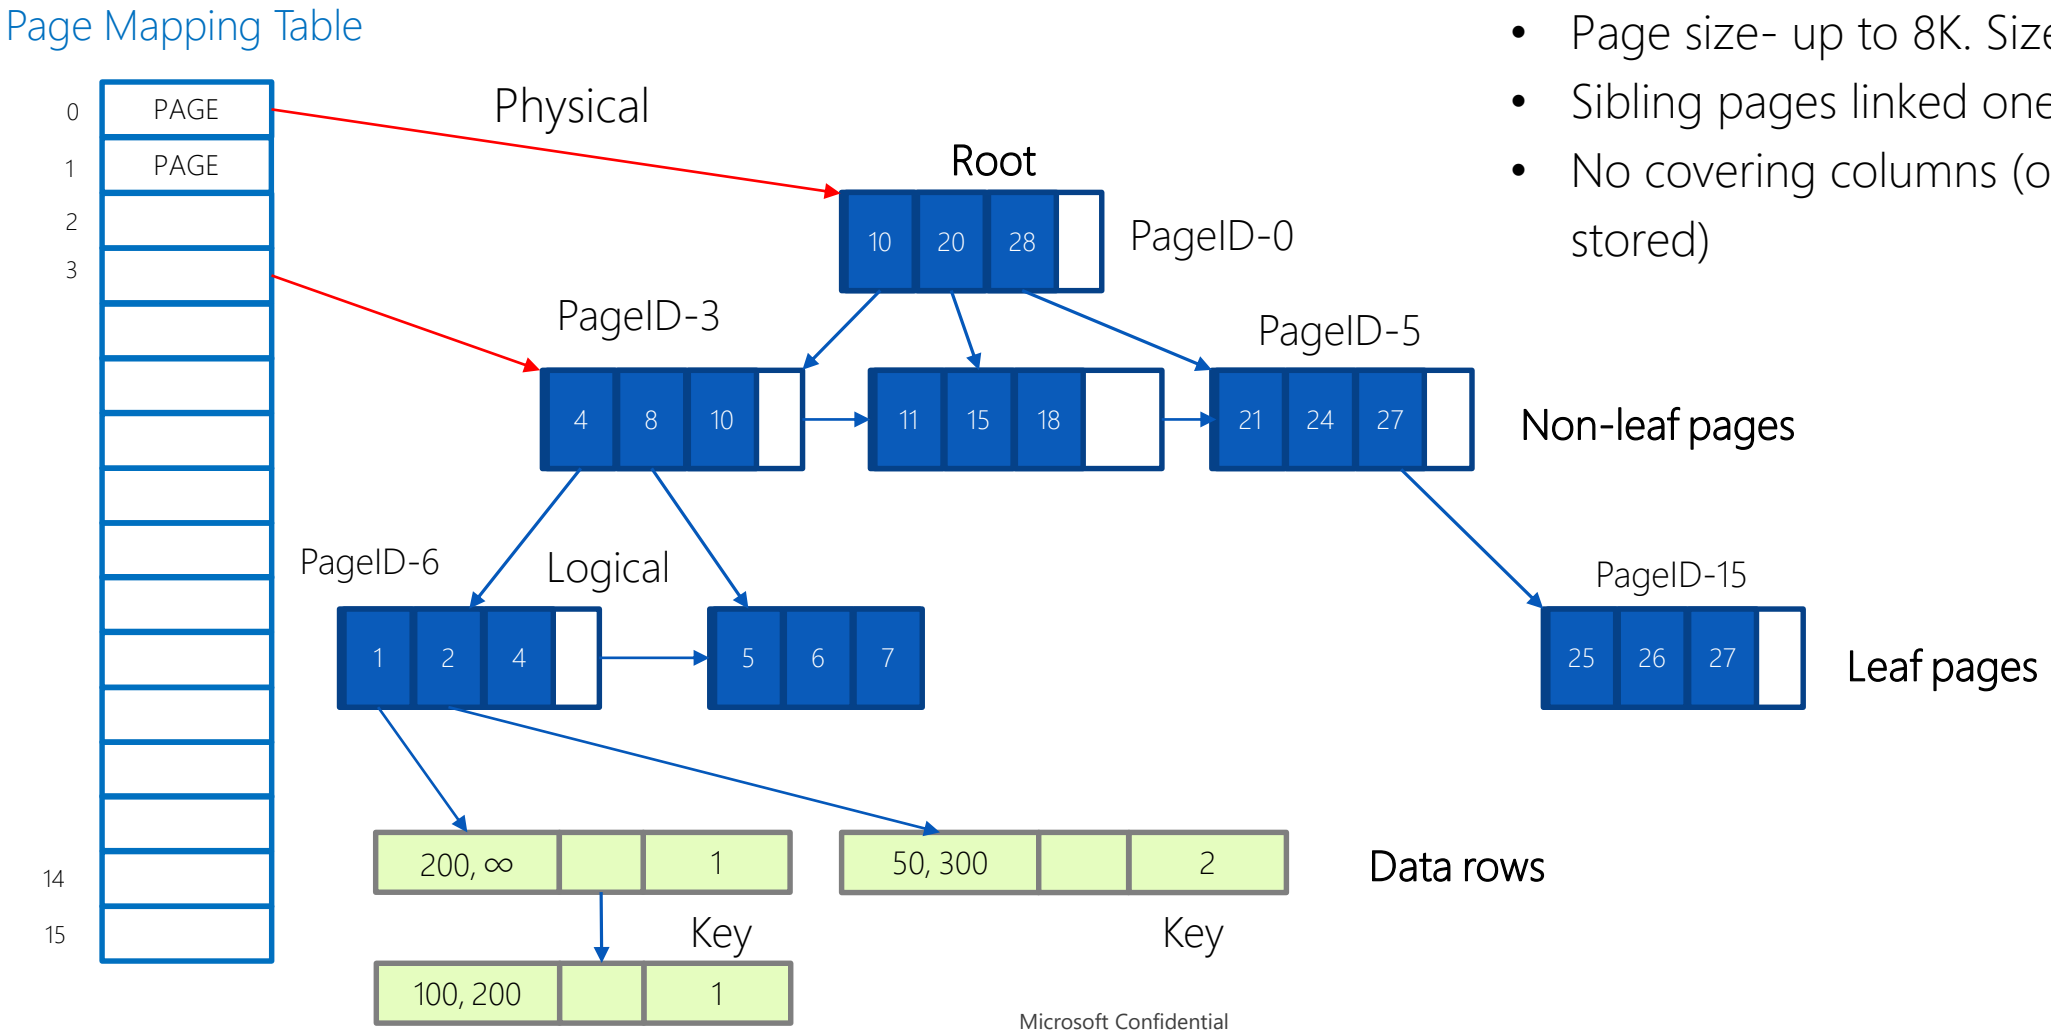

- No latch for page updates
- No in-place updates on index pages
- Page size- up to 8K. Sized to the row
- Sibling pages linked one direction
- No covering columns (only the key is stored)

Range Index Page Format

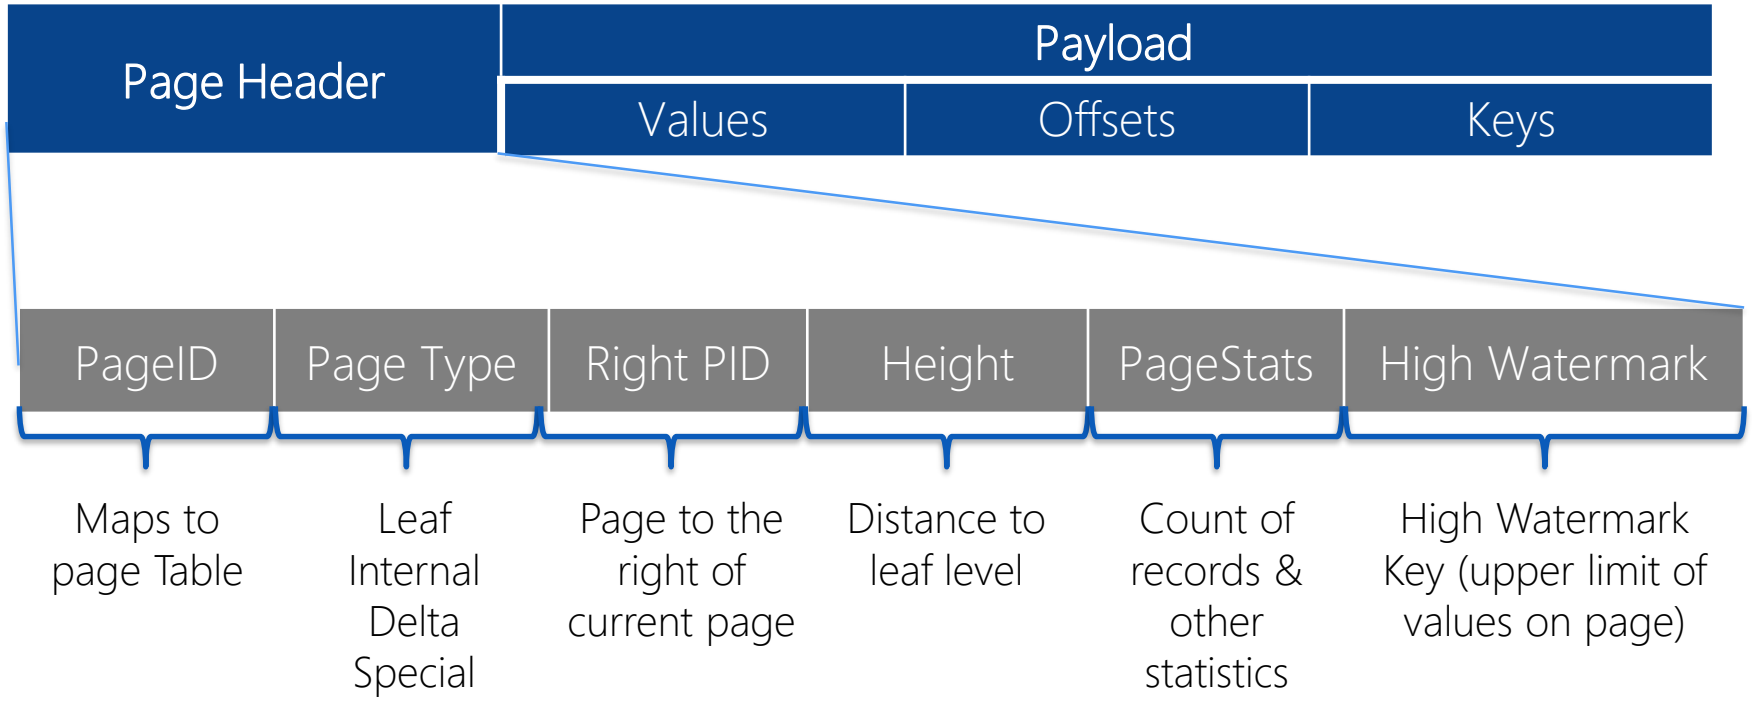

Non Clustered Index – Delta Consolidation

- Changes done through Delta records
- Leaf index page consolidated after 16 delta changes.

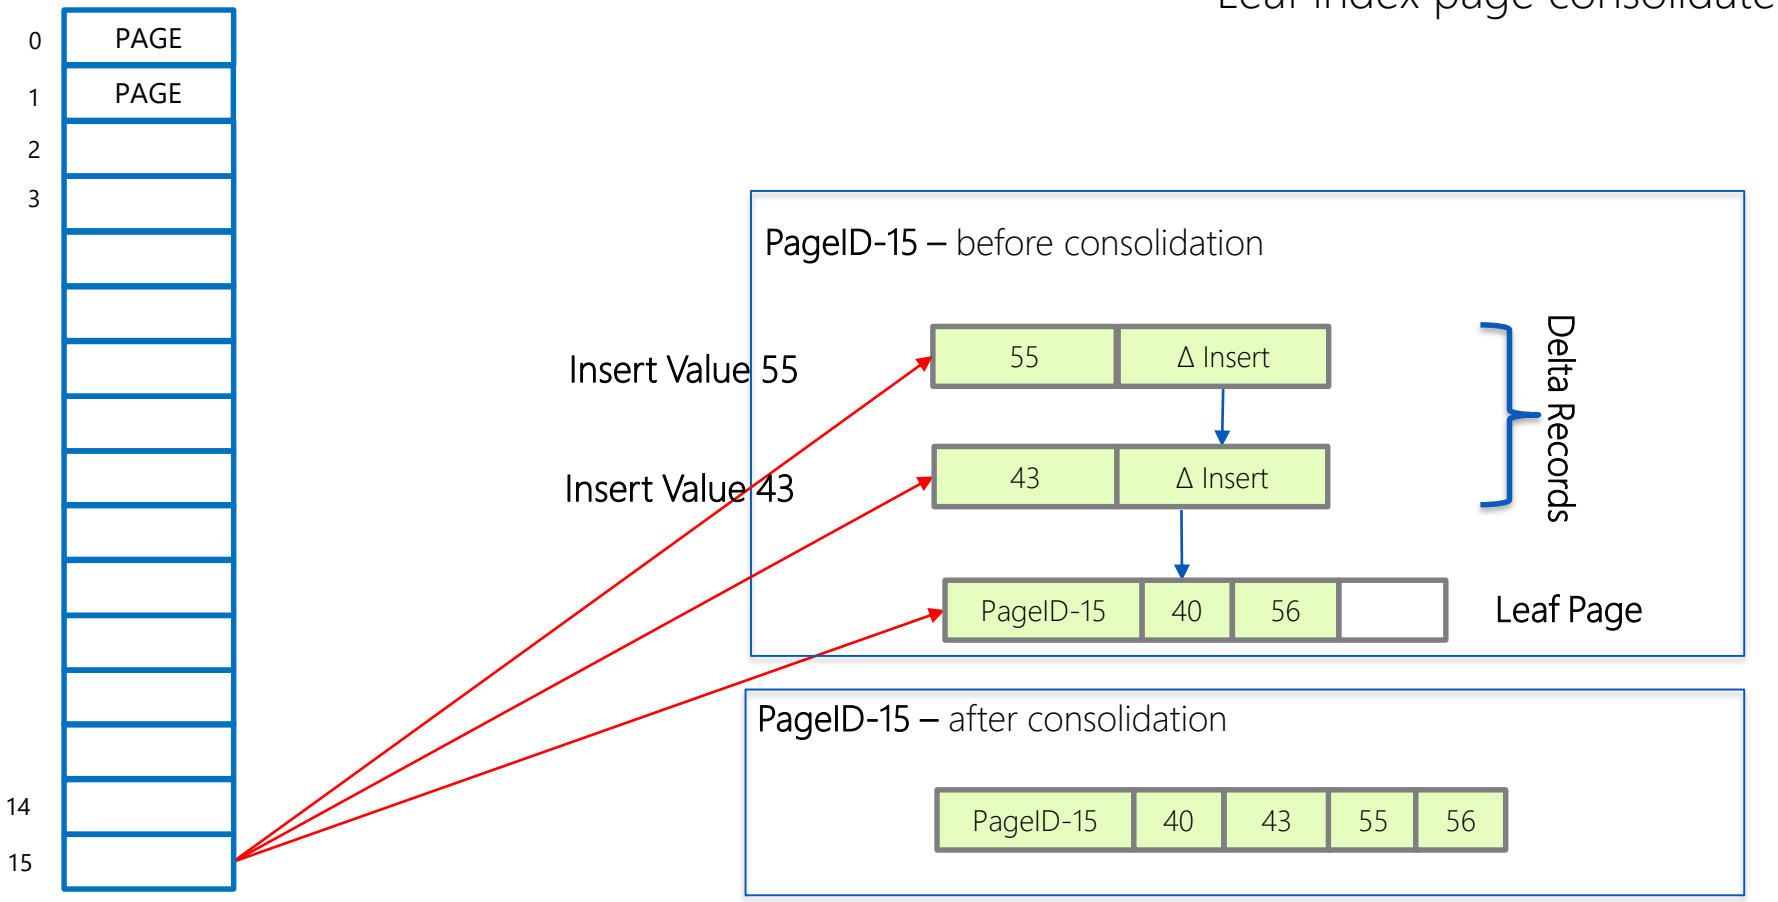

Non Clustered Index – Split Operation

Split operation requires two modifications:

- Splitting the page into 2 pages
- Updating the parent to point to the new child page

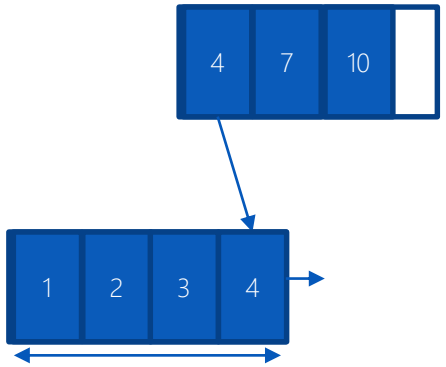

8k (leaf page full)

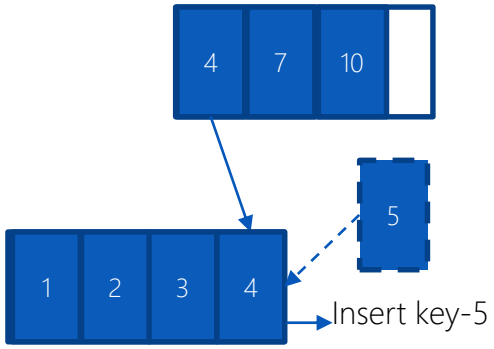

New insert forces split

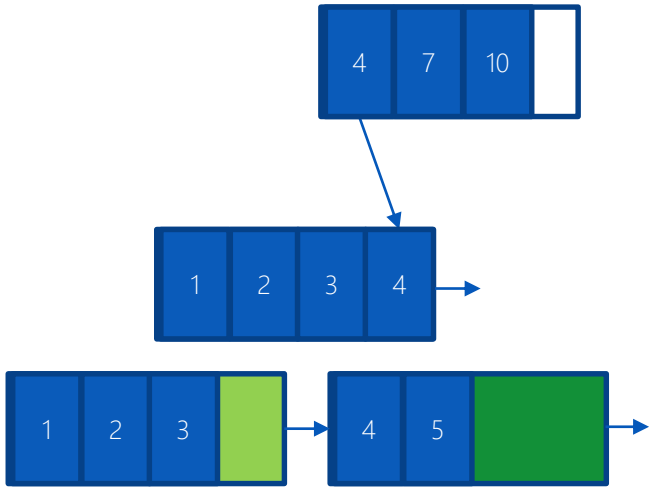

Rows split into two new pages

New internal Index page created

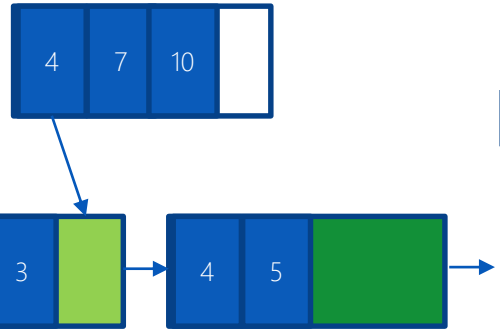

Old leaf page is removed

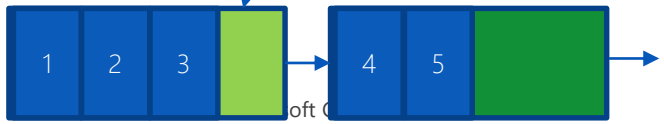

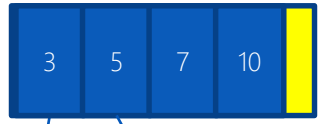

Old internal Index page is removed

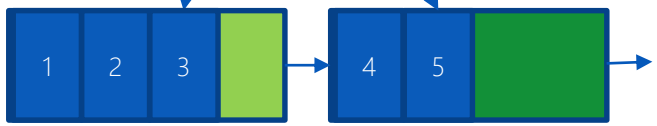

Range Index – Merge Operation

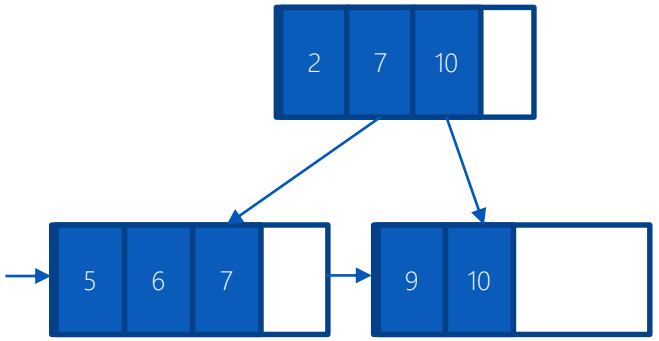

Leaf index pages

Operation - Delete row 10

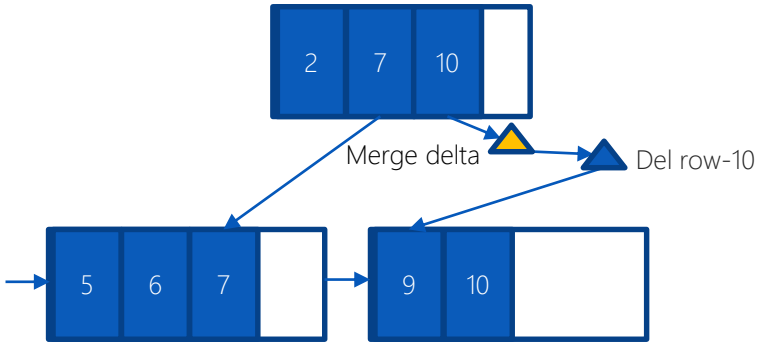

Leaf index pages

Row 10 marked deleted

STEP – 1

(Insert delta page for merge and for row-10)

New internal Index page created

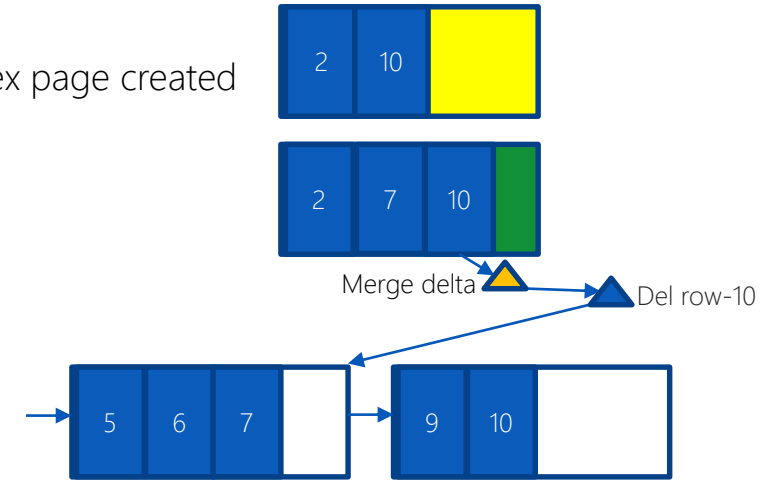

Leaf index pages

Row 10 marked deleted

STEP – 2

(Remove the parent row with Key-7 (requires new index page) and point key-10 to left leaf page)

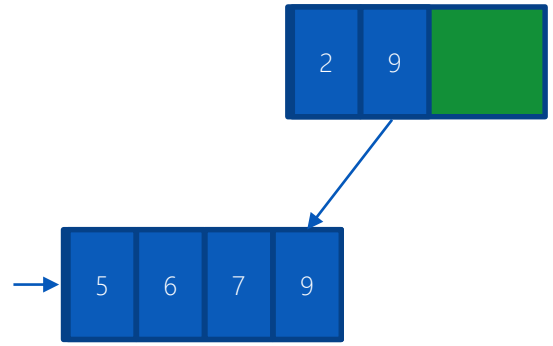

Leaf index pages merged into a new leaf page

STEP – 3

(Remove the delta pages and move key-9 to the left page)

Index Choice Considerations

Operation	Hash index	Nonclustered index (Range)	Disk-based index
Index Scan	Yes	Yes	Yes
Index seek on equality predicate(s) (=)	Yes (Full key required)	Yes (*)	Yes
Index seek on inequality (>, <, <=,'>=)	No (results in an index scan)	Yes (*)	Yes
Sort-order matching the index definition	No	Yes	Yes
Sort-order matching the reverse of the index definition	No	No	Yes

Yes → index can adequately service the request
No → index cannot be used successfully to satisfy the request.
(*) → For a non-clustered memory-optimized index, the full key is not required to perform an index seek

Durability

Memory optimized tables can be durable or non-durable

- Non-durable tables are ideal for transient data
- Default is durable

```
CREATE TABLE [dbo].[SalesOrder_MemOpt] (
[Order_ID] INT NOT NULL,
[Order_Date] DATETIME NOT NULL,
[Amount] FLOAT NOT NULL,
CONSTRAINT PK_SalesOrderID PRIMARY KEY NONCLUSTERED HASH (Order_ID)
)
WITH (MEMORY_OPTIMIZED = ON, DURABILITY = SCHEMA_AND_DATA);
```

Single memory optimized filegroup

- Sequential IO pattern (no random IO)
- Filegroup can have multiple containers to aid in parallel recovery
- Recovery depends on IO speed – recommend SSD or Fast SAS drives

Storage

Filestream container is the underlying storage mechanism

- Checksums and single-bit correcting ECC on files

Data files

- ~128MB in size, write 256KB chunks at a time
- Stores only the inserted rows (i.e. table content)
- Chronologically organized streams of row versions

Delta files

- File size is not constant, write 4KB chunks at a time (16MB if Server has > 16GB RAM)
- Stores IDs of deleted rows

Data and Delta Files

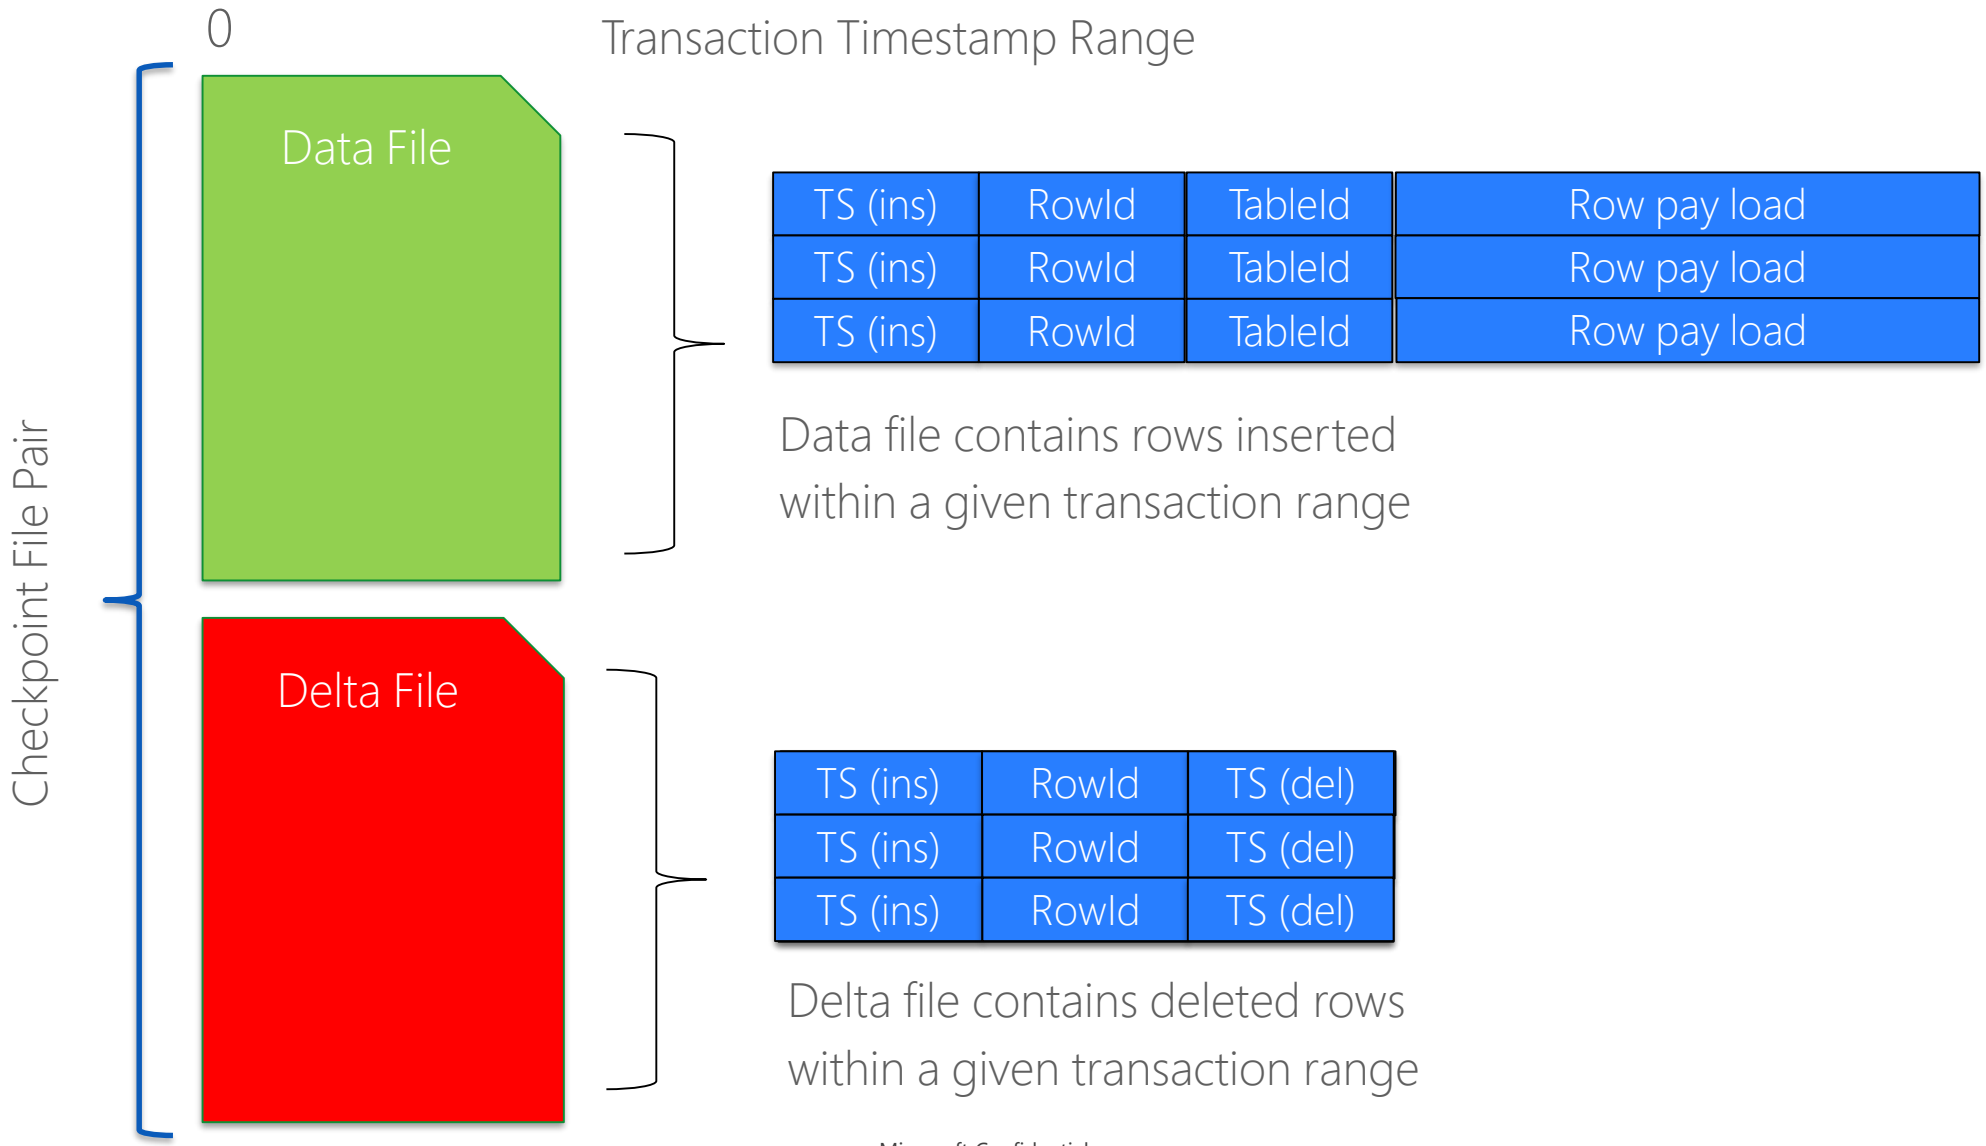

Checkpoint Process

Not tied to a "recovery interval"

"**Offline Checkpoint**" log records and flushes data to the data and delta files in the filestream container

Invoked in multiple ways:

- Manual Checkpoint – Explicit checkpoint command issued
- Size based – Log grown by 512MB
- Time Based – if the time since last checkpoint issued exceeds threshold

SE_MANAGE_VOLUME recommended for performance reasons (instant file initialization)

Troubleshooting: sys.dm_db_xtp_checkpoint_stats

Populating Data and Delta Files

SQL Transaction log (LogPool)

- Data file has pre-allocated size (128 MB)
- Engine switches to new data file when the current file is full
- Transaction does not span data files
- Once a data file is closed, it becomes read-only
- Row deletes are tracked in delta file
- Files are append only

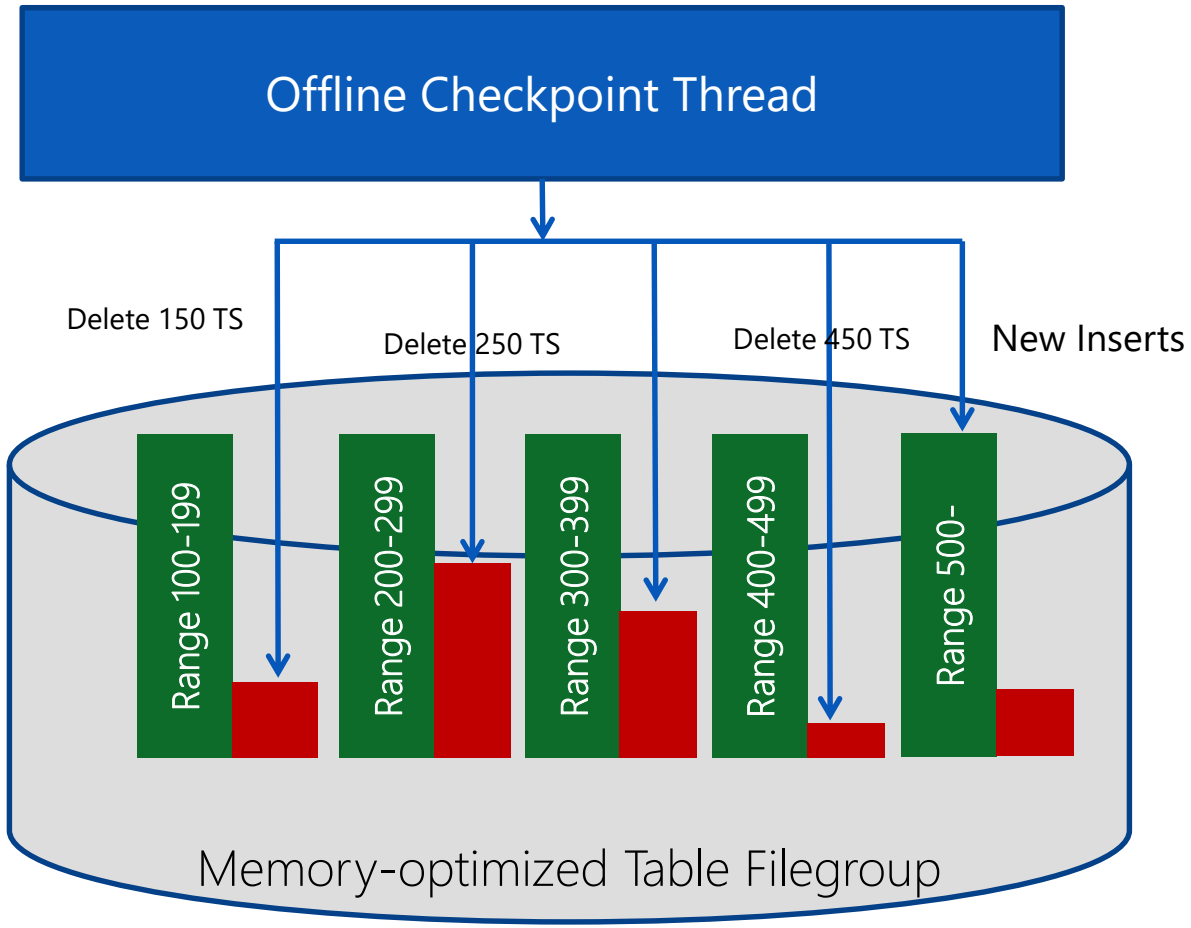

Data file with rows generated in timestamp range

IDs of Deleted Rows (height indicates % deleted)

Merge Operation

What is a Merge Operation?

- Merges two or more adjacent data/delta files pairs into 1 pair

Need for Merge

- Deleting rows causes data files to have stale rows
- DMV: *sys.dm_xtp_checkpoint_files* can be used to find inserted/deleted rows and free space

Benefits of Merge

- Reduces storage (i.e. fewer data/delta files) required to store active data rows
- Improves the recovery time as there will be fewer files to load

When is Merge done?

- 2 or more consecutive CFPs can be consolidated, after accounting for deleted rows, such that the resultant rows can fit into 1 CFP of ideal size
- single CFP can be self-merged the data file exceeds 256 MB and over half of the rows are deleted
- Manual Merge invoked calling ***sys.sp_xtp_merge_checkpoint_files***

Merge

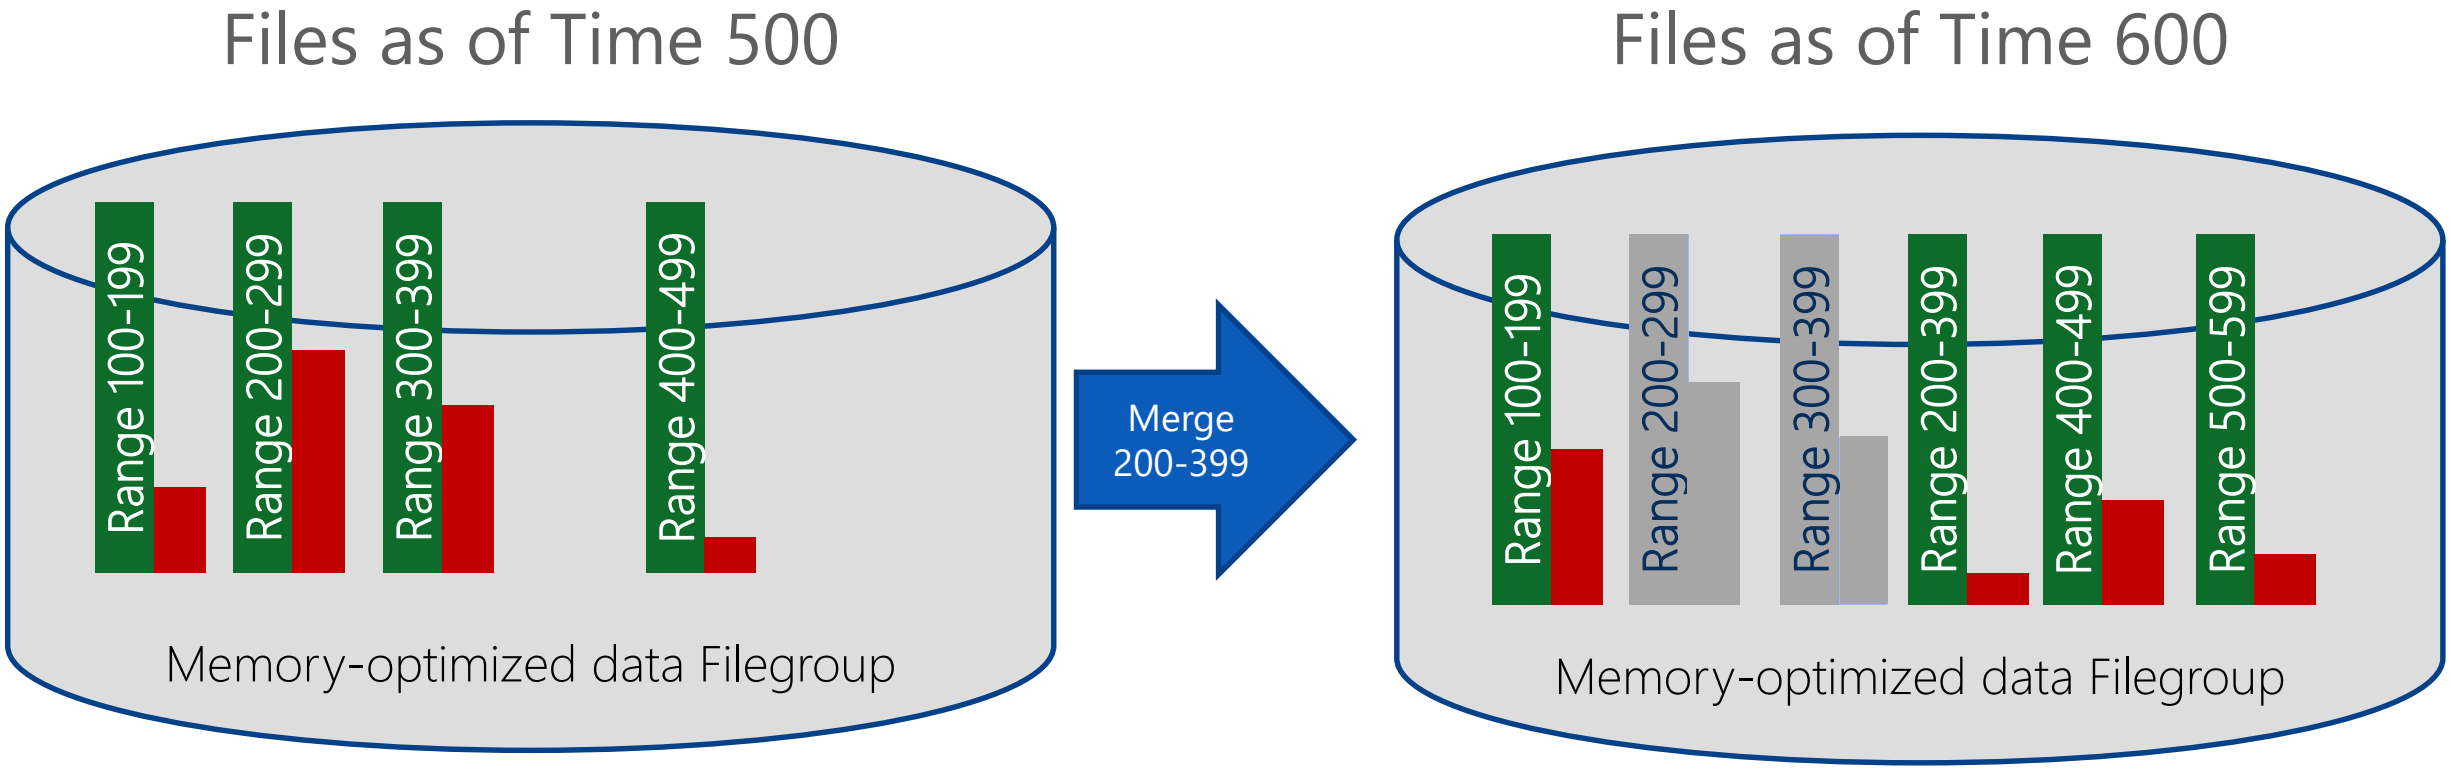
Files as of Time 500
Range 100-199
Range 200-299
Range 300-399
Range 400-499
Memory-optimized data Filegroup
Merge
200-399
Files as of Time 600
Range 100-199
Range 200-299
Range 300-399
Range 200-399
Range 400-499
Range 500-599
Memory-optimized data Filegroup

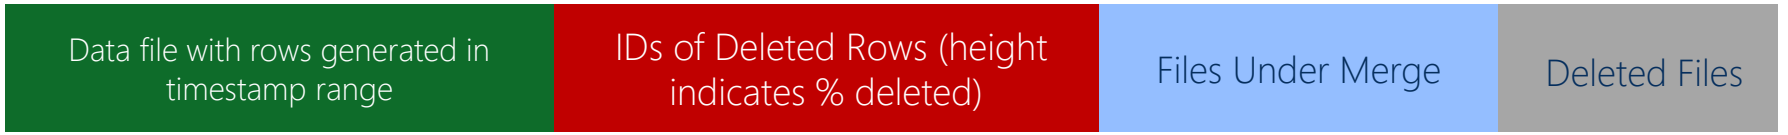
Data file with rows generated in timestamp range
IDs of Deleted Rows (height indicates % deleted)
Files Under Merge
Deleted Files

Checkpoint Files States

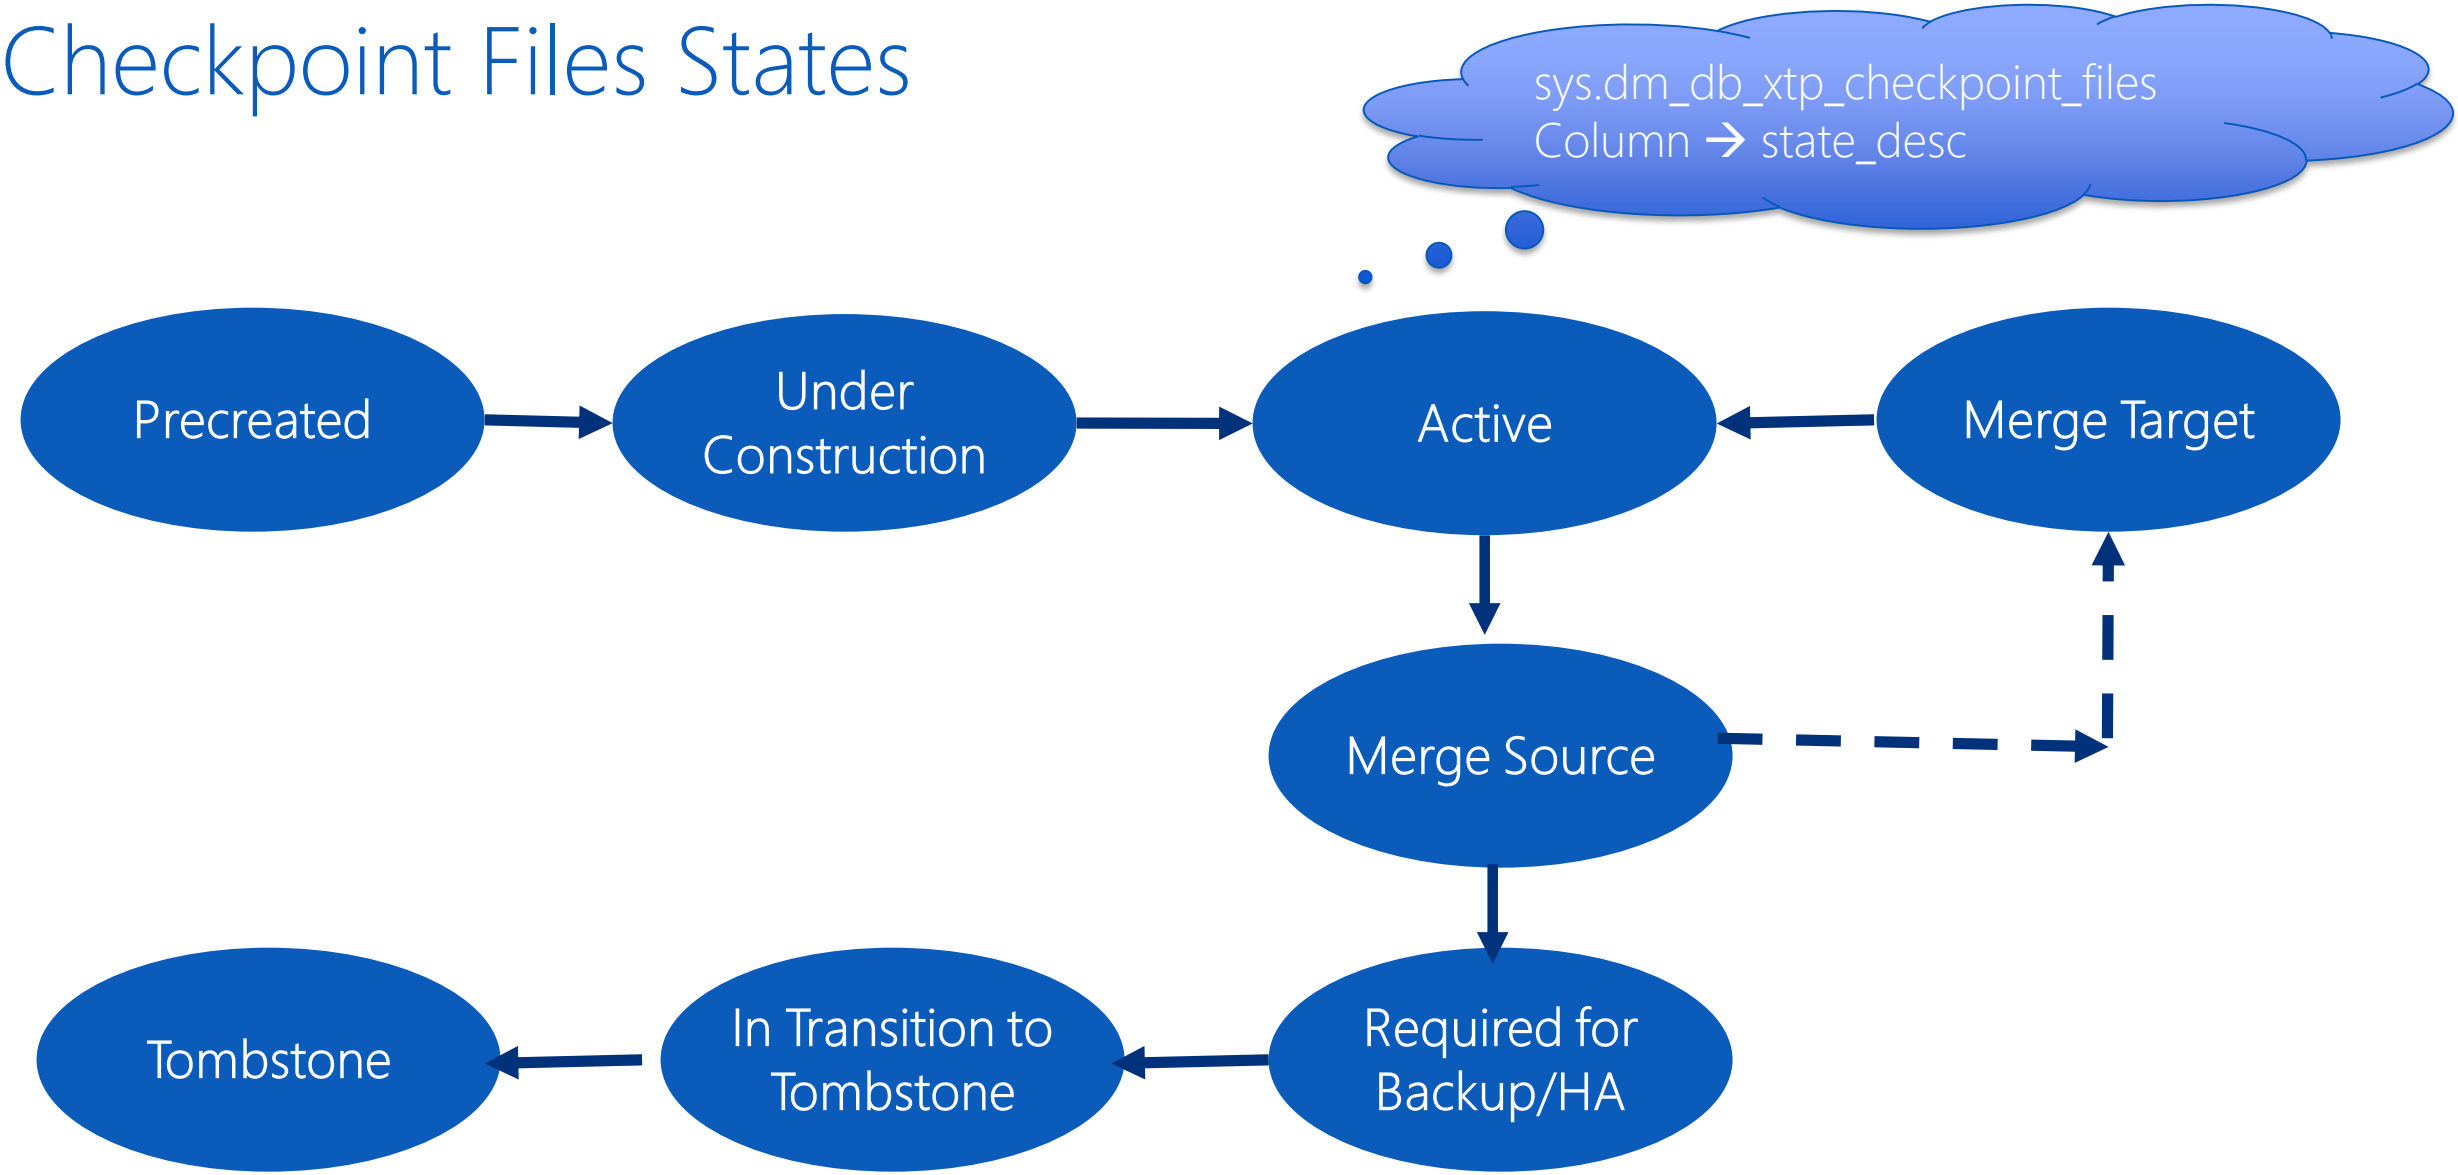

Logging

Uses SQL transaction log to store content

- Each In-Memory OLTP log record contains a log record header followed by memory optimized-specific log content

All logging for memory-optimized tables is logical

- No log records for physical structure modifications
- No index-specific / index-maintenance log records
- Log records are written only on a commit, no UNDO information is logged

Latency

- Latency is absolutely critical, recommend SSDs

Recovery Models

- All three recovery models are supported

Delayed Durability

Transaction commits logged asynchronously

Set at Database Level or Atomic Block

- **Disabled** – Normal behavior durability guaranteed.
- **Allowed** – Allowed at the DB Level, Transaction has to specify durability options, default is a durable transaction.

  COMMIT TRAN……. WITH (DELAYED_DURABILITY=ON)

- **FORCED** – Changes default durability for the DB to "delayed". Can be useful for applications bottlenecked on Log IO, that can tolerate some data loss on a failure.

```
CREATE PROCEDURE MyProc
WITH NATIVE_COMPILATION, SCHEMABINDING, EXECUTE AS OWNER
AS BEGIN ATOMIC WITH
(DELAYED_DURABILITY = ON,
TRANSACTION ISOLATION LEVEL = SNAPSHOT,
LANGUAGE = N'us_english')
        -- Insert T-SQL here…
END
```

```
ALTER DATABASE Hekaton
SET DELAYED_DURABILITY = FORCED
GO
```

Delayed Durability

Applicable for both memory-optimized and normal tables

Durability automatically managed by System behind the scenes

Transaction flush guaranteed if:

- Durable transaction is executed
- Manually execute sp_flush_log

All System transactions are durable

Note: Data loss possibility if the server goes down between the time a transaction is committed and its log record if flushed

Diagnostics

Log not being truncated

- sys.databases log_reuse_wait_desc='XTP_CHECKPOINT'.

DMV's / DMF's

- sys.dm_db_xtp_checkpoint_stats
- sys.dm_db_xtp_checkpoint_files
- sys.fn_dblog_xtp()/sys.fn_dump_log_xtp()

Manual merge

- sys.sp_merge_xtp_checkpoint_files

Database Recovery – Memory Optimized Tables

Analysis Phase	• Finds the last completed checkpoint
Data Load	• Instantiate memory optimized tables • Load from set of data/delta files from the last completed checkpoint • Parallel Load by reading data/delta files using 1 thread / file
Redo phase	• Apply the transaction log from last checkpoint • Concurrent with REDO on disk-based tables
No UNDO phase	• Since only committed transactions are logged

Parallel Data Load

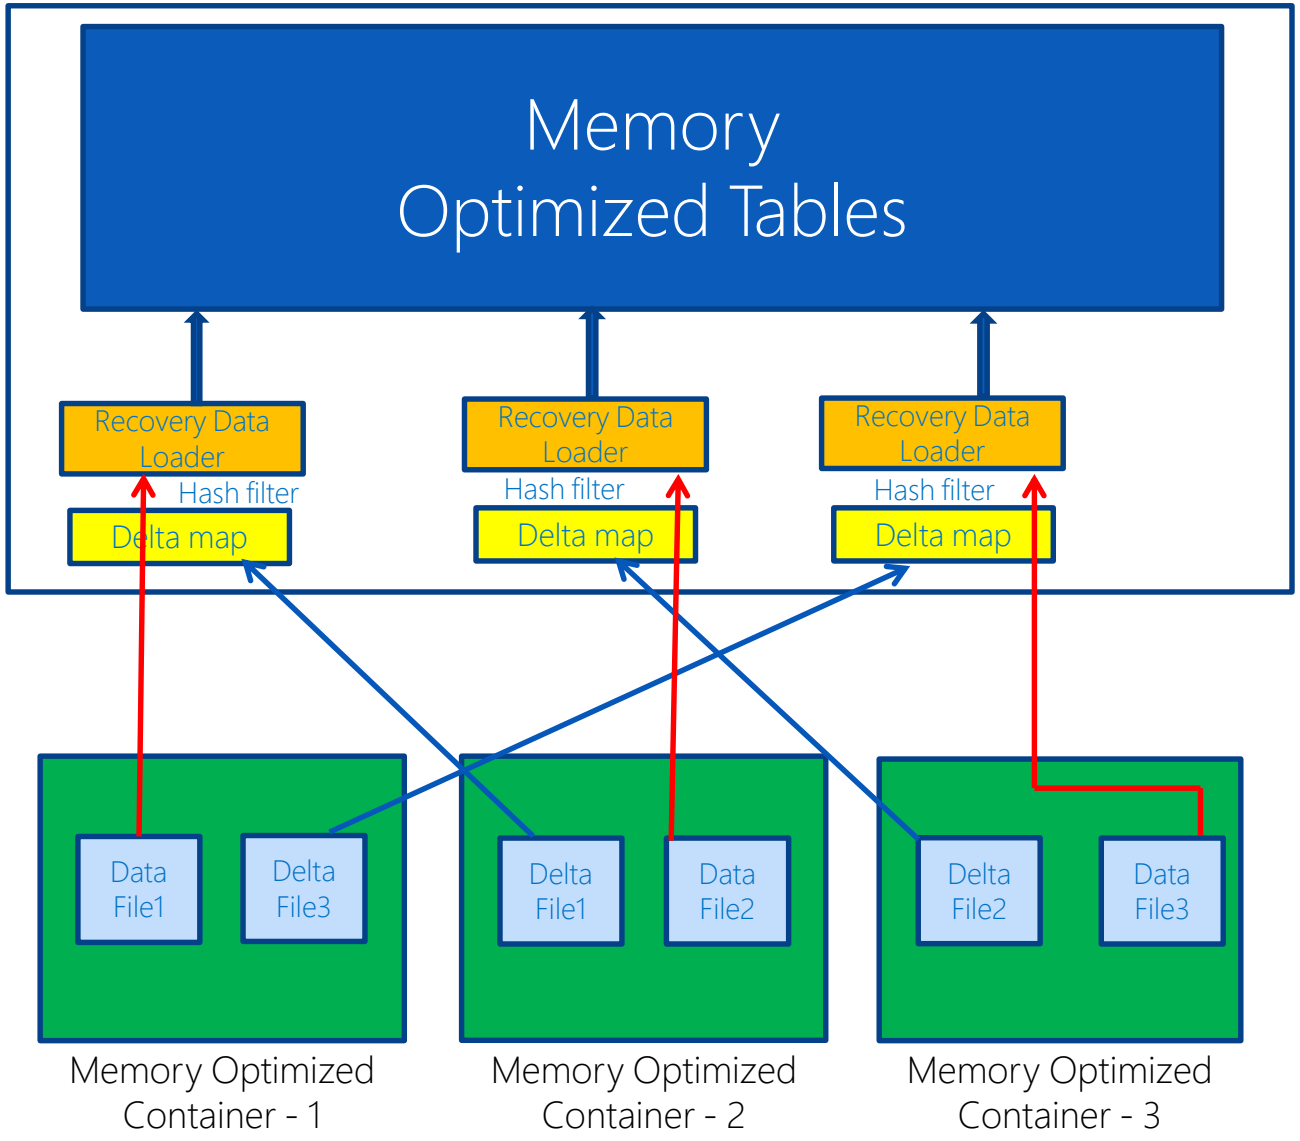

Impact on RTO

- Load speed of data
- Size of durable tables

Limit stands at 8192 data/delta file pairs per database

- If 8192 data files were fully utilized at 128MB a piece, it would be able to hold almost 1TB of data
- We recommend durable table database size to be 250GB or less

Recovery Progress and Diagnostics

Progress

- Extended Event xtp_recover_table
- Extended Event xtp_recover_done
- Errorlog

Recovery Failures

- Due to compilation
- Due to memory availability/configured to load the data
- Checksum failure on checkpoint files
- Table size limitation (checkpoint files limit is 8192)
- Other normal recovery failure reasons

Memory problem?

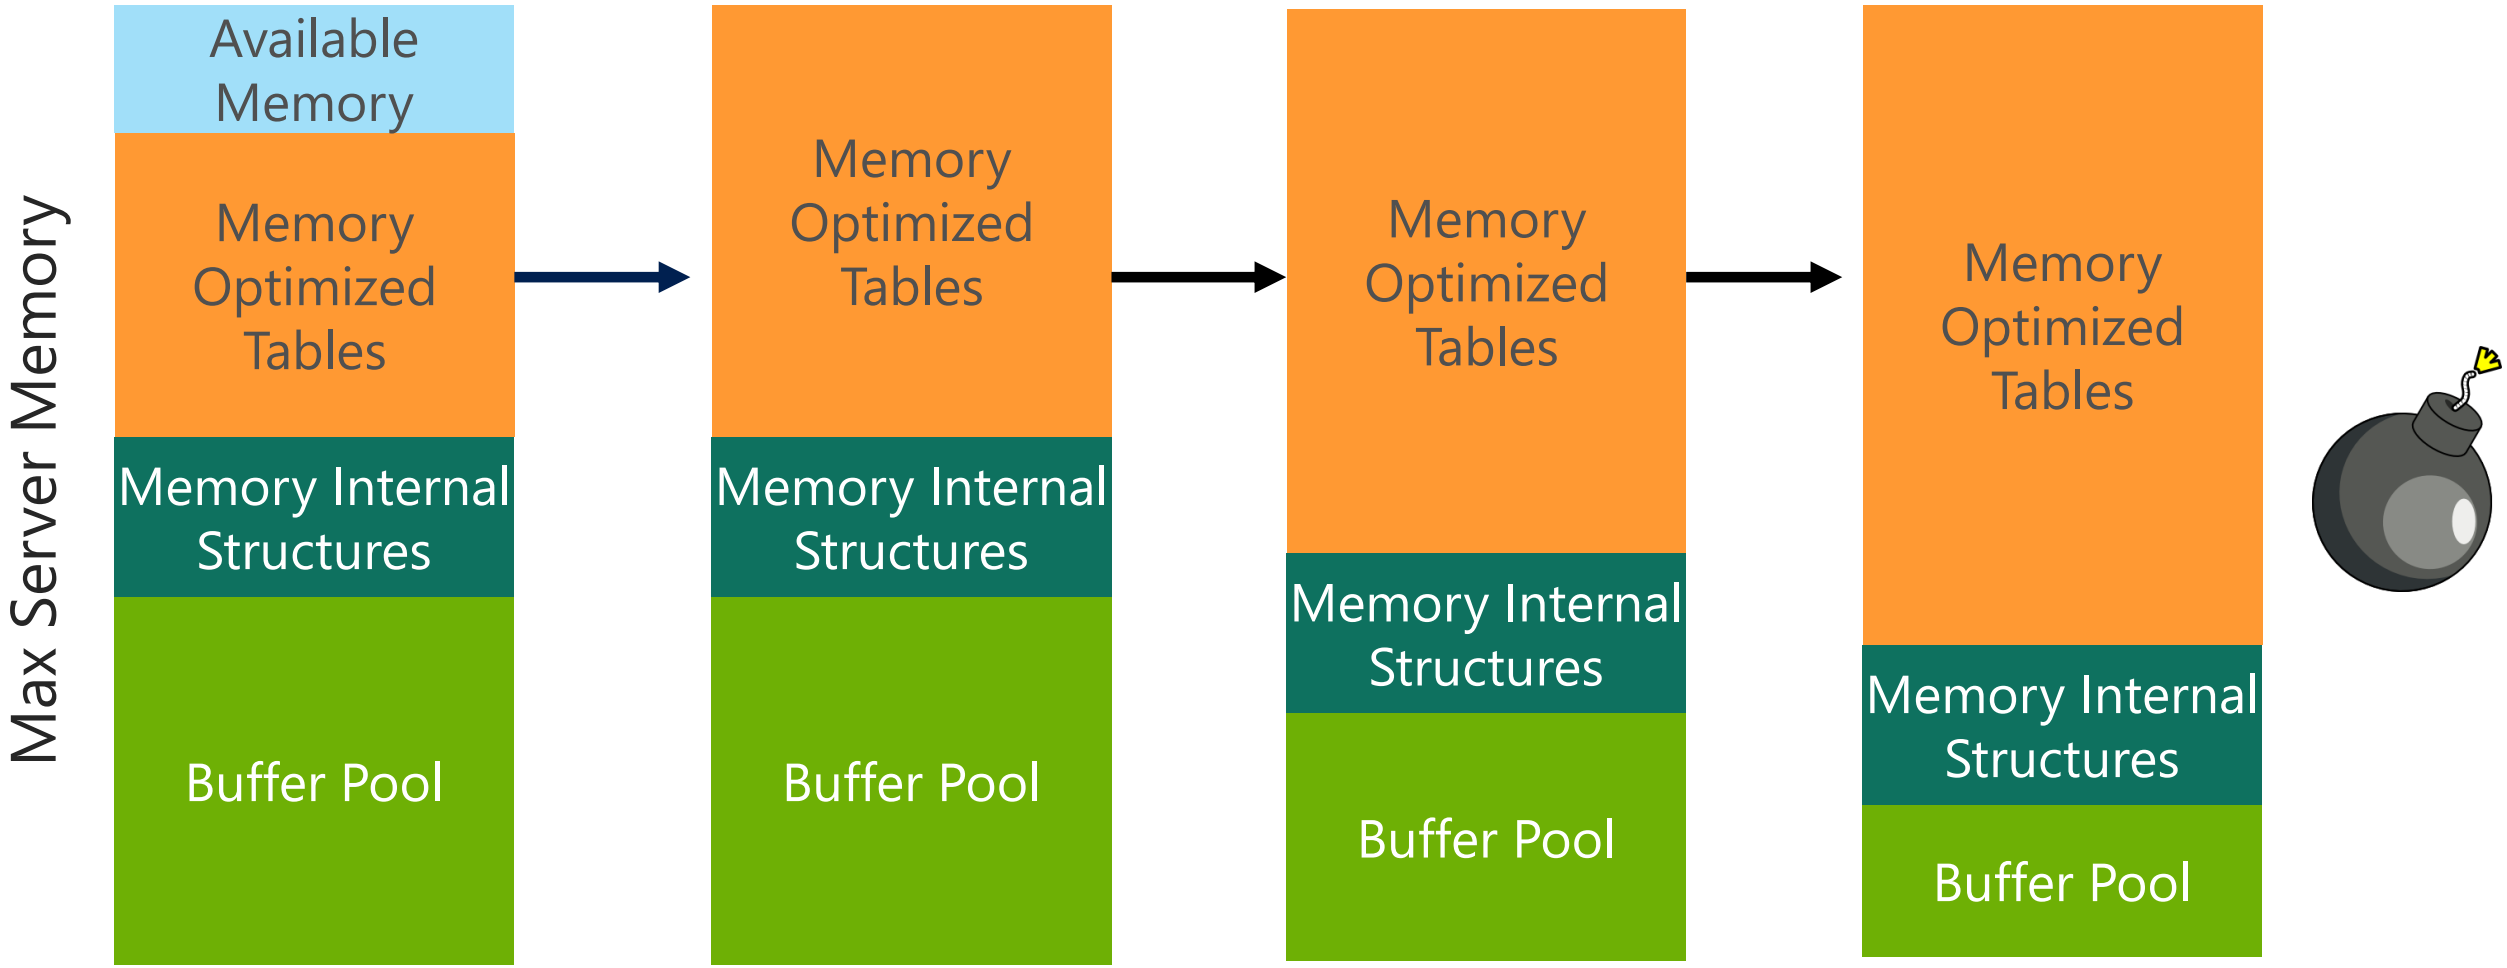

Memory Considerations

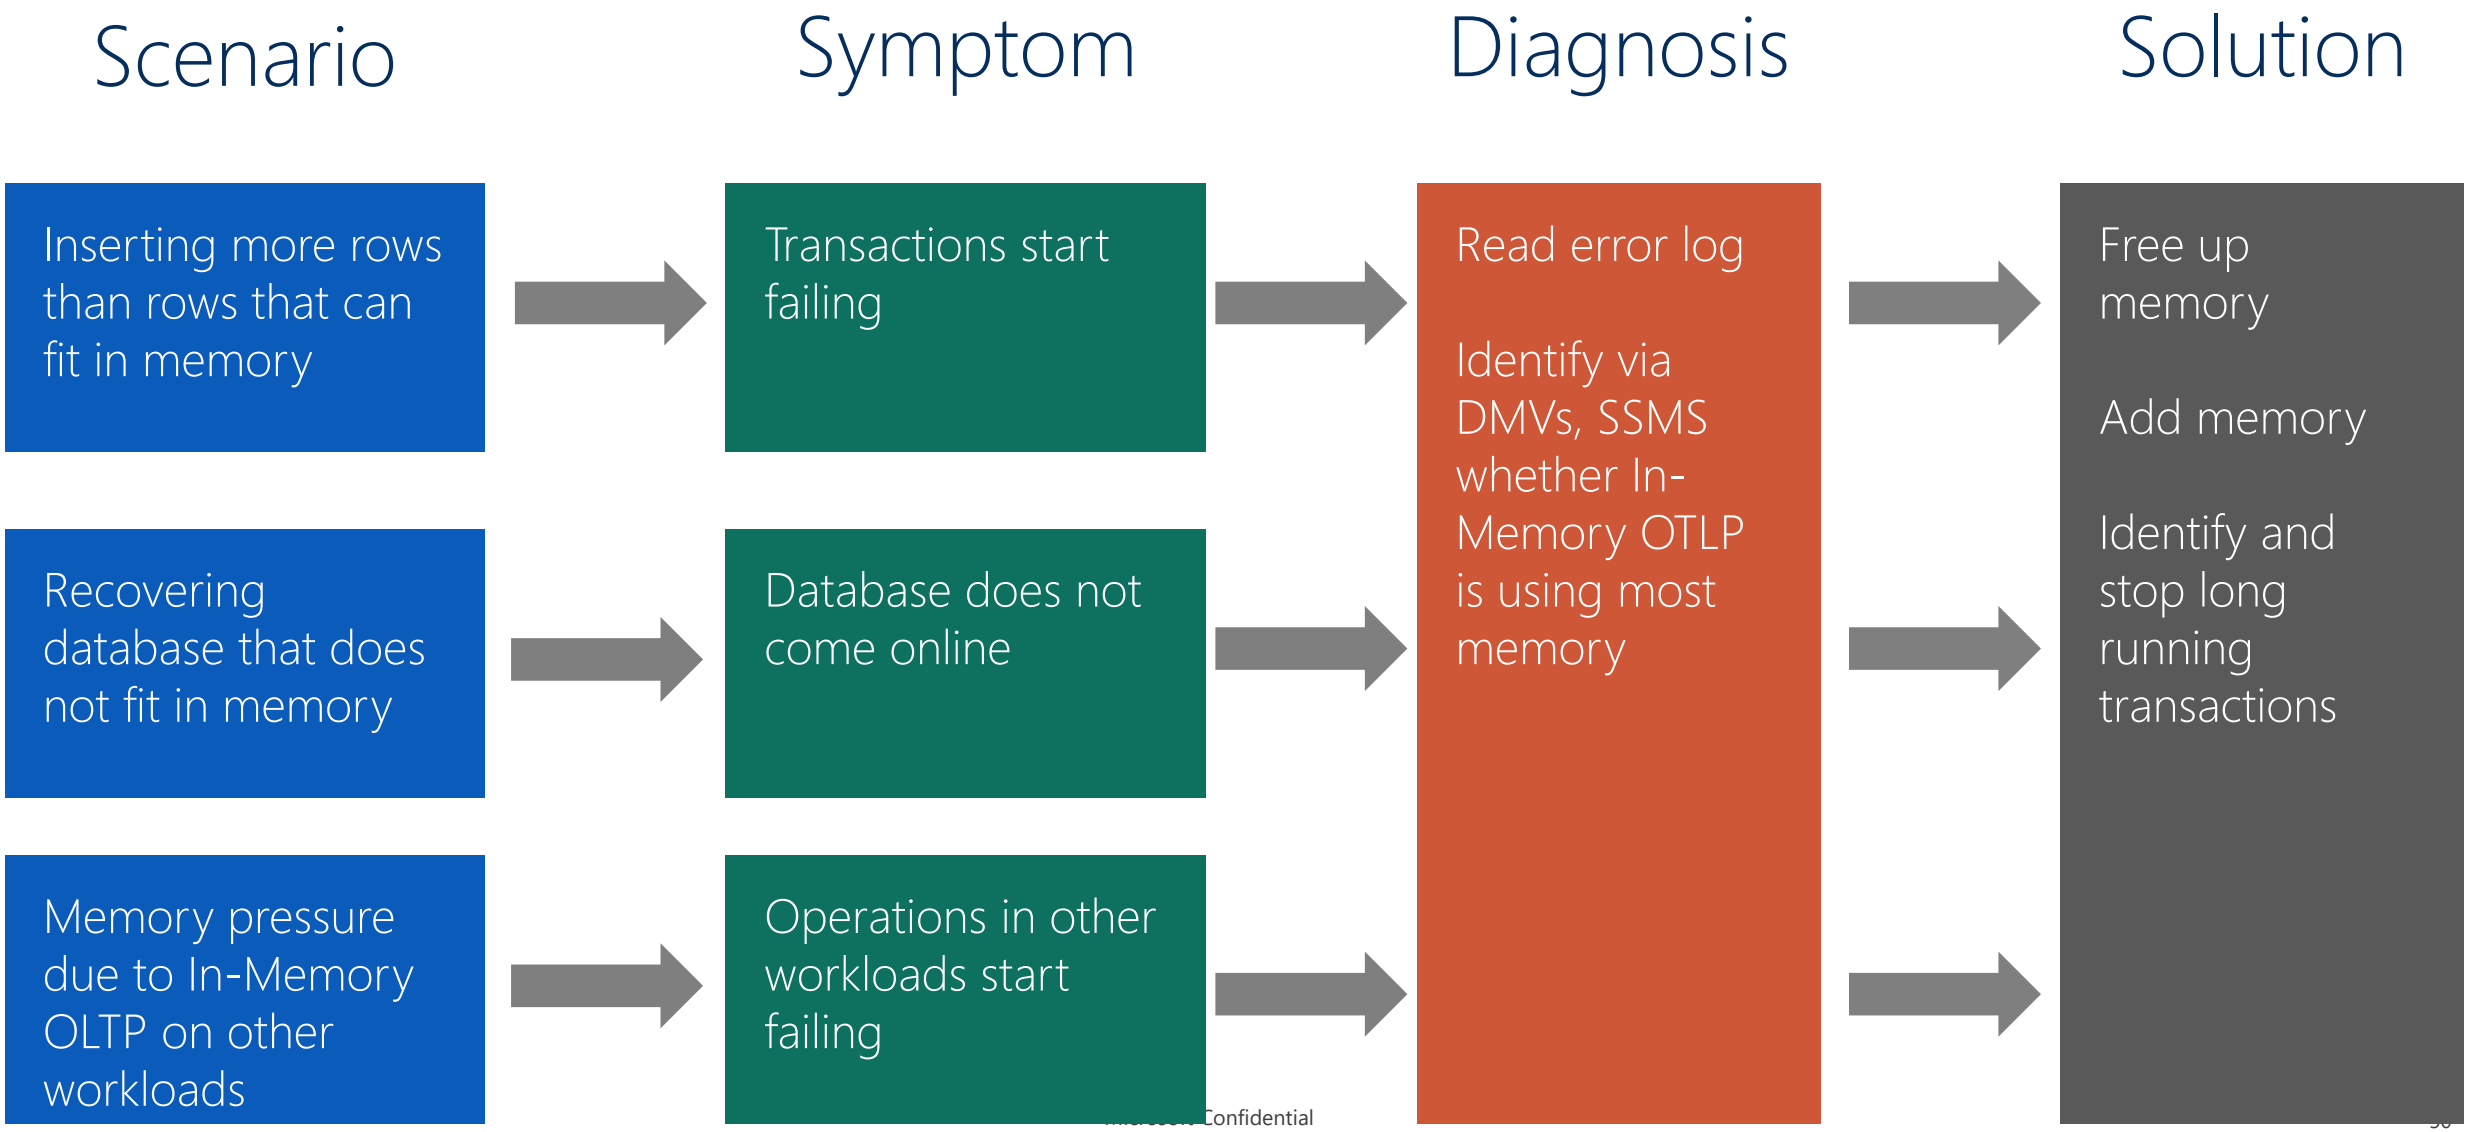

Memory Considerations

Data resides in memory at all times

- Configure SQL server with sufficient memory to store memory-optimized tables
- Failure to allocate memory will fail transactional workload at run-time
- Other SQL workloads can slow down to unacceptable performance

Sizing

- Table Size: rule of thumb 2* size of data
- Hash Indexes: 8KB * [actual bucket count]
- Non clustered index: variable size

Management

- Limit memory consumption using Resource Governor
- DMVs, Perfmon
- Freeing memory is not synchronous in most cases

Resource Governor Integration

Limits max memory allocation for In-Memory OLTP

Prevents performance degradation of regular SQL workloads

Implementation:

- Define dedicated pool for In-Memory OLTP
- Bind DB to resource pool defined
  → *sys.sp_xtp_bind_db_resource_pool*
- Set database offline and Online

Internal pool target memory	%available of dedicated pool before OOM notification
<= 8 GB	70%
<= 16GB	75%
<= 32GB	80%
<= 96GB	85%
>= 96 GB	90%

Memory Architecture

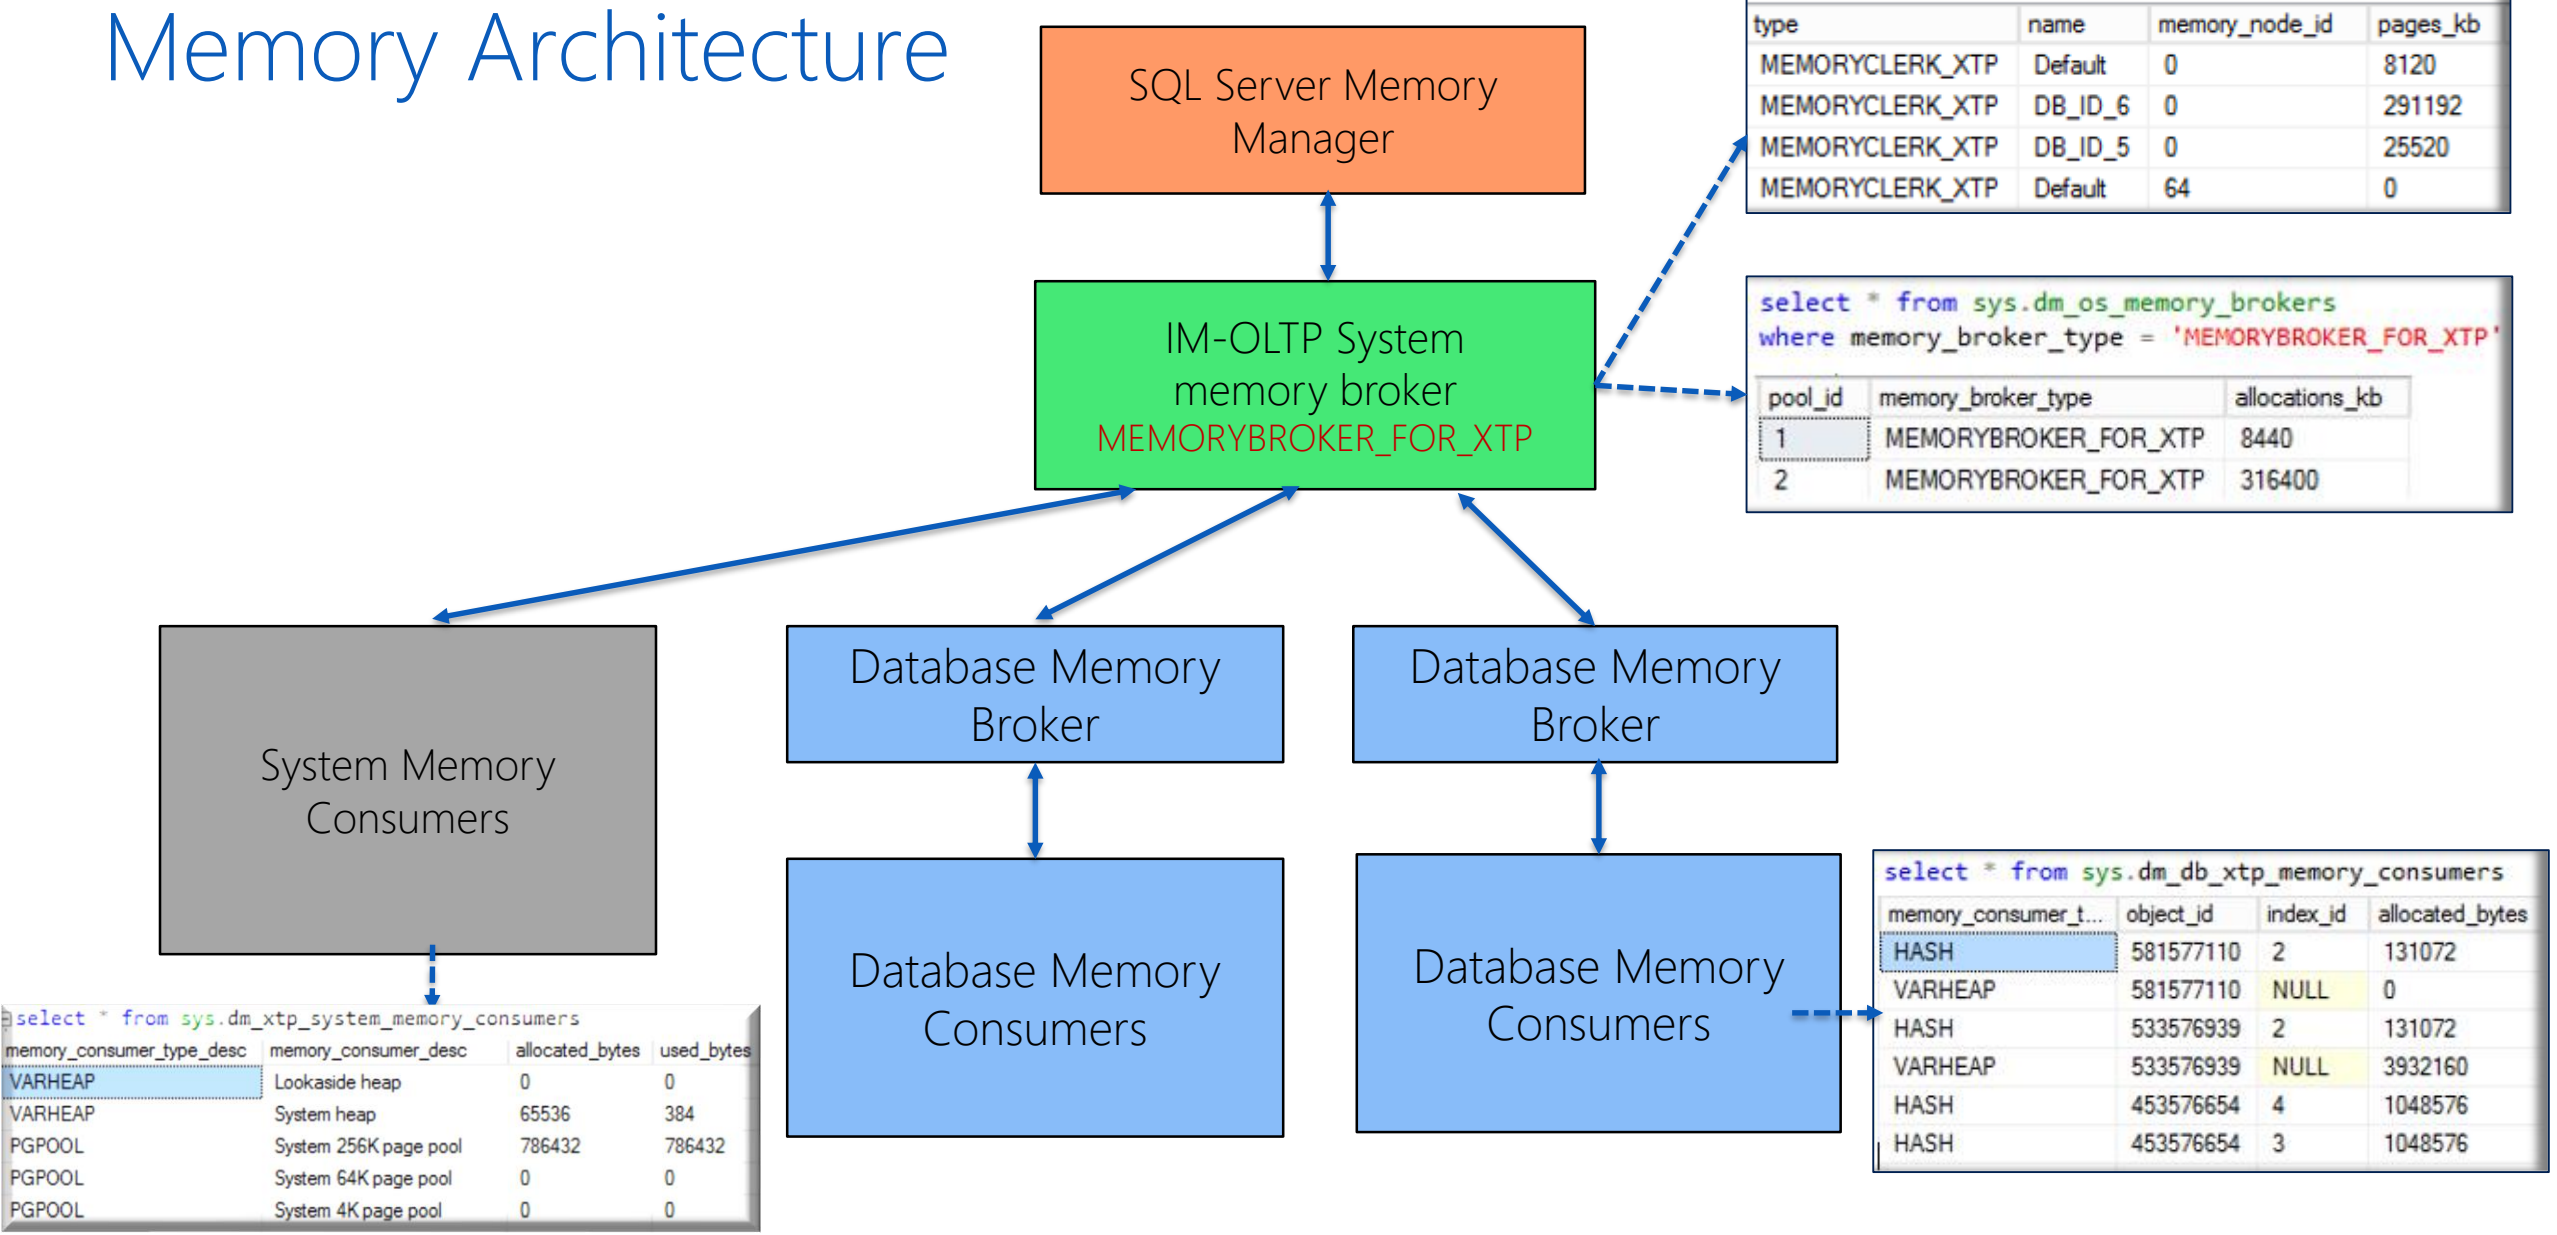

Monitoring

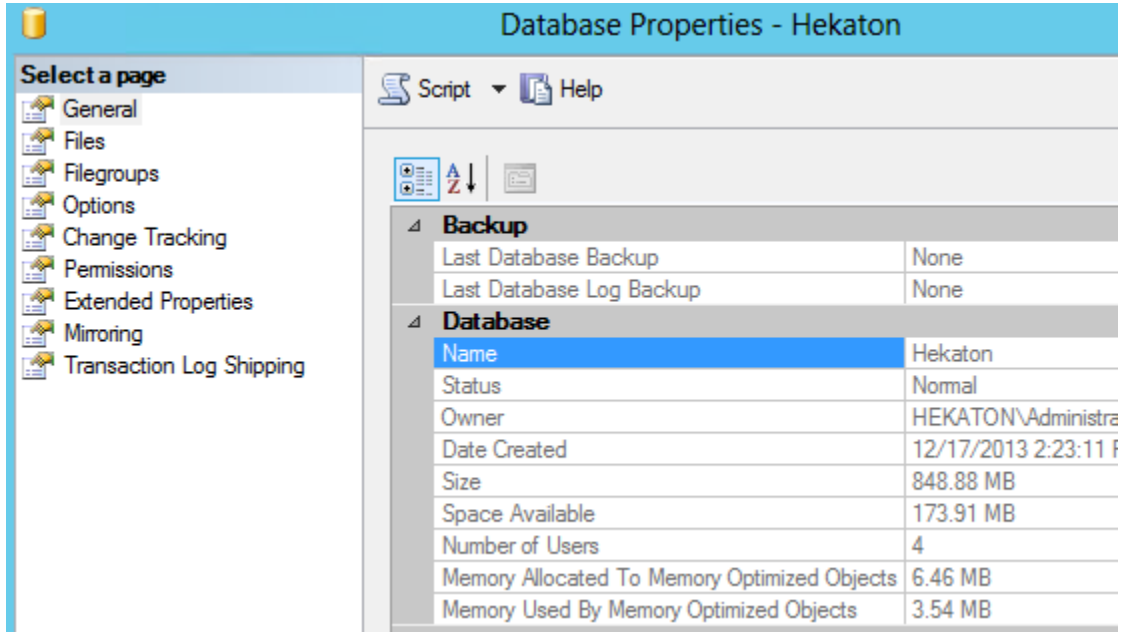

Memory Usage By Memory Optimized Objects [Hekaton]
on HEKATON at 1/13/2014 4:21:23 PM

Microsoft SQL Server 2014

This report provides detailed data on the utilization of memory space by memory optimized objects within the Database.

Total Memory Allocated To Memory Optimized Objects:	6.46 MB

Total Memory Usage By Memory Optimized Objects (MB)

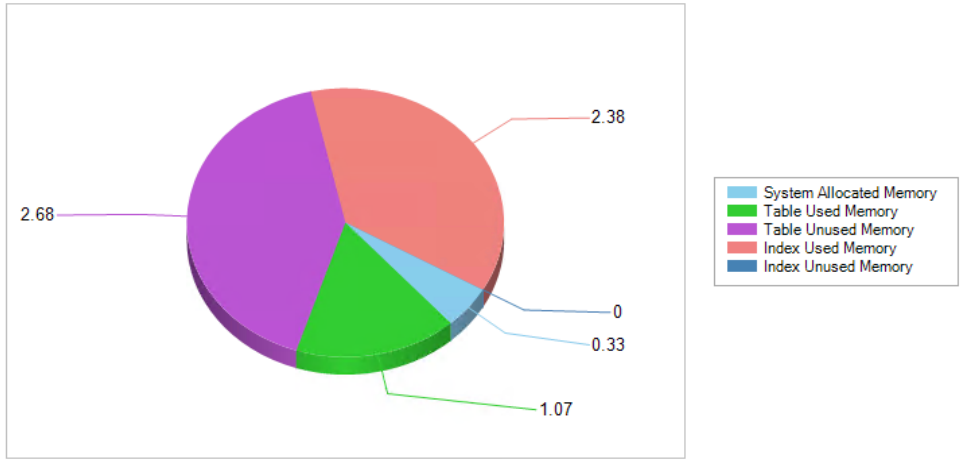

Memory Usage Details for Memory Optimized Tables (MB)

Table Name	Table Used Memory	Table Unused Memory	Index Used Memory	Index Unused Memory
Customer	0.00	0.00	2.13	0.00
SalesOrder_MemOpt_SchemaOnly	0.00	0.00	0.13	0.00
SalesOrder_MemOpt	1.07	2.68	0.13	0.00

Memory Usage	DMV
Memory Optimized Table size	sys.dm_db_xtp_table_memory_stats
Database IM-OLTP memory usage	sys.dm_db_xtp_memory_consumers
Non-Database IM-OLTP memory usage	sys.dm_xtp_system_memory_consumers
Overall IM-OLTP memory usage	sys.dm_os_memory_brokers sys.dm_os_memory_clerks

Multi-Version Lock Free Transactions

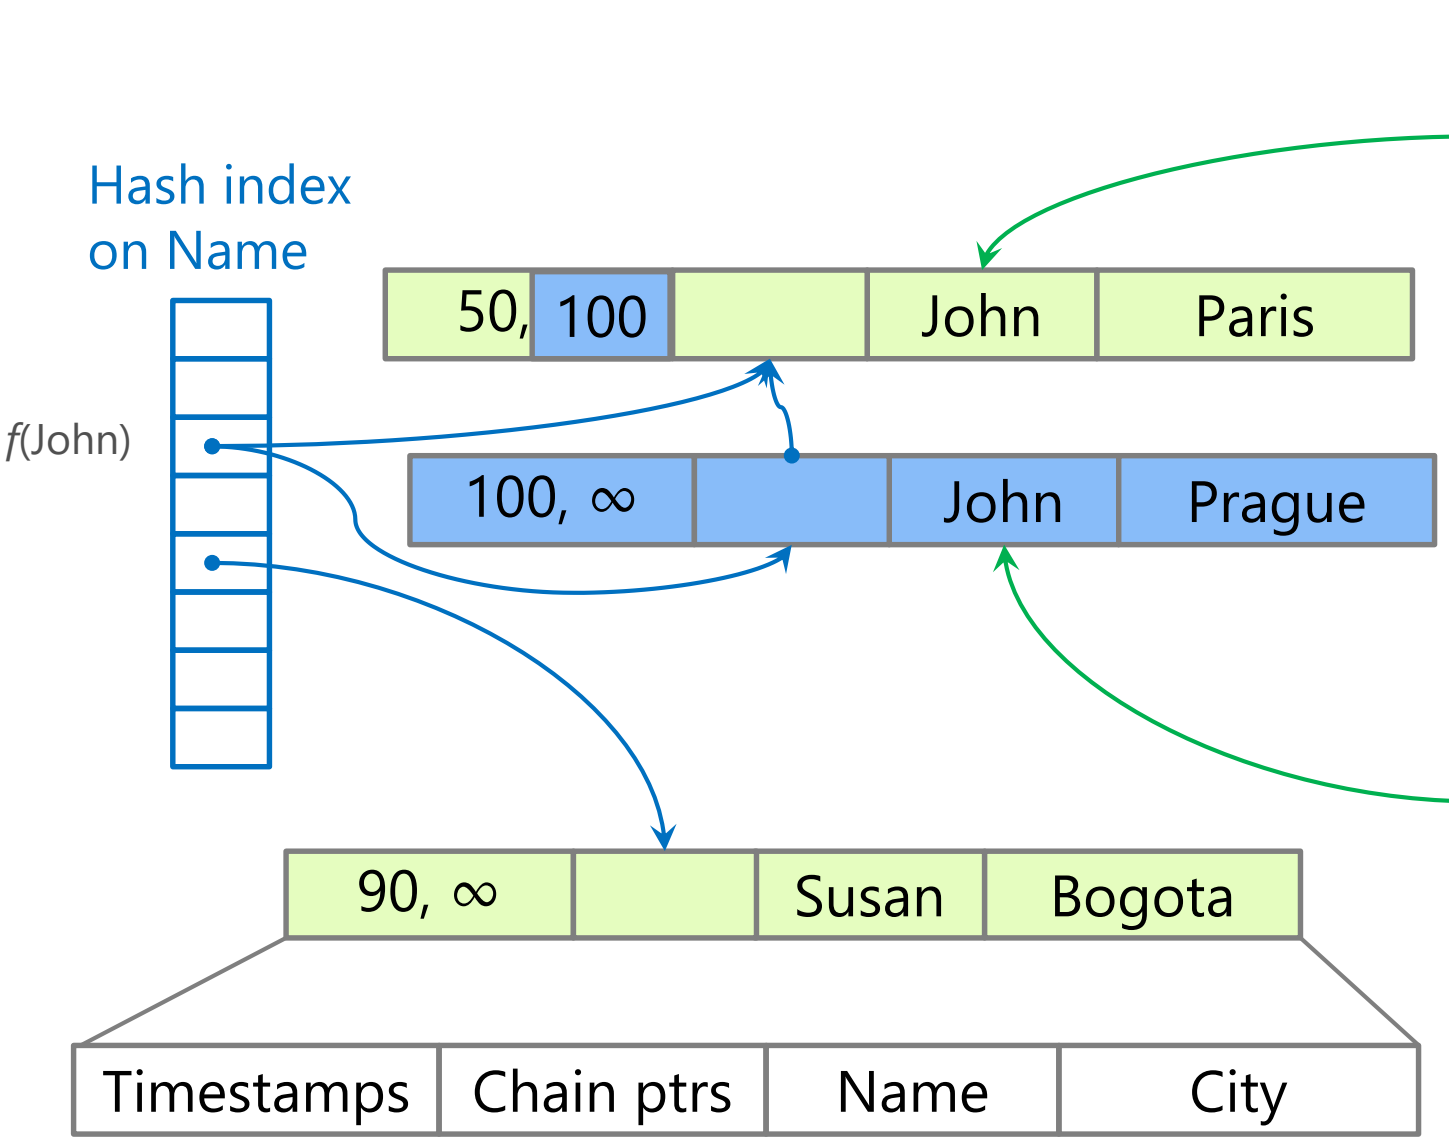

Transaction 99: Running compiled query

SELECT City WHERE Name = 'John'

Simple hash lookup returns direct pointer to 'John' row

Background operation will unlink and deallocate the old 'John' row after transaction 99 completes.

Transaction 100:

UPDATE City = 'Prague' where Name = 'John'

No locks of any kind, no interference with transaction 99

Garbage Collection

Stale Row Versions

- Updates, deletes, and aborted insert operations create row versions that (eventually) are no longer visible to any transaction
- Slow down scans of index structures
- Create unused memory that needs to be reclaimed (i.e. Garbage Collected)

Garbage Collection (GC)

- Analogous to version store cleanup task for disk-based tables to support Read Committed Snapshot (RCSI)
- System maintains ‘oldest active transaction’ hint

Cooperative Garbage Collection

- Scanners can remove expired rows when found
- Offloads work from GC thread
- Ensures that frequently visited areas of the index are cleaned regularly

A row needs to be removed from all indexes before memory can be freed

- Garbage collection is most efficient if all indexes are frequently accesses

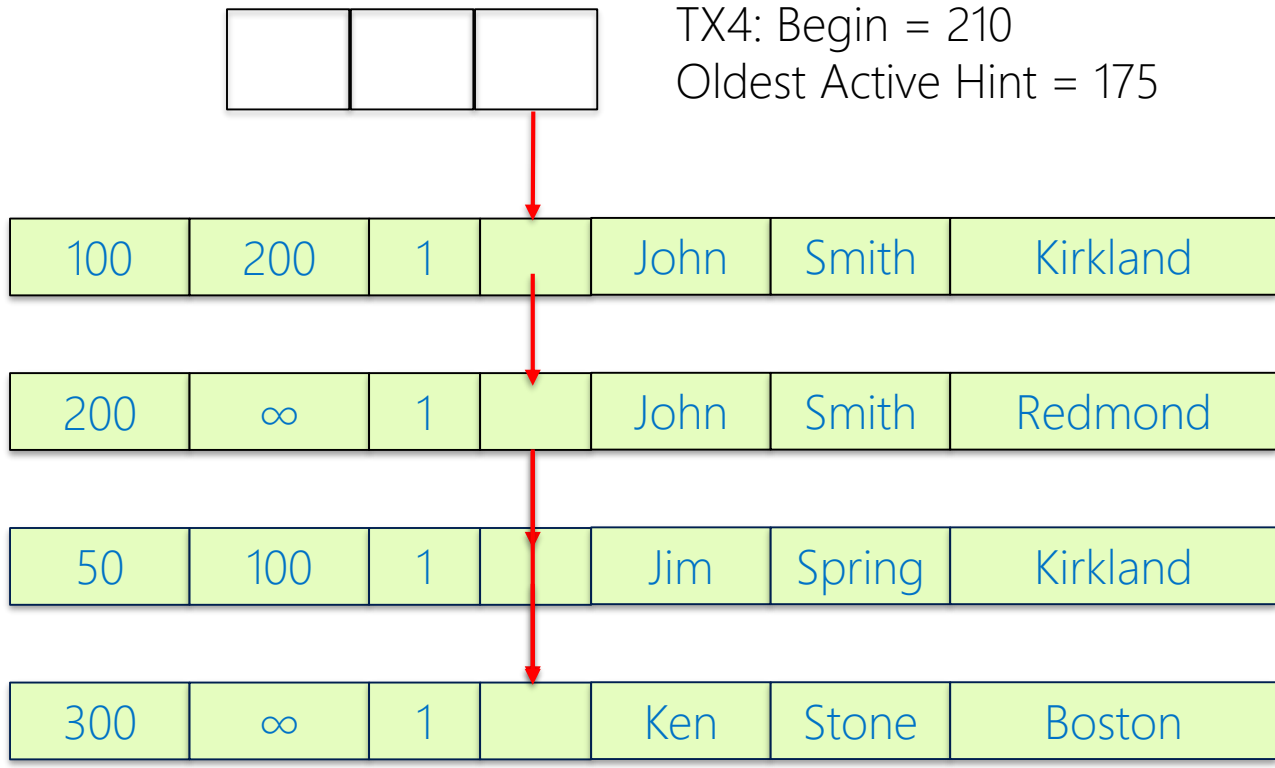

Garbage Cleanup

- Non-blocking, Cooperative, Efficient, Responsive, Scalable
- Active transactions work cooperatively and pick up parts of GC work
- A dedicated system thread for GC
- GC thread wakes every minute or when the number of committed transactions since it last ran exceed 1024

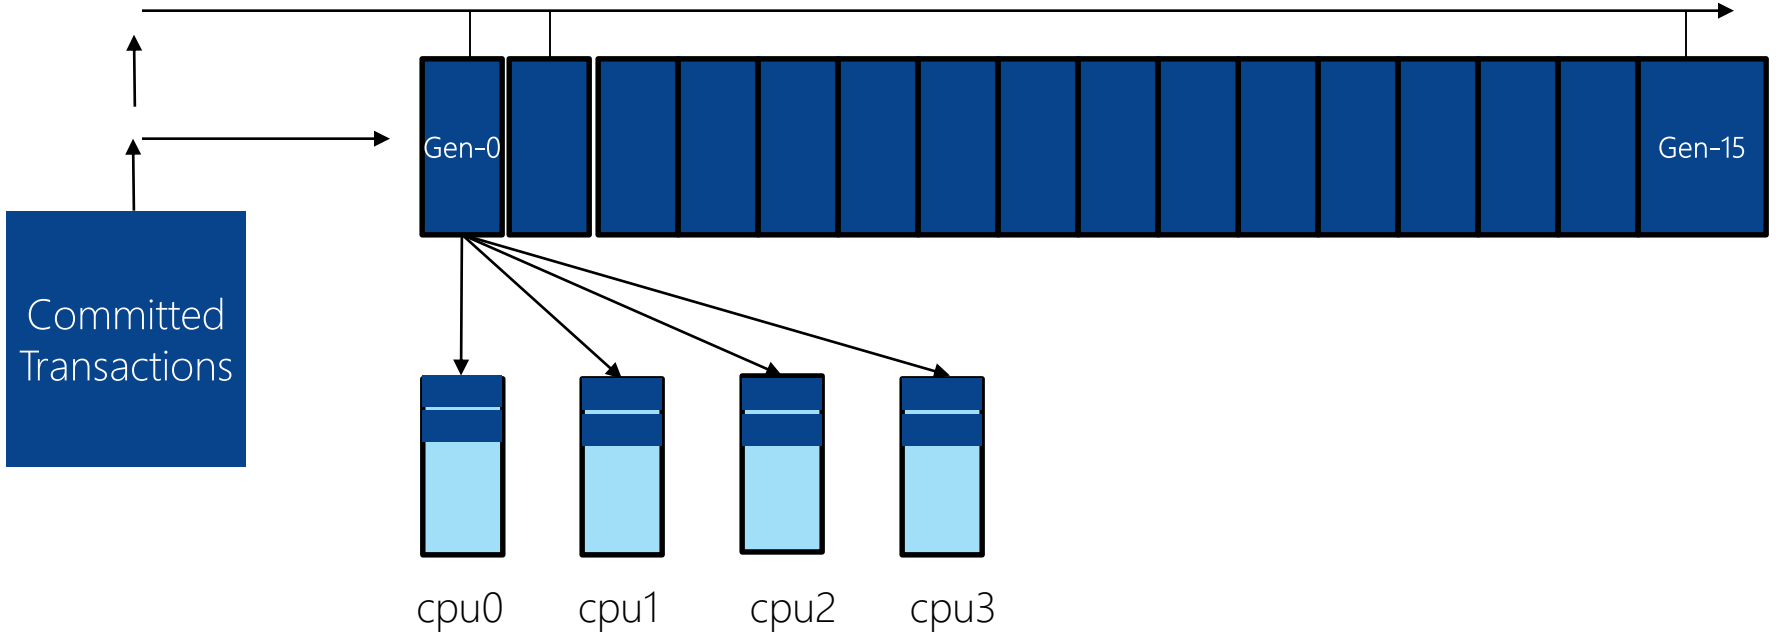

GC Diagnostics

DMV's

- sys.dm_xtp_gc_stats
- sys.dm_xtp_gc_queue_stats

Performance counters – XTP Garbage Collection

Extended Events – xtpengine.gc_*

Concurrency Control

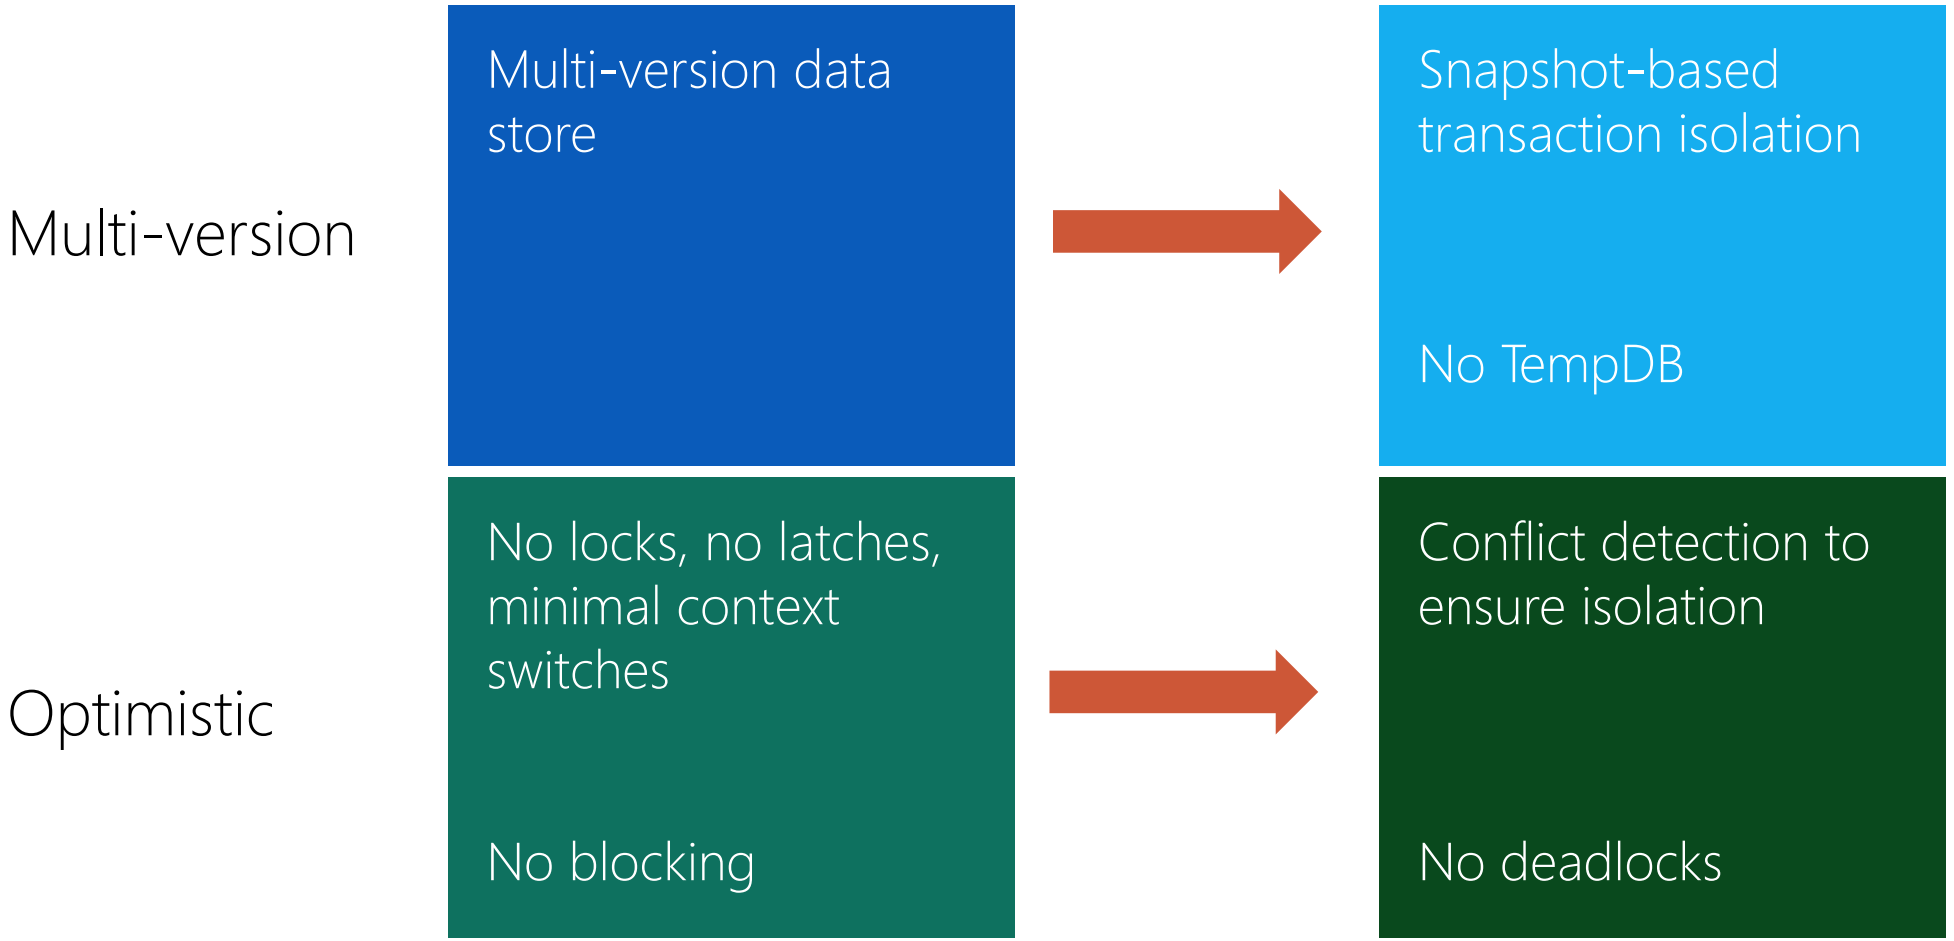

Isolation Levels Supported

SNAPSHOT	• Reads are consistent as of start of the transaction • Writes are always consistent
REPEATABLE READ	• Read operations yield same row versions if repeated at commit time
SERIALIZABLE	• Transaction is executed as if there are no concurrent transactions – all actions happen at a single serialization point (commit time)

Write Conflict

Time	Transaction T1 (SNAPSHOT)	Transaction T2 (SNAPSHOT)
1	`BEGIN`	
2		`BEGIN`
3		`UPDATE t SET c1='value2' WHERE c2=123`
4	`UPDATE t SET c1='value1' WHERE c2=123 (write conflict)`	

First writer wins

Transaction Processing

- Reads snapshot rows as of start of transaction
- Writes not visible to other transactions
- Transaction doomed if row modified is updated by others
- Commit dependency can exist

- Obtains End_TS for trans
- Determines if it can be c

- Transaction written to the Log
- Commit dependencies cleared

Snapshot Isolation	Repeatable Read	Serializable
• No Validation required	• Require "read stability" as of TS end time	• "read stability" • "phantom avoidance"

Validation Errors and Retry Logic

Err Number	Err Message
41302	The current transaction attempted to update a record that has been updated since the transaction started
41305	The current transaction failed to commit due to a repeatable read validation failure
41325	The current transaction failed to commit due to a serializable validation failure
41301	A previous transaction that the current transaction took a dependency on has aborted, and the current transaction can no longer commit

Failures causing transaction abort

Write conflicts, validation failures

Aborted transactions need to be retried

Solution: implement retry logic

Server-side retry avoids changes to client apps

Retry Logic for Transaction Failures

```
CREATE PROCEDURE usp_my_procedure @param1 type1, @param2 type2, ...
AS
BEGIN
 DECLARE @retry INT = 10
 WHILE (@retry > 0)
 BEGIN
  BEGIN TRY

   EXEC usp_my_native_proc @param1, @param2, ...

   SET @retry = 0
  END TRY
  BEGIN CATCH
   SET @retry -= 1
   IF (@retry > 0 AND error_number() IN (41302, 41305, 41325, 41301, 1205))
    IF (@@TRANCOUNT>0) ROLLBACK TRANSACTION
   ELSE
    THROW
  END CATCH
 END
END
```

Isolation Level Combinations Supported

Disk-based	Memory optimized	Usage recommendations
READCOMMITTED	SNAPSHOT	• Baseline combination – most cases that use READCOMMITTED today
READCOMMITTED	REPEATABLEREAD/ SERIALIZABLE	• Data migration • In-Memory OTLP only Interop
REPEATABLEREAD/ SERIALIZABLE	SNAPSHOT	• Memory-optimized table access is INSERT-only • Useful for data migration and if no concurrent writes on memory-optimized tables (e.g., ETL)

Unsupported isolation level combinations (V1)

Disk-based	Memory optimized
SNAPSHOT	Any isolation level
REPEATABLEREAD/ SERIALIZABLE	REPEATABLEREAD/ SERIALIZABLE