

Query Compilation and Execution

Introduction



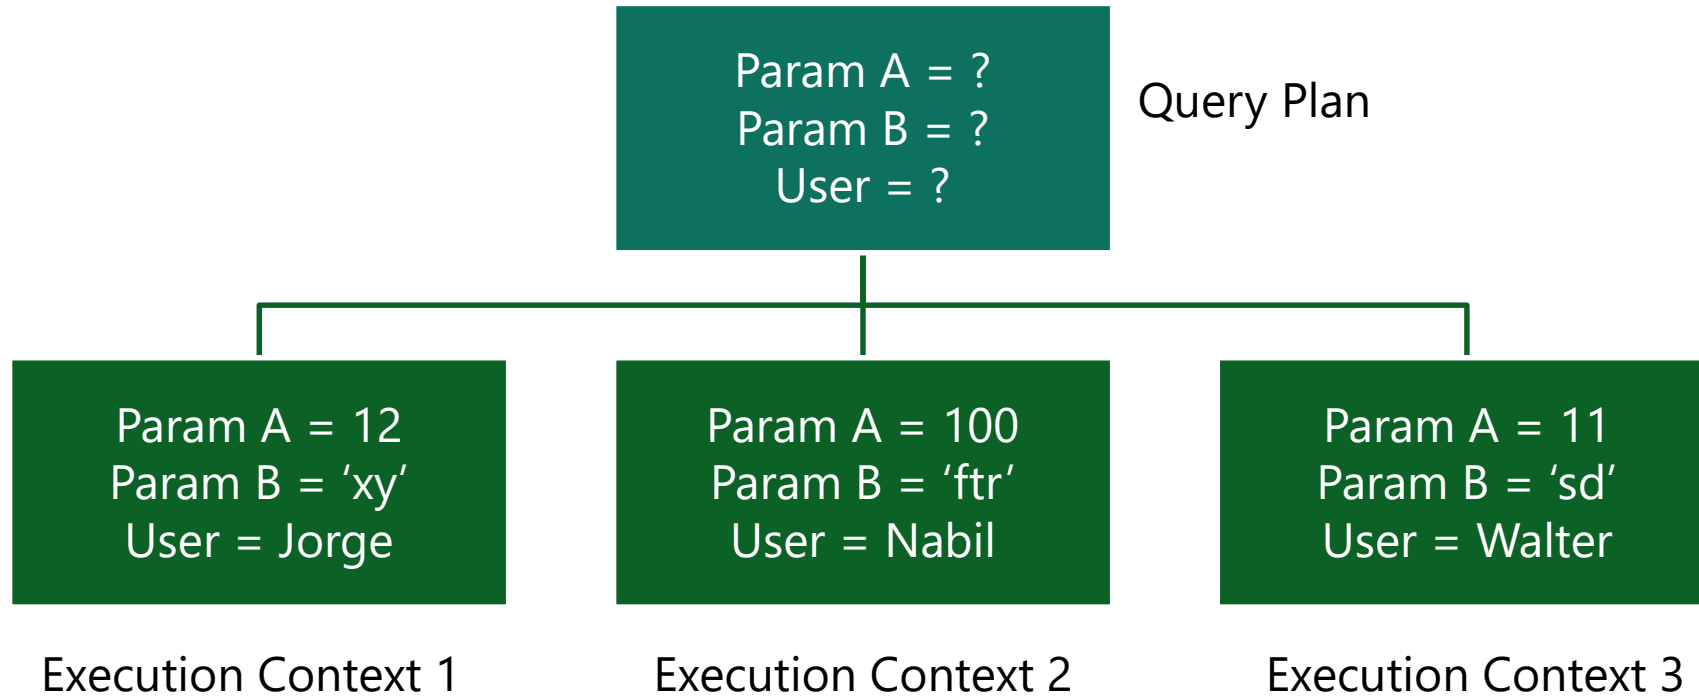
Query Compilation

- Query compilation is the process of choosing a good enough execution plan that the Optimizer has to act in the short amount of time
 - Parse a query into a tree representation
 - Normalize and validate the query
 - Evaluate possible query plans
 - Pick a good enough plan, based on cost

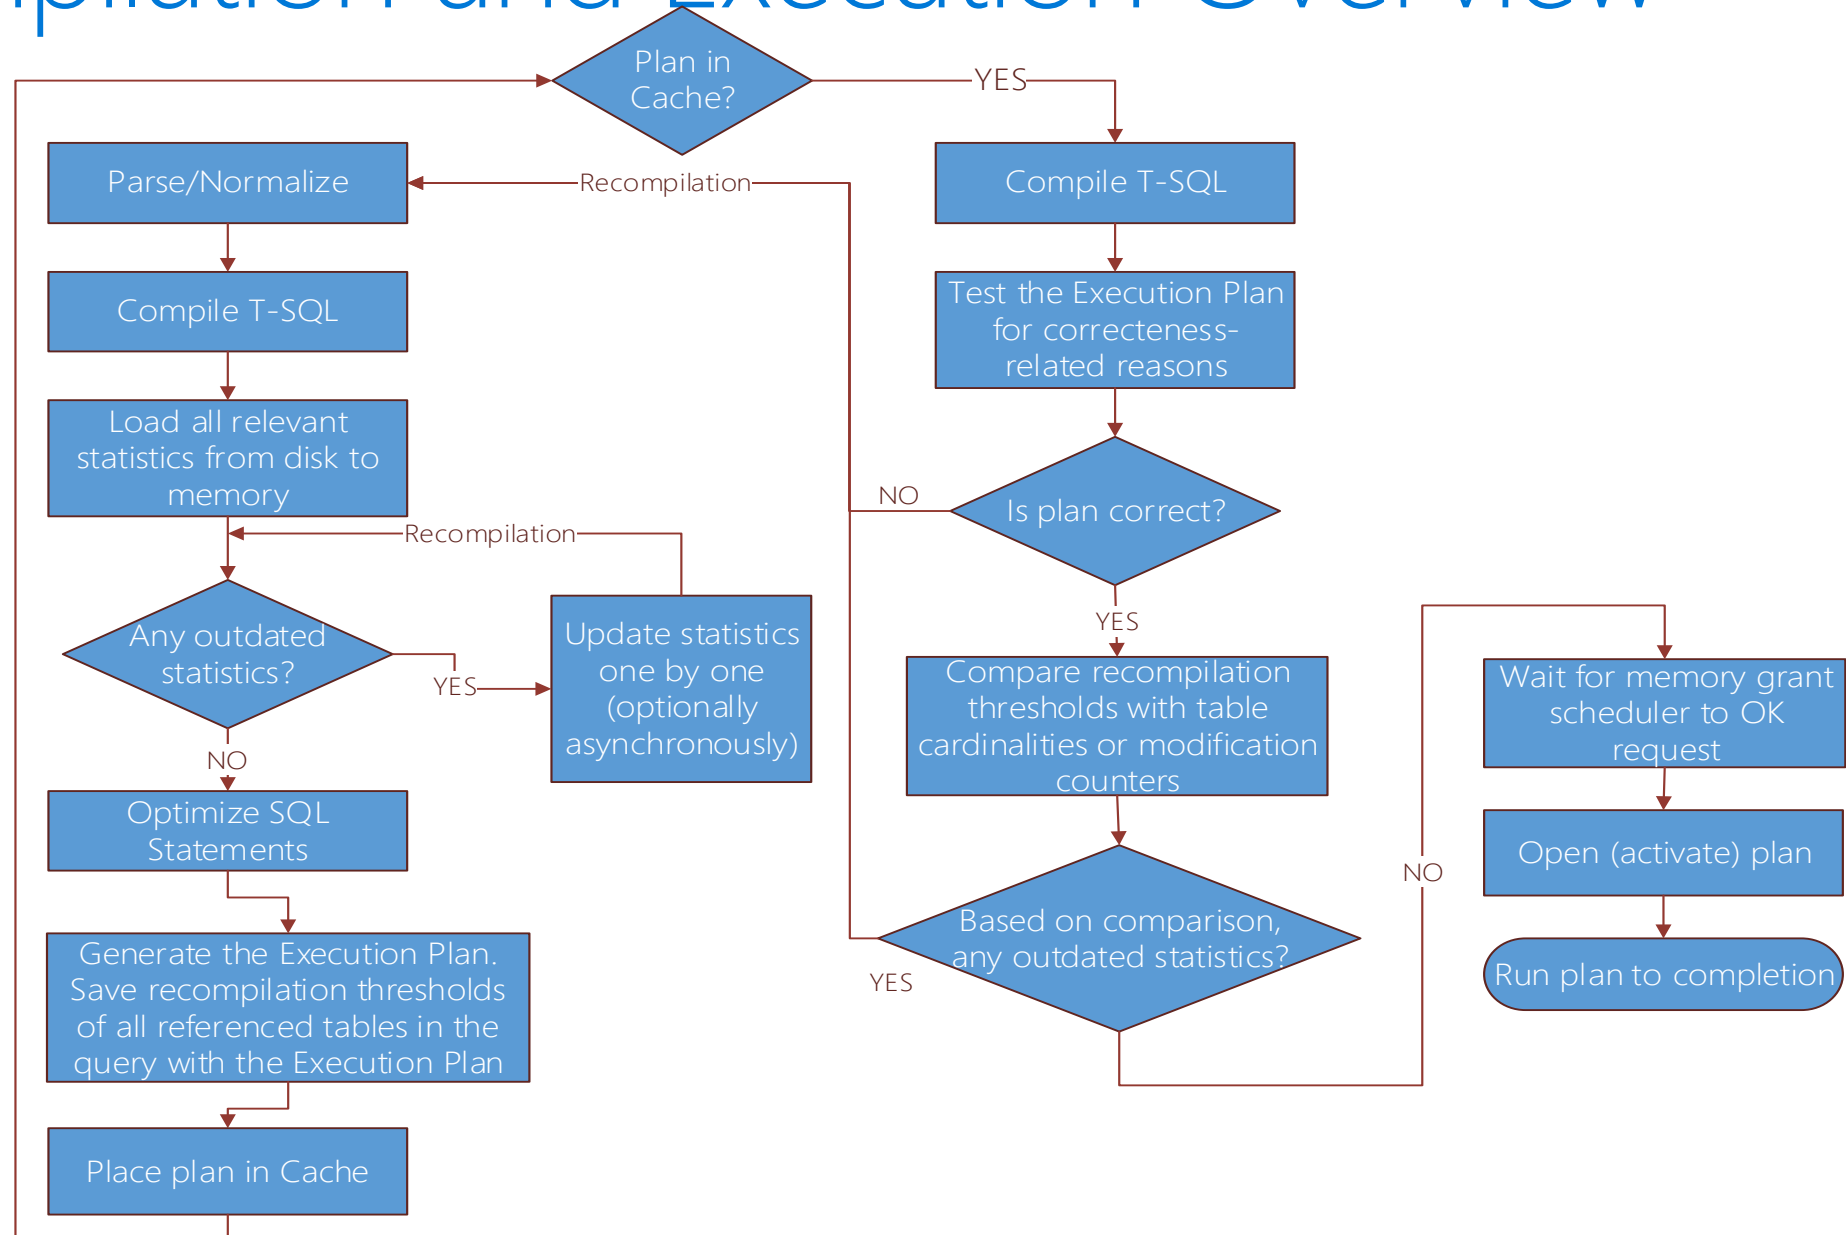
Query Execution

- Query execution is the process of executing the plan that is created during query compilation and optimization
 - Not necessarily performed directly after query compilation
 - May trigger a query recompile
 - Compilation versus recompilation
 - Query recompiles may occur because of correctness-related reasons or plan optimality related reasons

Query Plans and Execution Contexts



Compilation and Execution Overview

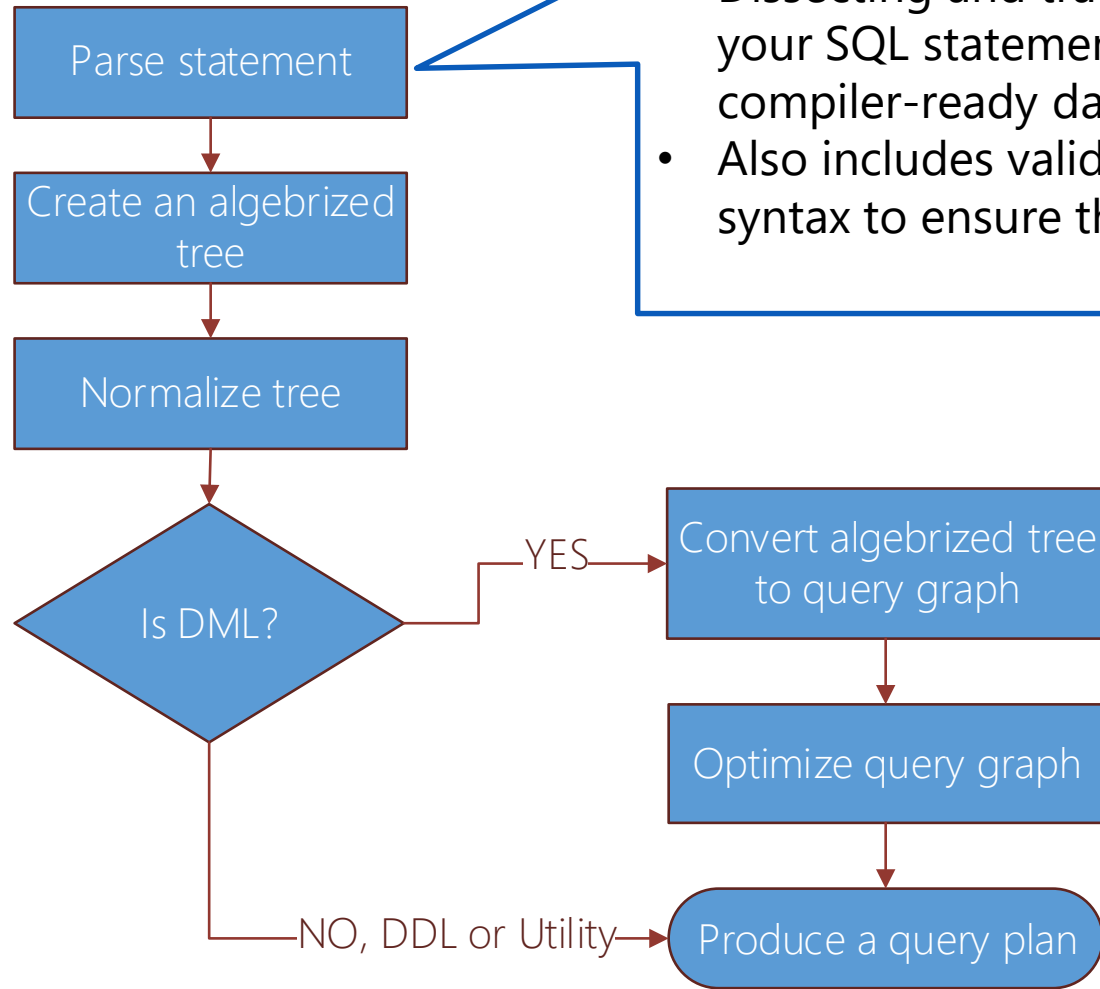


Query Processing

Query Plan Generation



First Stage - Compilation



- Dissecting and transforming your SQL statements into compiler-ready data structures
- Also includes validation of the syntax to ensure that it is legal

- Object binding, which includes verifying that the tables and columns exist, and expanding the views
- Loading the metadata information
- Syntax based optimizations

Second Stage – Optimization

Stage 1 - **Trivial Plan**

- Non-cost-based optimizer
 - Statistics are loaded and validated at this stage
 - This step generates plans for which there are no alternatives that require a cost-based decision
- Example: INSERT statement with a VALUE has only one possible plan

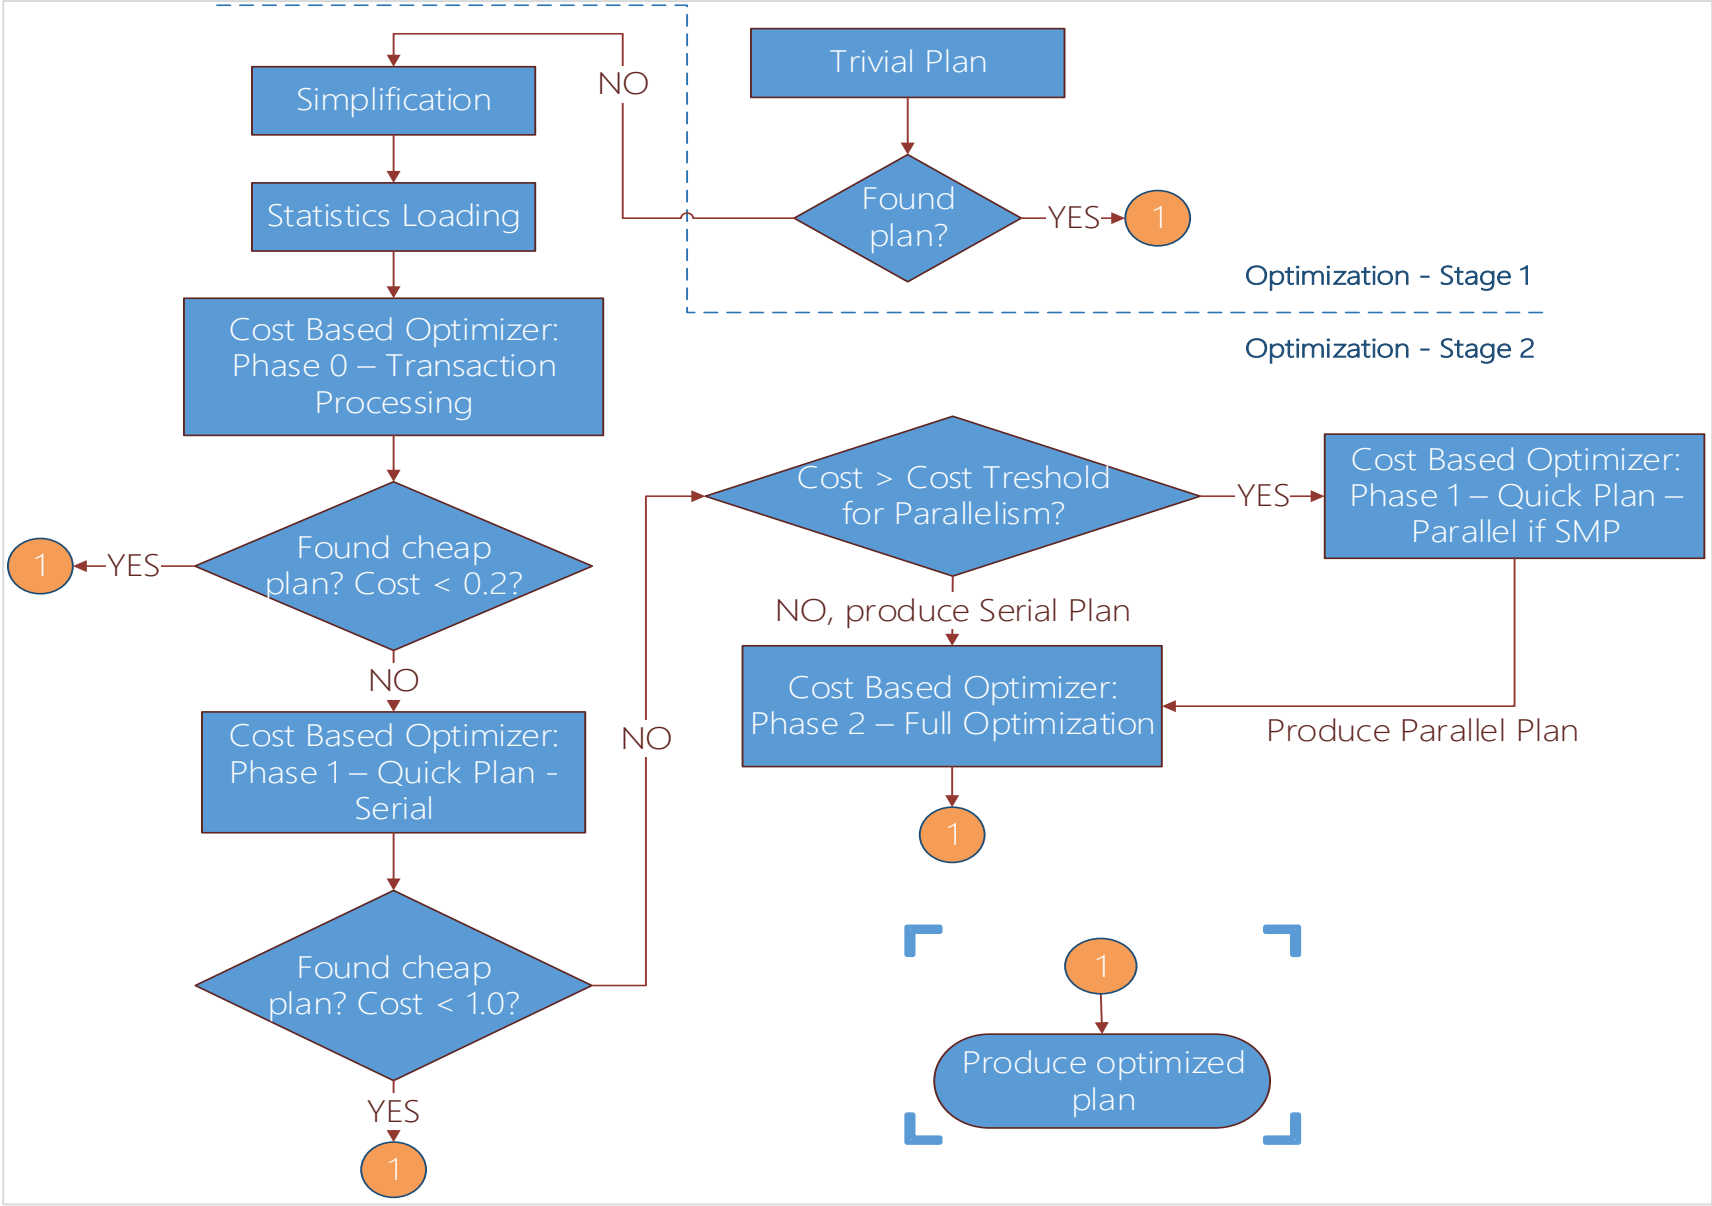
Stage 2 - **Simplification**

- Cost-based optimizer if previously unsuccessful
- Has three phases:
 - Phase 0 - Transaction Processing
 - Phase 1 - Quick Plan
 - Phase 2 - Full Optimization

Second Stage – Optimization (Continued)

- Query Optimizer may use various inputs (for example, stats, parameterized values) to reason about density/selectivity and cardinality
- Evaluates the cost of various plan alternatives and gives you the best one based on the provided information
- If it gets it wrong – you get what is perceived as *inefficient plan*
- Sources of inefficiency:
 - Bad cardinality estimation?
 - Look at plan
 - Parameter-sensitive plans?
 - Dynamic un-parameterized SQL Server?
 - Bad physical database design?
 - Missing indexes?

Second Stage – Optimization Overview



Query Processing

Statistics



Recompilation Threshold (RT)

The RT is a mechanism used by SQL Server to determine if a table has changed enough to force a recompile of a query plan to determine if a more efficient plan is available for the current data distribution

The *threshold crossing test* is performed to decide whether to recompile a query plan:

- $| \text{colmodctr}(\text{current}) - \text{colmodctr}(\text{snapshot}) | \geq \text{RT}$

If there are no statistics, or nothing is *interesting*, then table cardinality is used:

- $| \text{cardinality}(\text{current}) - \text{cardinality}(\text{snapshot}) | \geq \text{RT}$

Recompilation Threshold Calculation

Permanent table

- If $n \leq 500$, $RT = 500$.
- If $n > 500$, $RT = 500 + 0.20 * n$

Temporary table

- If $n < 6$, $RT = 6$.
- If $6 \leq n \leq 500$, $RT = 500$
- If $n > 500$, $RT = 500 + 0.20 * n$

Table variable

- RT does not exist

With TF2371:

- RT when $\text{colmodctr} > \text{SQRT}(\text{table cardinality} * 1000)$

n = table rows (cardinality) or colmodctr of the leading column of the statistics object

Query Processing

Optimizations



What QP Searches and Considers When Optimizing?

- Join reordering
- Outer joins
- Sub-queries
- Aggregation
- Stars and snowflakes
- Join elimination
- Materialized views
- Index plans
- Update plans

- Halloween protection
- Empty tab simplification (Integrity constraints)
- Partitioned tables
- Parallelism
- Remote queries
- Correlation elimination
- Sub-query elimination


Join Reordering

- SQL Server join-paths between tables may differ from the actual written Transact-SQL form, known as **join reordering**
- The goal is to reduce the row-count of each join as early as possible

Join Reordering (Continued)

- JOIN ordering uses:
 - Heuristics
 - Statistics
 - Indexing
- Queries can be represented as trees, in accordance with the JOINed tables:
 - Linear trees

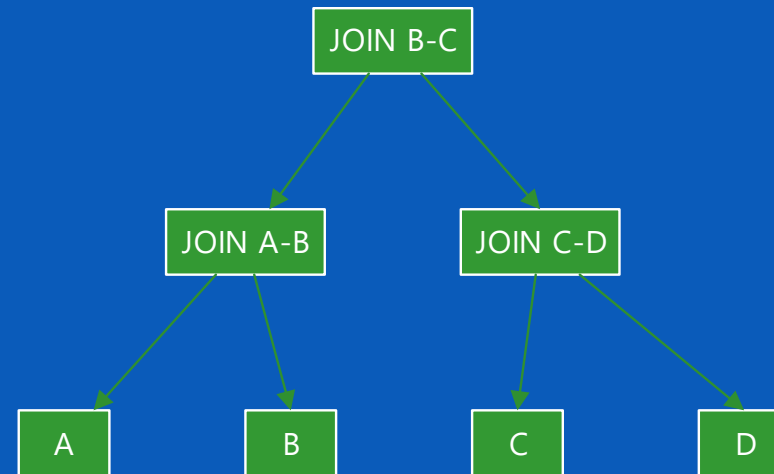
```
SELECT *  
FROM tA JOIN tB ON tA.c1 = tB.c1  
JOIN tC ON tB.c1 = tC.c1  
JOIN tD ON tC.c1 = tD.c1
```



Join Reordering (Continued)


- JOIN ordering uses:
 - Heuristics
 - Statistics
 - Indexing
- Queries can be represented as trees, in accordance with the JOINed tables:
 - Linear trees
 - Bushy trees

```
SELECT *  
FROM (tA JOIN tB ON tA.c1 = tB.c1)  
JOIN (tC JOIN tD ON tC.c1 = tD.c1)  
ON tB.c1 = tC.c1
```




Join Reordering (Continued)

```
SELECT *  
FROM tA JOIN tB ON tA.c1 = tB.c1  
JOIN tC ON tB.c1 = tC.c1  
JOIN tD ON tC.c1 = tD.c1
```




```
SELECT *  
FROM (tA JOIN tB ON tA.c1 = tB.c1)  
JOIN (tC JOIN tD ON tC.c1 = tD.c1)  
ON tB.c1 = tC.c1
```



Join Reordering – Permutations


```
SELECT *  
FROM tA JOIN tB ON tA.c1 = tB.c1  
JOIN tC ON tB.c1 = tC.c1  
JOIN tD ON tC.c1 = tD.c1
```



Nr. Tables	Linear Tree
1	1
2	2
3	6
4	24
5	120
6	720
7	5,040
8	40,320

Join Reordering – Permutations (Continued)

```
SELECT *  
FROM (tA JOIN tB ON tA.c1 = tB.c1)  
JOIN (tC JOIN tD ON tC.c1 = tD.c1)  
ON tB.c1 = tC.c1
```



Nr. Tables	Bushy Tree
1	1
2	2
3	12
4	120
5	1,680
6	30,240
7	665,280
8	17,297,280

Query Parallelism

- Used by SQL Server to reduce the run-time of a query
- CPU cost is generally higher than a serial plan
- Queries are parallelized by horizontally partitioning data and assigning a thread to each partition
- The degree of parallelism (DOP) is determined at the time of execution based on resource availability and Resource Governor settings