

# Buffer Pool Challenge

## OLTP workload characteristics

- High volumes which access to small amount of data
- Generates many random IOPS to HDD

## Buffer Pool Challenge

- Long wait times for PAGE\_IO\_LATCH

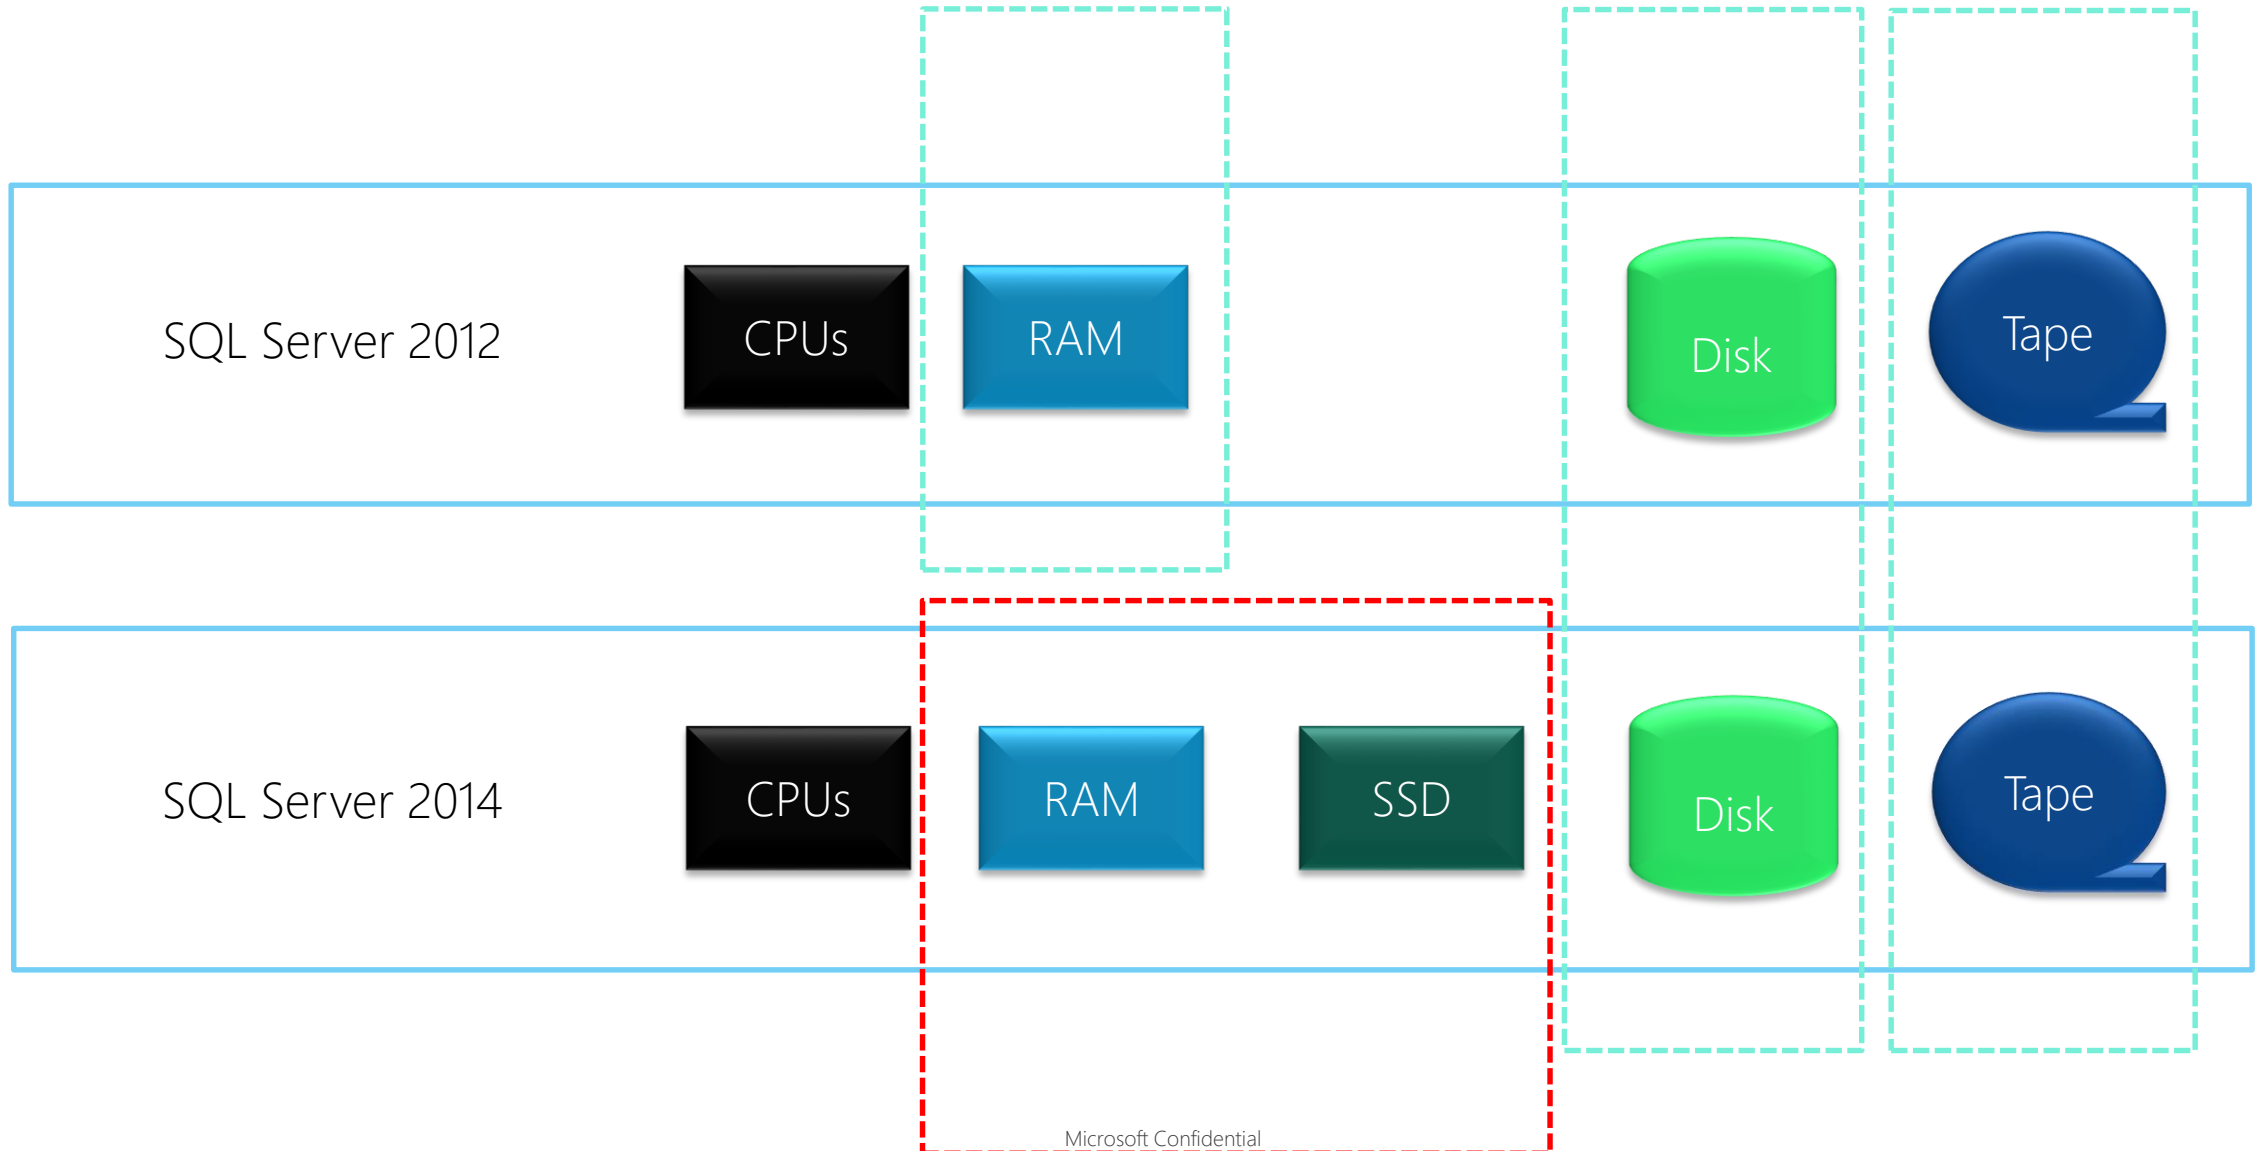
# Why Buffer Pool Extensions?

OLTP workloads  
can have IO  
bottlenecks due  
to Random Disk  
access

Need an  
affordable way to  
increase OLTP  
workload  
performance

Adding memory  
can be prohibitive  
due to hardware  
limitations & cost

# New Layer For Data Placement



# Terminology

L1 cache: Buffer Pool on RAM

L2 cache: Buffer Pool Extension on SSD

Buffer: An 8KB page in memory



# Benefits

Use of non-volatile drives (SSD) to extend buffer pool

- Reduced IO Latency
- Increased Random IO throughput.

NUMA-Aware large page and BUF array allocation

Improve OLTP query performance with no application changes

No risk of data loss (using clean pages only)

Optimized for OLTP workloads on commodity servers

# Syntax

New configuration option - BUFFER POOL EXTENSION

Need to specify the size equal or greater than BUFFER POOL

ALTER SERVER CONFIGURATION

SET **BUFFER POOL EXTENSION** ON (FILENAME = 'os\_file\_path\_and\_name', SIZE = {SIZE KB/MB/GB}) [;]



New option  
- BUFFER POOL EXTENSION



Required:  
Path and file name



Required:  
Size of BPE

# Commands

- Check if enabled

```
SELECT * FROM sys.dm_os_buffer_pool_extension_configuration
```

- Disable

```
ALTER SERVER CONFIGURATION  
SET BUFFER POOL EXTENSION OFF
```

- Which pages are in Buffer Pool Extensions?

```
SELECT * FROM sys.dm_os_buffer_descriptors  
WHERE is_in_bpool_extension = 1
```

- Need to turn configuration OFF and ON to change value

# Extension File and Device

## No risk of data loss in extension file

- It only contains clean buffers

## If SSD device is inaccessible

- SSD Buffer Pool Extension is disabled
- Need to re-enable

## On a Cluster

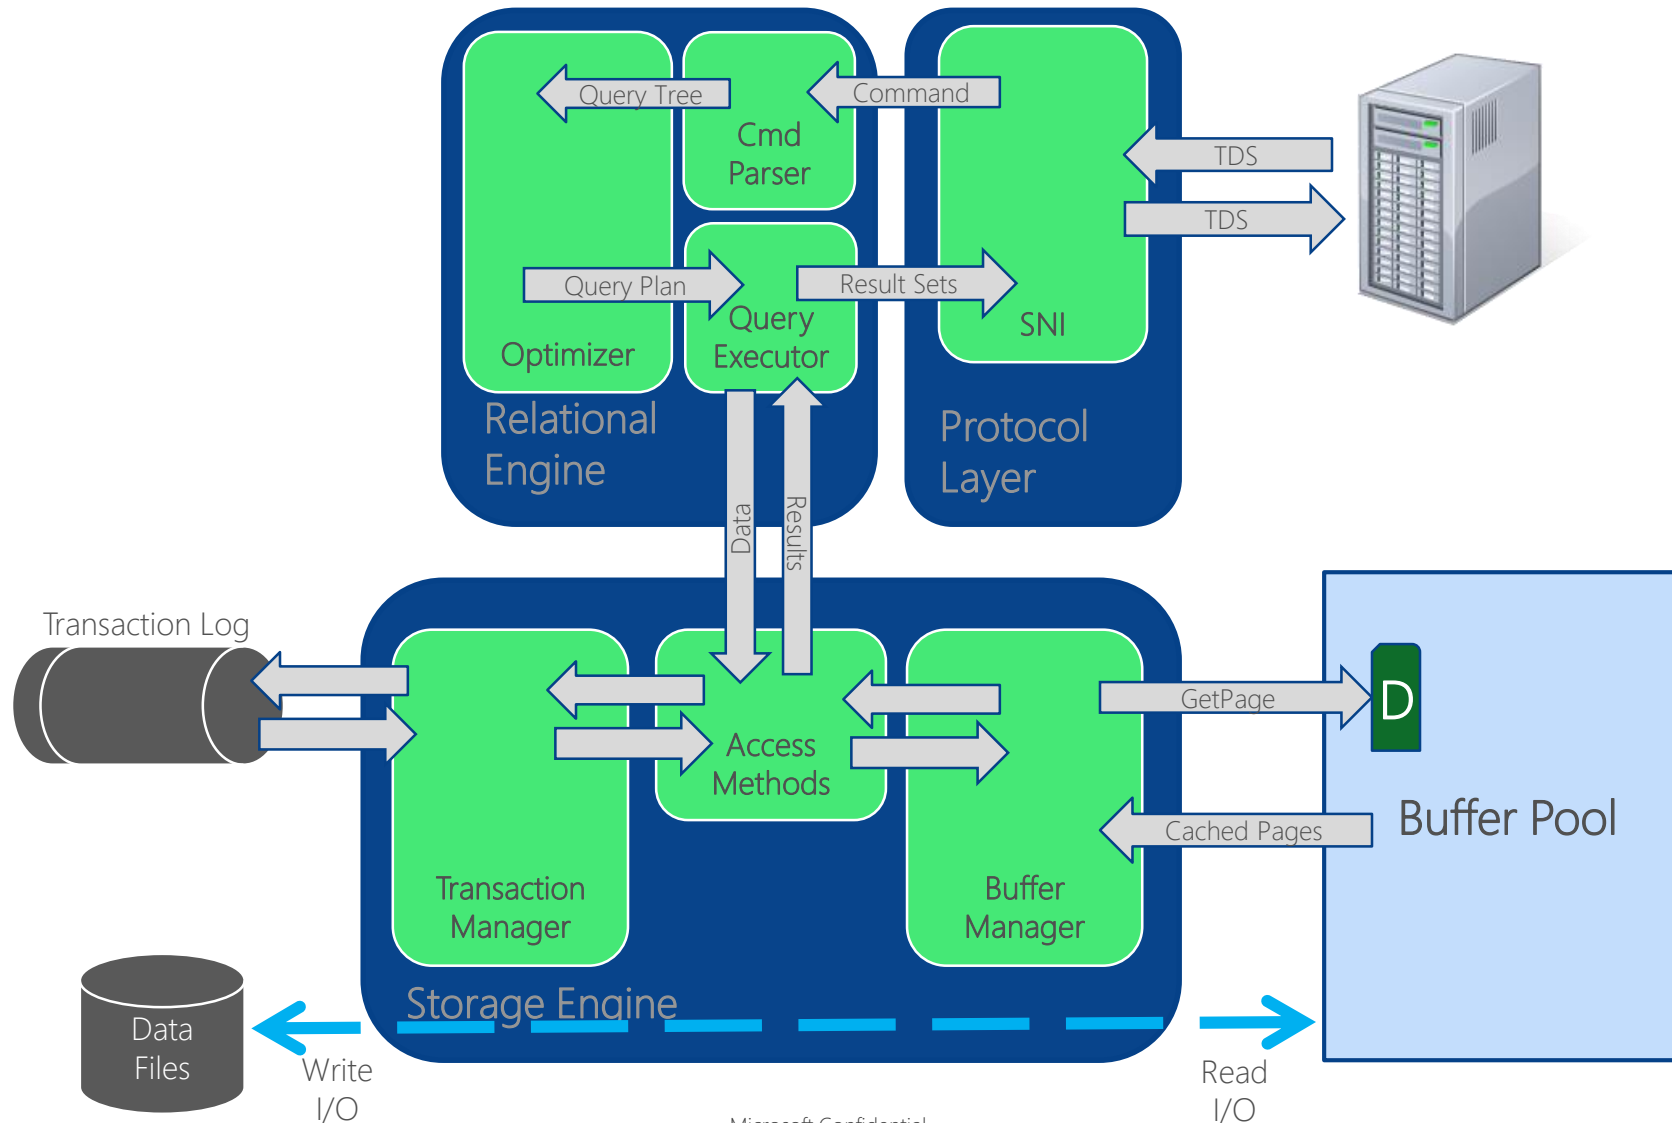
- Ensure it is sized equally on both nodes




## Limitations

- Memory Optimized Tables
- ColumnStore Index


# Buffer Pool Manager



# Buffer Pool Manager – With Extensions



# SSD Buffer Pool Extension Internals

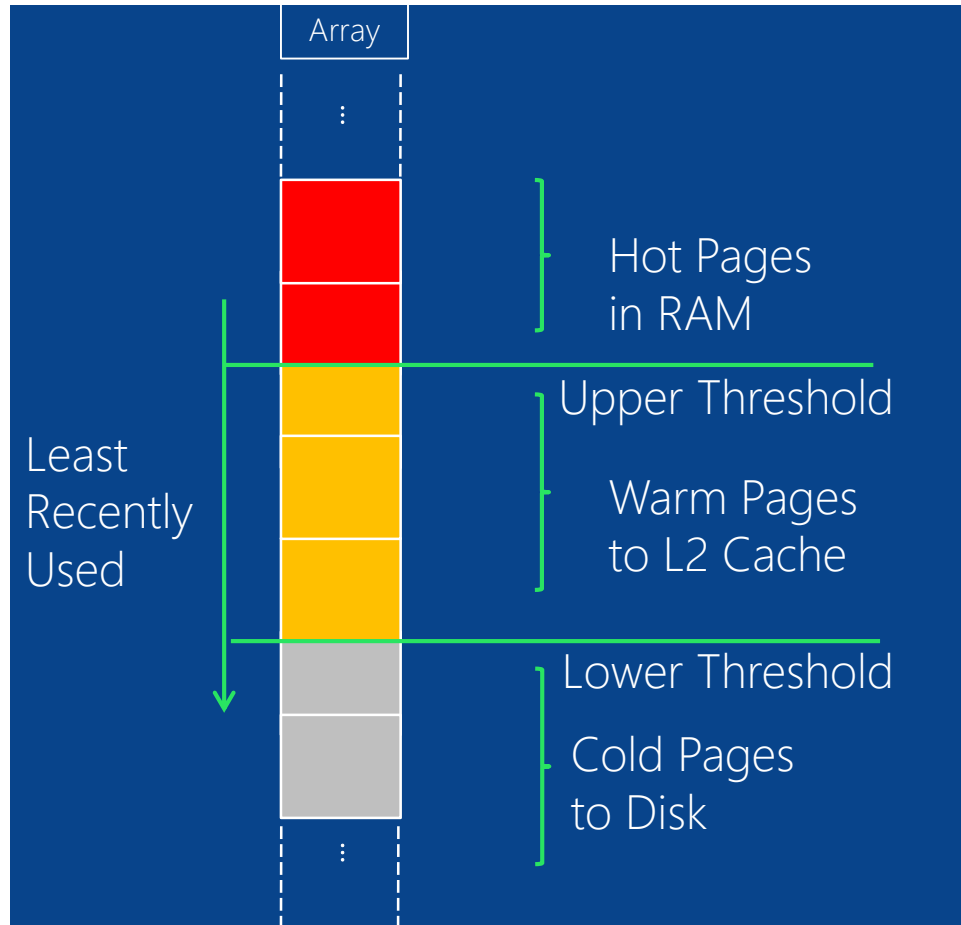


Managing the pages on the SSD buffer pool extension

It is managed as a 8 KB page

Using BUF Array

# Data Rotation Algorithm



Data placement corresponding to usage frequency

- LRU algorithm

Upper Threshold > Reference counts

- Buffer Pool on RAM (L1 cache)

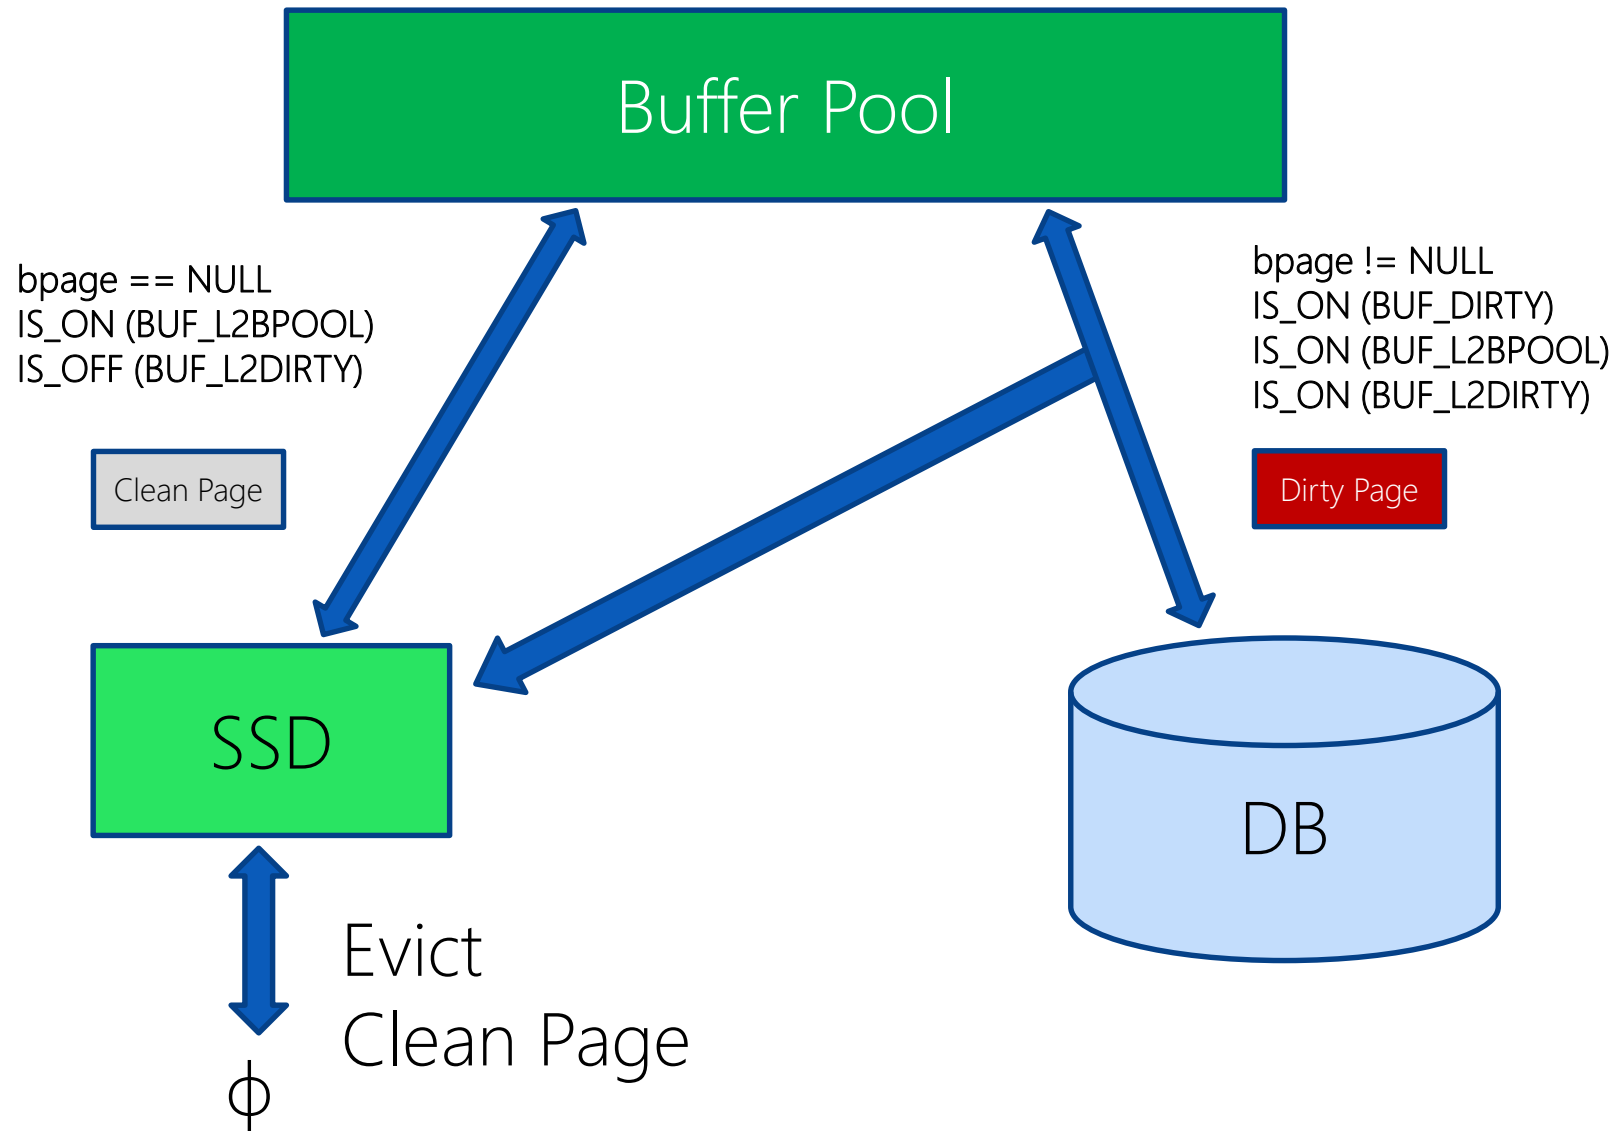
Upper Threshold > Reference counts > Lower Threshold

- Buffer Pool on SSD (L2 cache)

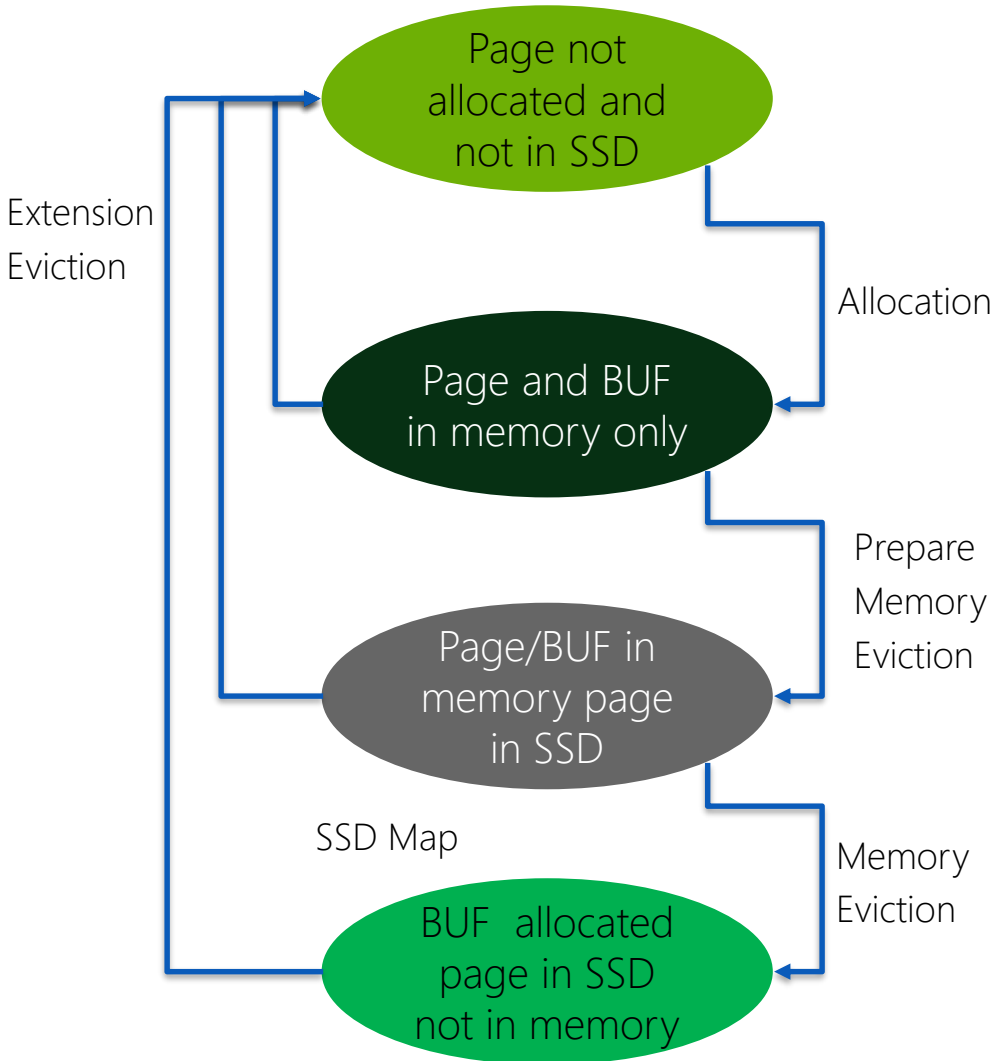
Lower Threshold < Reference counts

- On mechanical disk

# Implementation



# Page Eviction



## Memory Eviction:

- Migrates a page from virtual memory to the L2 Cache

## Extension Eviction:

- Writes dirty pages out to spinning media, frees page from RAM and SSD

## Performance Counters

### SQL Server Buffer Manager

- Extension page writes/sec
- Extension page reads/sec
- Extension outstanding IO counter
- Extension page evictions/sec
- Extension allocated pages
- Extension free pages
- Extension page unreferenced time
- Extension in use as percentage



# xEvents and DMVs

## xEvents:

- `buffer_pool_extension_pages_written`
- `buffer_pool_extension_pages_read`
- `buffer_pool_extension_pages_evicted`
- `buffer_pool_page_threshold_recalculated`

## DMV:

- `sys.dm_os_buffer_descriptions`
- `sys.dm_os_buffer_pool_extension_configuration`

# Demonstration: Buffer Pool Extensions

# Best Practices

## Enough on-board memory

- No benefit from Buffer Pool Extension

## Target Workload

- OLTP - Read heavy preferred
- A large workload
- Multiple smaller workloads (consolidation)
- Systems that have memory pressure
- Sweet spot: SSD storage sized 4x-10x times of RAM size

# Managed Lock Priority - What is it trying to solve?

## Blocking by online DDL

- Partition switching - short duration Sch-M lock acquired
- Online index rebuilds - short table S and Sch-M lock acquired

## Concurrency and availability

- Blocking transactions need to complete before DDL
- Switch will block new transactions
- Timeouts, workload slowdowns
- Impact to mission critical OLTP workloads

# Without Managed Lock Priority



# Syntax

New clause in existing T-SQL DDL for ALTER TABLE and ALTER INDEX

## Syntax

**<low\_priority\_lock\_wait>::=**

```
{  
    WAIT_AT_LOW_PRIORITY ( MAX_DURATION =  
<time>[MINUTES],  
  
    ABORT_AFTER_WAIT = { NONE | SELF | BLOCKERS } )  
}
```

NONE - current behavior  
SELF - abort DDL  
BLOCKERS – abort user blockers

```
ALTER TABLE stgtab SWITCH PARTITION 1 TO parttab  
PARTITION 1
```

```
    WITH (WAIT_AT_LOW_PRIORITY (MAX_DURATION=  
60 minutes,  
    ABORT_AFTER_WAIT=BLOCKERS))
```

```
ALTER INDEX clidx ON parttable REBUILD  
WITH ( ONLINE=ON (WAIT_AT_LOW_PRIORITY  
(MAX_DURATION= 300,  
    ABORT_AFTER_WAIT=SELF)) )
```

Examples

# ABORT\_AFTER\_WAIT Summary

## Kill All Blockers (BLOCKERS)

- Blocking user transactions killed
- MAX\_DURATION in minutes

Requires ALTER ANY CONNECTION permission to kill sessions

## Switch to Normal Queue (None)

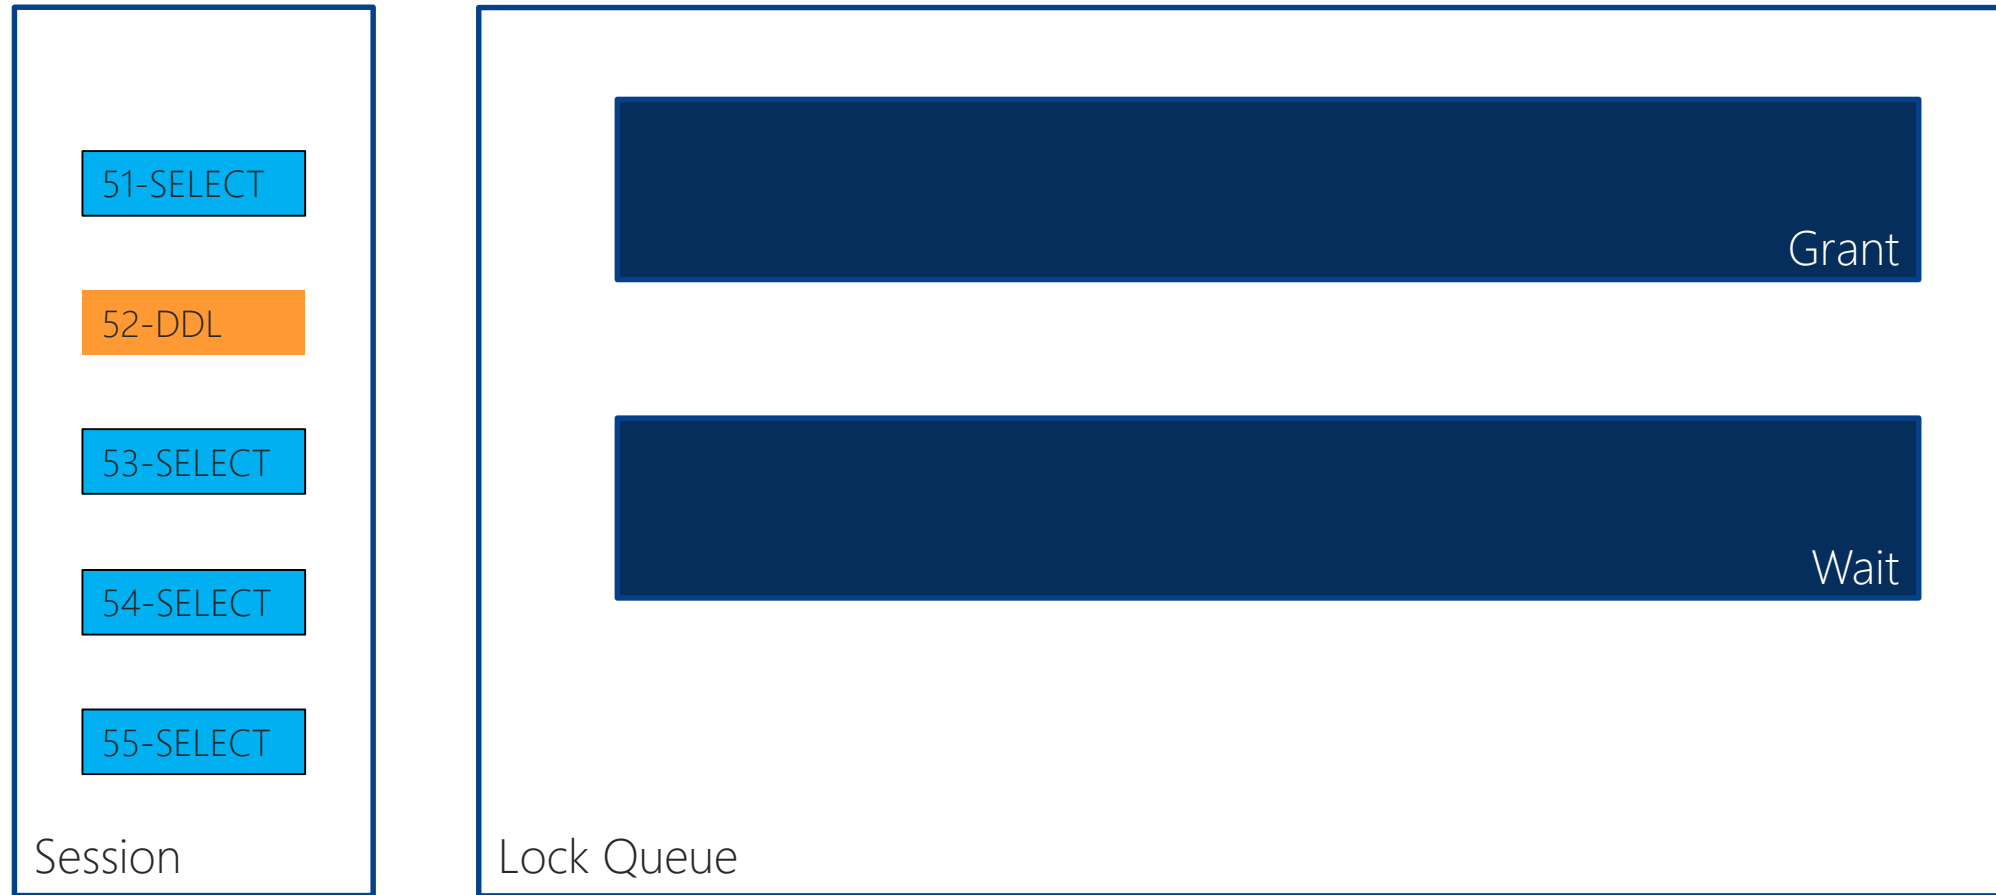
- Wait for blockers
- MAX\_DURATION in minutes
- Regular FIFO lock queue

## Exit DDL after wait (SELF)

- Wait for blockers
- MAX\_DURATION in minutes
- Terminates DDL after wait time

If no blockers, lock granted immediately, DDL completes

# Behavior with Managed Lock Priority





# Managed Lock Priority Summary

## Benefits

- Managed by DBA for both partition switch and Online Index Rebuild
- Lock request placed in lower priority queue
- Decision to wait or kill self or blockers
- Executed immediately if no blockers

## Partition Switch

- Blockers killed at source and destination
- Sch-M lock (source and destination)

## Online Index Rebuild

- MAX\_DURATION applies to every lock request
- Time reset for every S & Sch-M lock
- Only Sch-M lock conflict for read only workloads

# Best Practices

## ABORT\_AFTER\_WAIT=BLOCKERS

- Long running queries/transactions
- May not be beneficial in workloads with short transactions

## Locking

- SWITCH acquires locks on source and target tables
- Wait on the second lock affects concurrency on first
- Dormant tables

## DDL Transaction

- Nested transaction holds locks longer

## Transaction Log Growth

- Log truncation affected by MAX\_DURATION
- Extra logging Online Index Rebuild DDL operation
- Can lead to system performance issues

# Single Partition Online Index Rebuild

## Accessibility

- One or more partitions can be rebuilt
- Table accessible for DML and query operations
- Short term locks beginning and end of the index rebuild

## Lock Priority

- Use Managed Lock Priority with Single Partition Online Index Rebuild

## Availability

- Reduced down time for mission critical workloads

## Resource savings

- CPU, memory and disk space
- Log space usage reduced

# Managed Lock Priority Diagnostics

## DMVs

- sys.dm\_tran\_locks "request\_status" → LOW\_PRIORITY\_CONVERT, LOW\_PRIORITY\_WAIT, or ABORT\_BLOCKERS
- sys.dm\_os\_wait\_stats "wait\_type" → ...LOW\_PRIORITY and ..ABORT\_BLOCKERS

## Extended Events

- lock\_request\_priority\_state
- process\_killed\_by\_abort\_blockers
- ddl\_with\_wait\_at\_low\_priority
- sqlserver.progress\_report\_online\_index\_operation

## Error Log

- Abort session diagnostics in SQL Server Error Log
- Deadlock diagnostics in deadlock graph

# SQL Server & Statistics

## What are Statistics?

- Distribution of values
- One or more columns

## Why are they important?

- Plan choices
- Query performance

## How are they maintained?

- Automatic update triggered by at least 20% of the data change
- For Larger Tables, TF 2371 ([Auto Update Stats Changes](http://blogs.msdn.com/b/saponsqlserver/archive/2011/09/07/changes-to-automatic-update-statistics-in-sql-server-traceflag-2371.aspx) - <http://blogs.msdn.com/b/saponsqlserver/archive/2011/09/07/changes-to-automatic-update-statistics-in-sql-server-traceflag-2371.aspx>)

# Challenges Prior to SQL Server 2014

- Statistics become stale
- SCANs or samples the entire table
- Ascending keys
- One stats blob for whole table

# Incremental Statistics

KB 2970386 -  
Metadata  
Corruption Fix

A feature targeted for partitioned tables

- Assuming most updates in newest partitions

## CREATE STATISTICS

- One statistical object for each partition
- Binary merge of partial statistics
- A complete tree of statistics persisted

## UPDATE STATISTICS

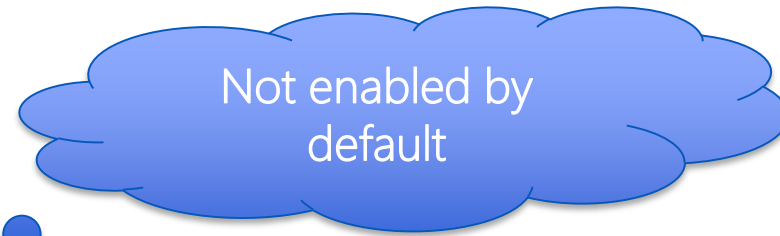
- Rebuild statistics on a subset of partitions
- Merge with rest of the statistics tree

# Enabling Incremental Statistics

```
CREATE STATISTICS stat ON tbl (col1, col2)  
WITH INCREMENTAL = ON
```

```
UPDATE STATISTICS tbl (stat)  
WITH RESAMPLE ON PARTITIONS (1, 3, 5)
```

```
UPDATE STATISTICS tbl (stat)  
WITH INCREMENTAL = { ON | OFF }
```




```
ALTER DATABASE dbName  
SET AUTO_CREATE_STATISTICS ON (INCREMENTAL = { ON | OFF })
```

```
CREATE INDEX myind ON tbl(col)  
WITH (STATISTICS_INCREMENTAL = ON)
```




# Incremental Stats Example


Create Incremental Statistics on a 4 partition table



Add partition 5 and update statistics



Merge partitions 4 and 5, and update statistics



# Query compilation

Still uses global table stats for query optimization

- No per-partition statistics for optimization yet
- Consider Filtered Statistics

Threshold for auto stats update and query recompile

- To update global : 500 + 20% of **average** partition size (non-empty)
- To update per-partition : 20% data change in a partition

Default sample rate for incremental stats

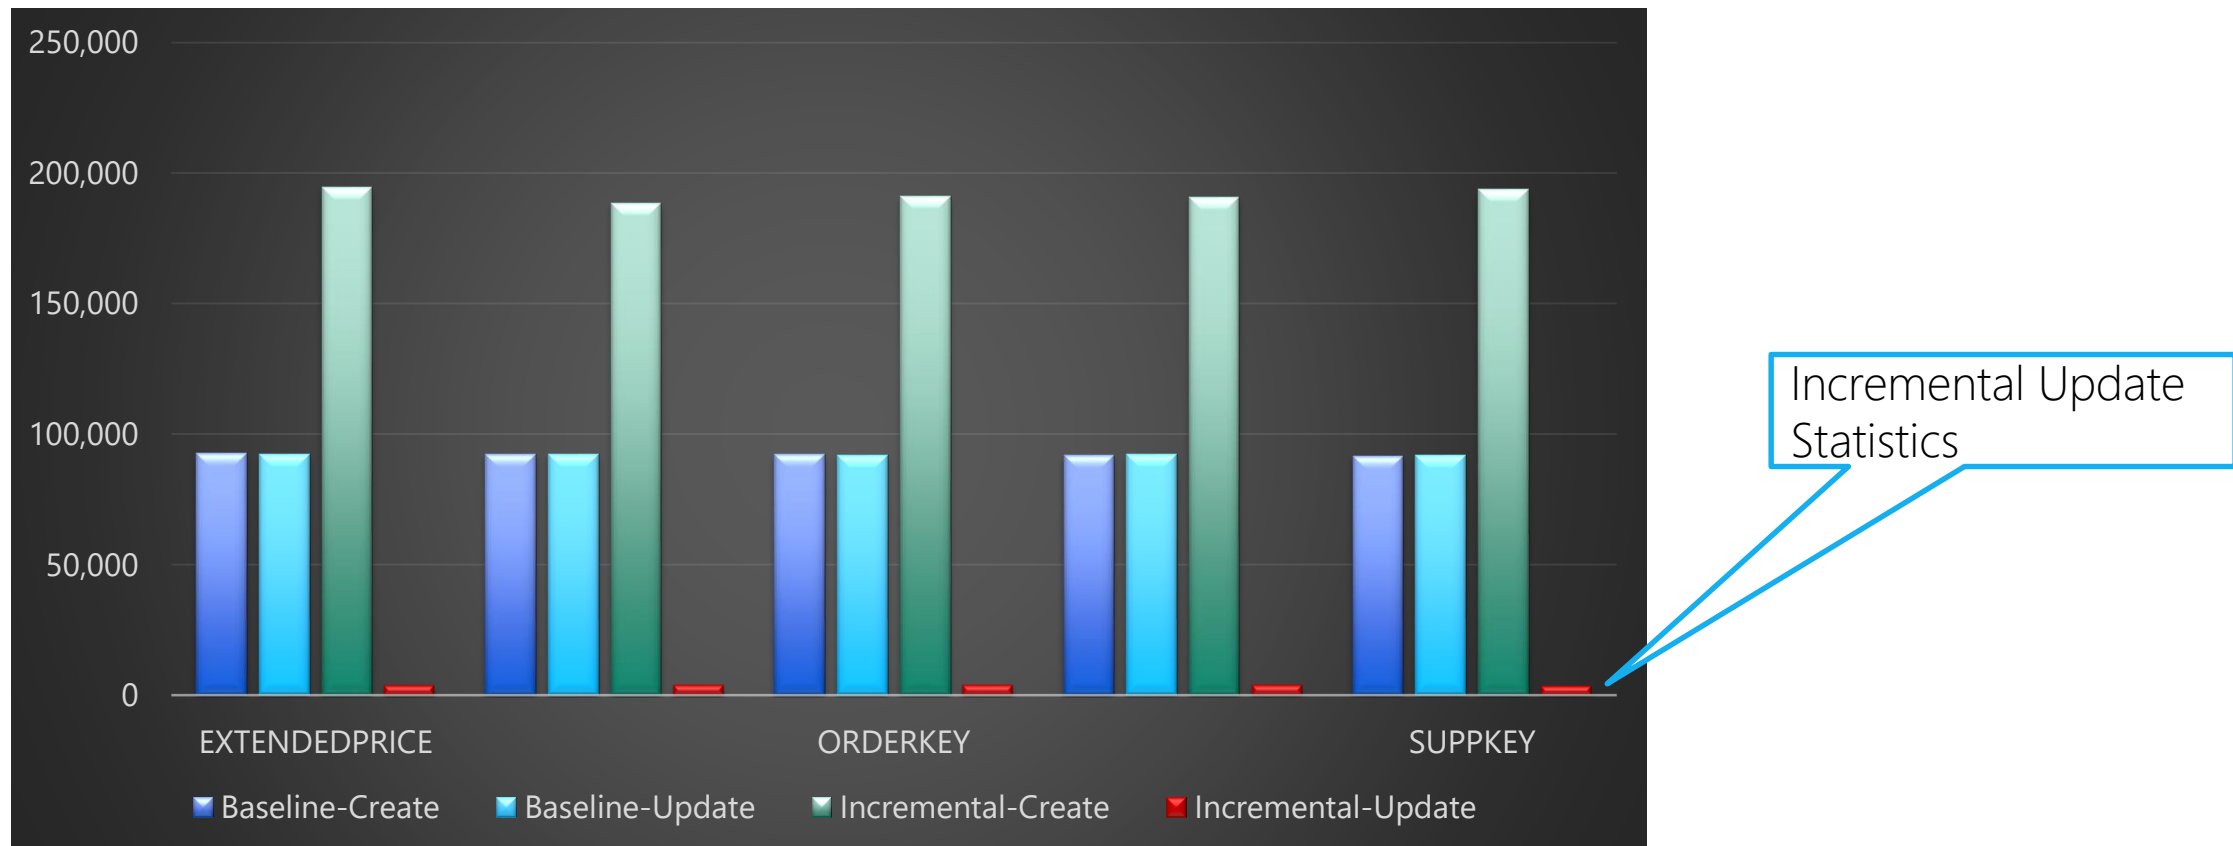
- Auto create : same default sample rate as regular stats
- Auto update : using most recent sample rate

# Performance Results- Sample

Database **TPCH 100G**

Table **LINEITEM** **600 million** rows in **2,500** partitions

Measure elapsed time for create/update statistics with **sample size = 1 million rows**

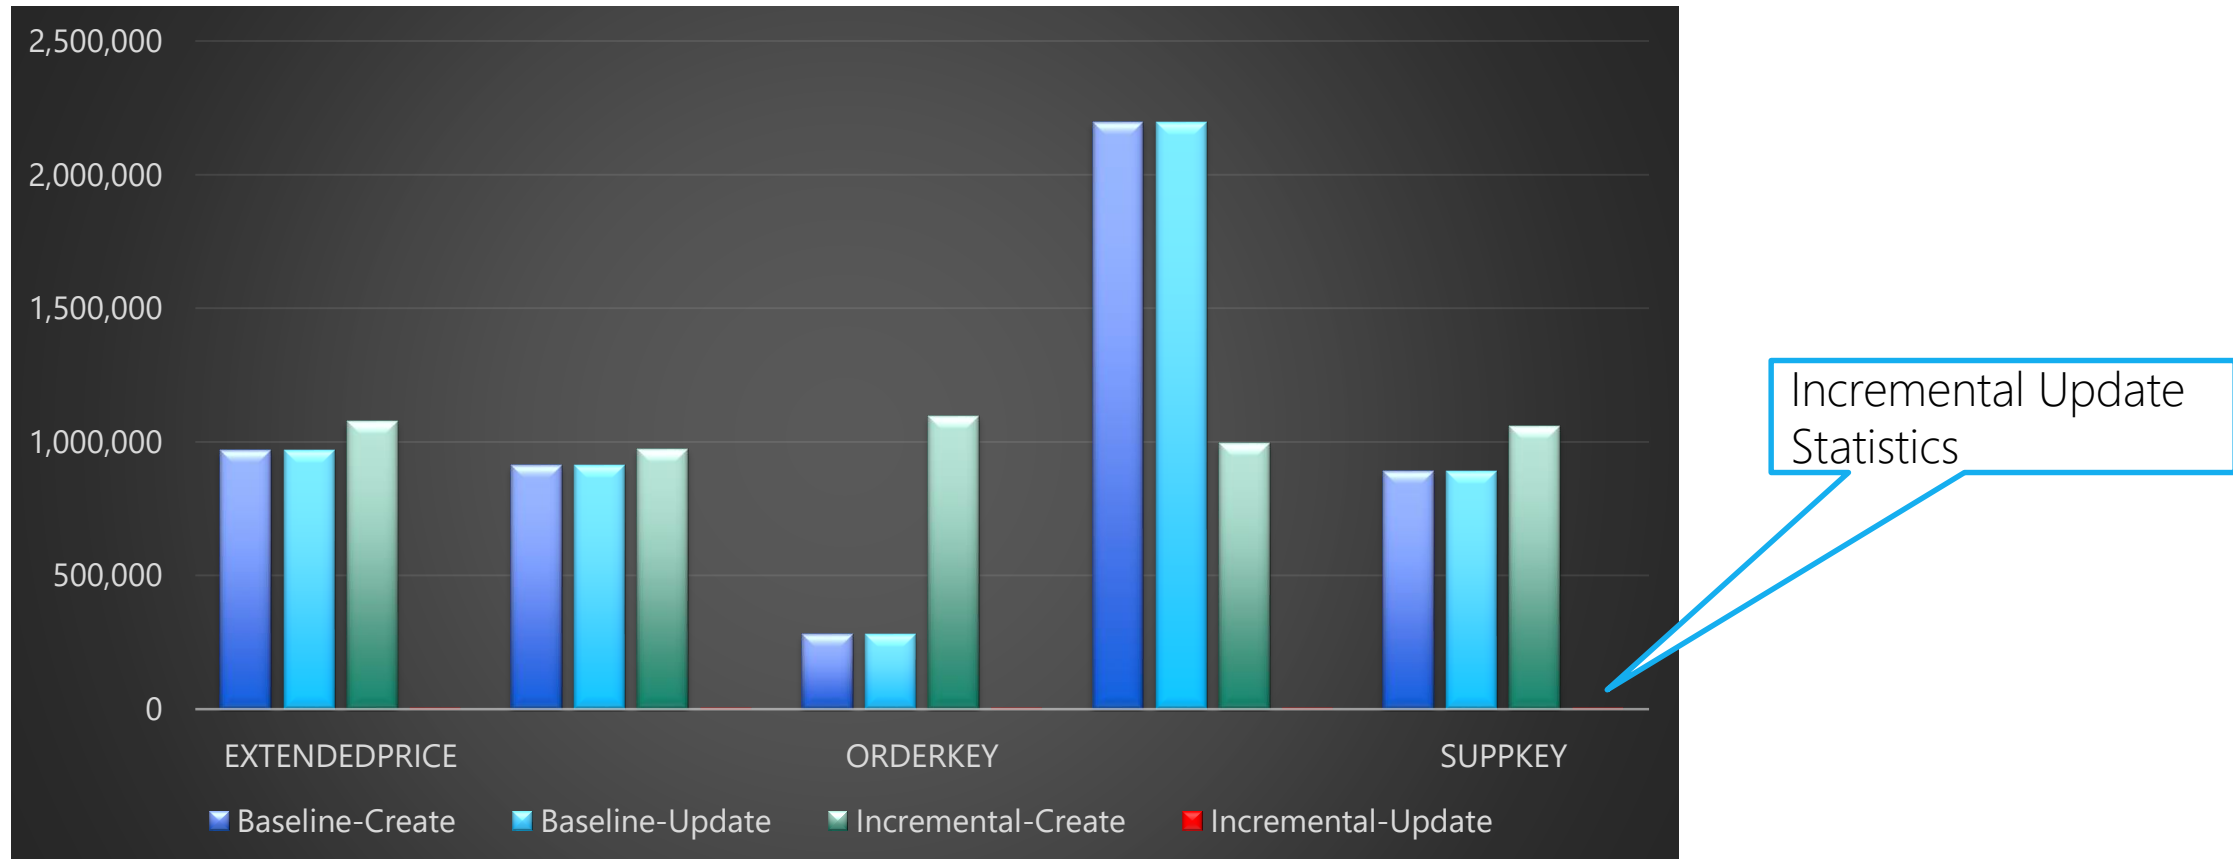


# Performance Results - FULLSCAN

Database **TPCH 100G**

Table **LINEITEM** **600 million** rows in **2,500** partitions

Measure elapsed time for create/update **FULLSCAN** statistics



# Limitations

Indexes that are not partition-aligned with the base table

AlwaysOn readable secondary databases

Read-only databases

Filtered indexes

Views

Internal tables

Spatial indexes or XML indexes

# Incremental Statistics Diagnostics

## sys.stats

- sys.stats has a column *is\_incremental* to indicate whether a statistic is incremental

## sys.databases

- sys.databases has a column *is\_auto\_create\_statistics\_incremental* to indicate whether auto created stats will be incremental

## sys.dm\_db\_stats\_properties

- Accurately track the number of rows changed in a table

# What is Cardinality Estimation?

Estimate of number of rows returned by each operation

- Based on statistics metadata
- Fixed estimates (assumptions) when no histograms are available

Query optimization is inherently unpredictable and very sensitive to cardinality estimations

Cardinality estimation influences the cost, which is an indicator of response time

Poor estimations can result in inefficient plans

# Why Change Cardinality Estimation?

## Business Case

- CE engine was the same since 7.0
- Many changes under trace flag ([Trace flag 4199 for Optimizer](#))
- Goal was to architecturally enhance the quality of cardinality estimation for a broad range of queries and workloads: OLTP, DW and DSS
- Create better plans for most cases, especially complex queries
- Make performance smoother and more predictable

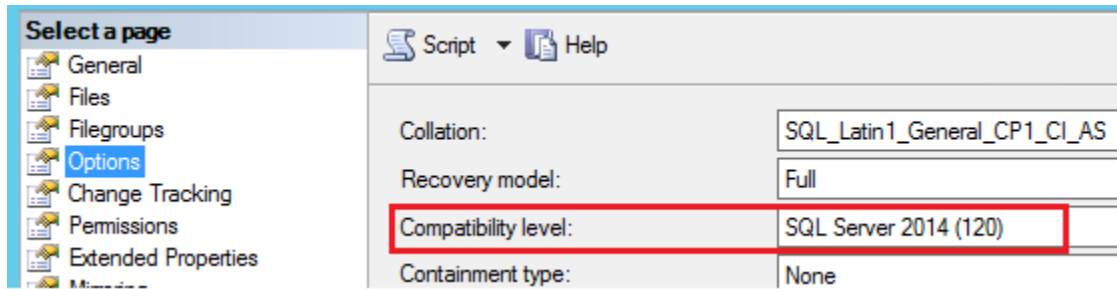
## Ramifications

- Can result in worse plans in edge cases
- Can be disabled, new CE is side-by-side



# New Cardinality Estimation

## New CE model

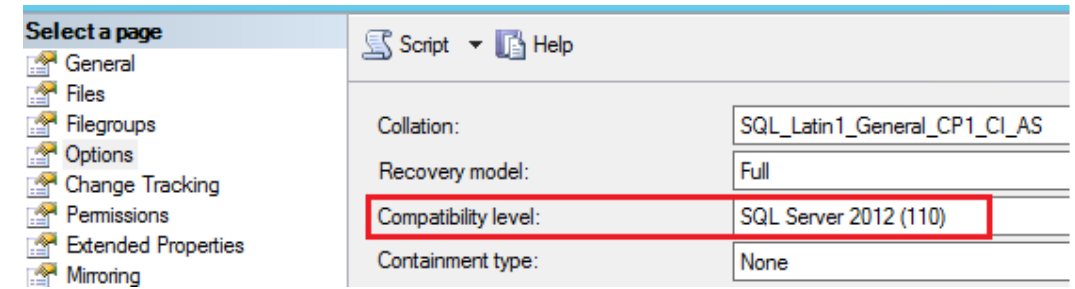


Select a page

Script Help

Collation:	SQL_Latin1_General_CP1_CI_AS
Recovery model:	Full
Compatibility level:	SQL Server 2014 (120)
Containment type:	None

## Old CE model



Select a page

Script Help

Collation:	SQL_Latin1_General_CP1_CI_AS
Recovery model:	Full
Compatibility level:	SQL Server 2012 (110)
Containment type:	None

# New Cardinality estimation – Trace Flags

New CE model

```
SELECT * FROM FactCurrencyRate  
WHERE DateKey = 20101201  
OPTION (QUERYTRACEON 2312)
```



or Server Level TF 2312

Old CE model

```
SELECT * FROM FactCurrencyRate  
WHERE DateKey = 20101201  
OPTION (QUERYTRACEON 9481)
```



or Server Level TF 9481

# Precedence

Query Level Trace Flag



Server or Session Trace Flag



Database Compatibility Mode

# Ascending Key Problem

## Example

- `SELECT item, category, amount FROM dbo.Sales AS s WHERE Date = '2013-12-19'`
- Stats last updated 18<sup>th</sup> December, 2012

## Old CE

- Latest data inserted not in the histogram
- Cardinality estimation comes up as 1
- Trace flag 2389 can help ([Ascending Keys and Auto Quick Corrected Statistics](#))

## New CE

- Assumption is that the value exists
- Estimate cardinality using average frequency
- Trace flag 2389 no longer relevant

# What's New - Exponential Back Off

Change in calculating combined density with single-table predicates

1000 Rows. 200 "Microsoft" , 50 XRacers.

```
SELECT * FROM Cars
WHERE Manufacturer = 'Microsoft' AND Model = 'XRacer'
```

Old CE (Independence)	New CE (Exponential Backoff)
$p_0 \times p_1 \times p_2 \times p_3 \times \dots$	$p_0 \times p_1^{1/2} \times p_2^{1/4} \times p_3^{1/8}$
$0.05 * 0.2 * 1000 = 10$	$0.05 * \text{sqrt}(0.2) * 1000 = 22.36$


Density = 1/# Distinct Values

# What's New – Join Changes

```
SELECT t1.c, t1.d, t2.c  
FROM t1 INNER JOIN t2  
ON t1.a = t2.a AND t1.b = t2.b;
```

Equijoins with multiple predicates

- Take the lesser of the *Distinct Value Counts* between two inputs (10)
- Multiply by *Average Frequencies* from both inputs ( $10 * 1 * 14.4 = 144$ )



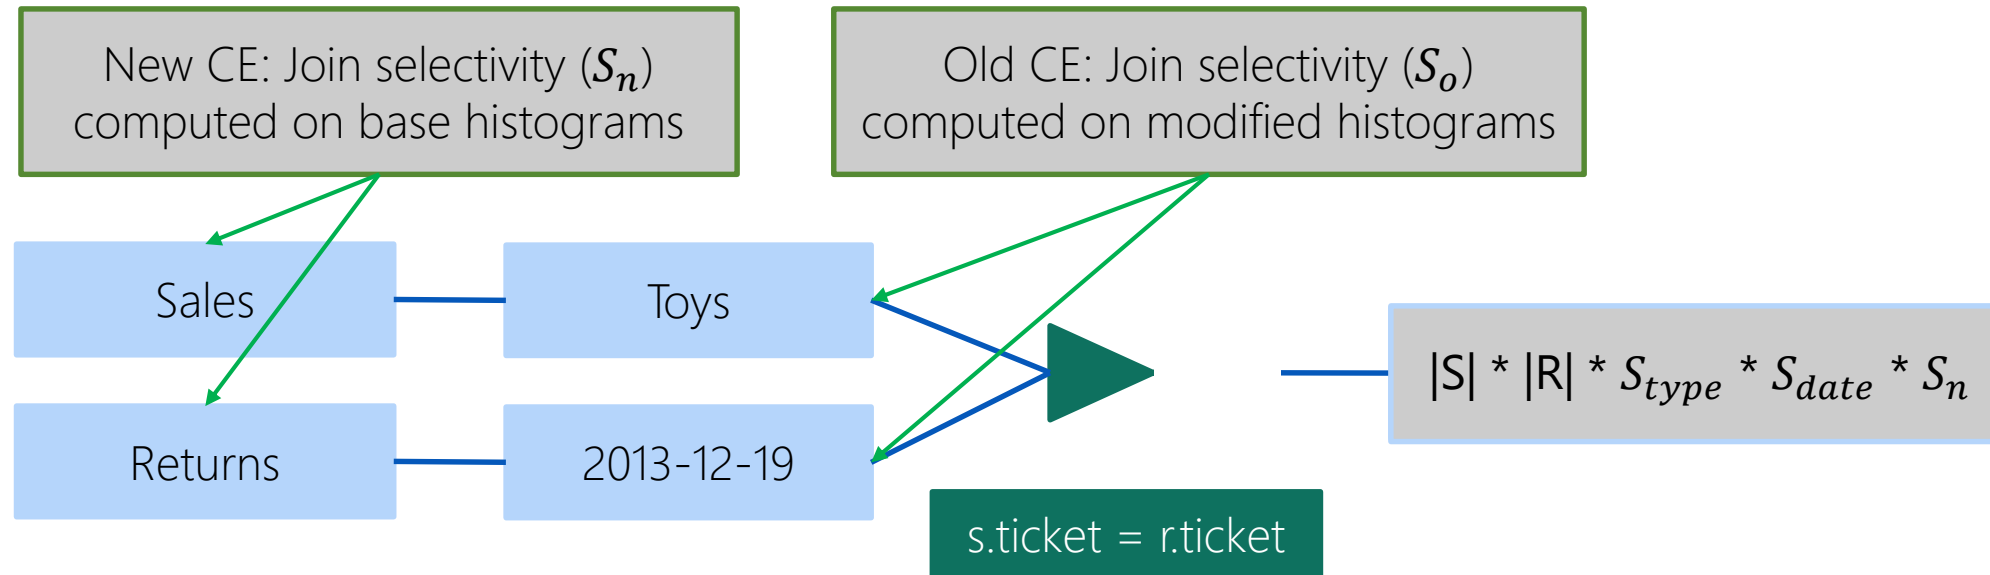
Joins use Base Containment in new CE v/s Simple containment in old CE

- Selectivity applied to Base tables
- Then selectivity of additional predicates applied


# What's New – Containment Assumption

```
SELECT s.ticket, s.customer, r.store
FROM dbo.Sales AS s
INNER JOIN dbo>Returns AS r ON s.ticket = r.ticket
WHERE s.type = 'toy' AND r.date = '2013-12-19';
```

Equijoins with multiple predicates



# Impact of using New CE





# When to use New CE?

Customers with query predictability issues

Data warehouse workloads can hugely benefit per testing

Test your workload with the new CE, there are edge cases!

Supportability features to help analyze Query Plans and Optimizer choices

- Extended Event *query\_optimizer\_estimate\_cardinality*
- Extended Event *query\_optimizer\_force\_both\_cardinality\_estimation\_behaviors*

# Diagnostics

## Extended Events

- query\_optimizer\_estimate\_cardinality
- query\_optimizer\_force\_both\_cardinality\_estimation\_behaviors

## Interpreting Extended Event

- Stats collection
- Input\_relation
- CalculatorName

## XML Query Plans

- StatsID