

# SQLintersection

## Eliminating Low Hanging Fruit

Jonathan Kehayias  
Principal Consultant  
[SQLskills.com](https://SQLskills.com)

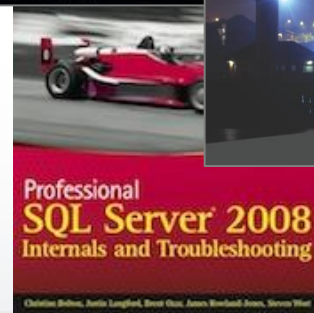

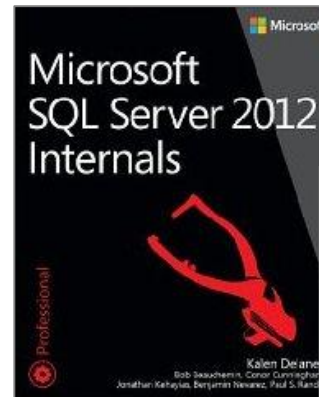


SQL  
*intersection*

# Jonathan Kehayias



- Consultant/Trainer/Speaker
- Principal Consultant, [SQLskills.com](http://www.SQLskills.com)
  - Email: [Jonathan@SQLskills.com](mailto:Jonathan@SQLskills.com)
  - Blog: <http://www.SQLskills.com/blogs/jonathan>
  - Twitter: @SQLPoolBoy
- SQL Server MVP since October 2008
- Microsoft Certified Master: SQL Server 2008



# Overview

- **Row-By-Agonizing-Row processing**
- **Develop and test against realistic scale**
- **Profile your workload during development**
- **Use appropriate data types**
- **Sargability matters**

# Row-By-Agonizing-Row Processing

- **Certain constructs force SQL Server into RBAR (row-by-agonizing-row) processing of results**
- **Well known:**
  - WHILE loops
  - Cursors
- **Less well known:**
  - Scalar user defined functions
  - Correlated subqueries

# Cursors and Loops

- **Characteristics**

- Explicit cursor declaration
- WHILE loop
- SqlDataReader in application

- **Problems**

- Row based processing over Set Based

- **Replace with**

- Appropriate set based operation
- Move looping code into SQLCLR or Middle Tier
- Consume and dispose of SqlDataReader as quickly as possible

# Correlated Sub-Queries

- **Characteristics**

- Refer to the outer query in the inner query in SELECT statement
- SELECT Statement used as column value in UPDATE

- **Problems**

- May cause row-by-row processing to occur
- Performance decreases exponentially as row count increases

- **Replace with**

- Table Join
- Derived Table Join
- Cross Joined Table Valued Function

# Scalar User Defined Functions

- **Characteristics**

- Encapsulate common code blocks/business logic in a single call.
- If columns are passed as parameters it is not inline

- **Problems**

- Cause row-by-row processing to occur
- Performance decreases exponentially with data access

- **Replace with**

- Inline expressions
- Derived Table Join
- Cross Joined Inline Table Valued Function

# Develop and Test Against Realistic Scale

- **Development databases often do not contain realistic datasets which can hide/mask potential performance problems**
  - Key Lookups on small data sets may become index scans on larger data sets
  - Missing index impacts may be hidden by data residing in memory for small data sets
  - RBAR problems are often hidden until data sizes scale up
- **Testing a single execution in isolation is not load testing**
  - Testing needs to be performed at scale through load generation to measure accumulated effects
  - Only testing at scale can identify “*death by 1000 cuts*” problems



# Profile Your Workload During Development

- Learn to use Extended Events (2012+) or SQL Trace to profile your workload during testing
- Know the important events to watch for during development
  - Statement/Batch/RPC completed events
  - SP completed/Module End events – (procedure/trigger/function executions)
  - Execution warnings (sort, hash, missing join predicate)
- Profiling during development can uncover nasty RBAR issues and performance effecting side effects of trigger executions
- Be aware of “*observer overheads*” but not typically a problem with development/test workloads

# Use Appropriate Data Types

- Understand the storage costs and implications of data types during schema design – especially for keys

Data type	Range	Storage
<b>bigint</b>	-9,223,372,036,854,775,808 to 9,223,372,036,854,775,807	8 Bytes
<b>int</b>	-2,147,483,648 to 2,147,483,647	4 Bytes
<b>smallint</b>	-32,768 to 32,767	2 Bytes
<b>tinyint</b>	0 to 255	1 Byte

# Use Appropriate Data Types (2)

- **Consider precision requirements for dates and times**
  - DATETIME2 = 6, 7, or 8 bytes with nanosecond precision
    - 6 bytes for precisions less than 3
    - 7 bytes for precisions 3 and 4
    - All other precisions require 8 bytes
  - DATETIME = 8 bytes with fractions of second precision
    - Fraction of seconds rounded to .000, .003, or .007 seconds
  - SMALLDATETIME = 4 bytes with minute precision
  - DATE = 3 bytes
- **Don't store dates or times as CHAR, VARCHAR, NCHAR, or NVARCHAR**

# Sargability Matters

- A query is sargable (Search ARGument ABLE) if an index seek can be used to speed up the execution of the query
- Anti-patterns to sargable expressions include:
  - Functions in the WHERE clause
  - Implicit/Explicit data type conversions on a column
  - Leading wildcard expressions with LIKE '%<SearchTerm>'
  - Catch all queries and search procedures

# Functions on WHERE Clause Columns

- **Characteristics**

- Used to change the data stored to match criteria being checked
- Conversion of data to a different type

- **Problems**

- Causes Table/Index Scan over Seek

- **Replace with**

- Appropriate Table Design to support business needs
- Indexed/Persisted Computed Column
- Indexed View
- Other coding paradigm to

# Implicit/Explicit Column Conversions

- **Characteristics**

- Column data type is of lower precedence than filtering parameter / joining column data type
- Common in LINQ to SQL/EF and other ORMs

- **Problems**

- Causes Table/Index Scan over Seek

- **Replace with**

- Higher precedence column data type
- Matching data type for filtering parameter

# Catch-All Search Queries

- **Characteristics**

- Used to search across multiple columns using parameters
- Not all parameters require input values
- WHERE clause similar to (`@Param1 IS NULL OR Column1 = @Param1`)

- **Problems**

- No optimized execution plan
- Causes Table/Index Scan over Seek

- **Replace with**

- Separate search procedures for different parameters passed
- Parameterized Dynamic SQL

# Review

- **Row-By-Agonizing-Row processing**
- **Develop and test against realistic scale**
- **Profile your workload during development**
- **Use appropriate data types**
- **Sargability matters**



# Questions?

# SQLintersection

Don't forget to complete an online evaluation!

## Eliminating Low Hanging Fruit

Your evaluation helps organizers build better conferences  
and helps speakers improve their sessions.



**SQL**  
*intersection*

**Thank you!**

# Our Spring Show Lineup Speaks for Itself!

May 21-24, 2017 in Orlando, Florida



## ■ 40 SQL Sessions, 3 SQL keynotes, and 7 workshops:

- 2 workshop on Saturday, May 20
  - *PowerShell for the DBA from 0-60 in a Day* with Ben Miller
  - *SQL Server 2014 and 2016 New Features and Capabilities* with Tim Chapman and David Pless
- 2 workshops on Sunday, May 21
  - *Performance Troubleshooting using Waits and Latches* with Paul Randal
  - *Azure for the SQL Server DBA* with Tim Radney
- 3 workshops on Thursday, May 25
  - *Finding and Fixing Performance Problems in SQL Server* with Erin Stellato and Jonathan Kehayias
  - *Advanced SQL Server High Availability* with Brent Ozar
  - *Cortana Intelligence Suite-Microsoft R for Architects* with Buck Woody

Show focuses on versions 2008 and higher with special sessions on new features in SQL 2014 / 2016 plus sessions about Azure and vNext!



**Fantastic evening events  
interact + meet the speakers  
get YOUR questions answered!**

## ■ Industry-experts and Microsoft speakers

Sessions on performance tuning, troubleshooting, coding / development, query

tuning, architecture, new features + vNext, plus much more – learn real-world solutions & bring back **immediate ROI**

© SQLIntersection. All rights reserved.  
<http://www.SQLIntersection.com>